


Neural Networks


What are they?
How do they work?
Where am I?
Please, somebody, help me!

What is a neural network?



A neural network is just a *thing* that takes some input values and gives back some outputs

What is a neural network?




The input to a network can be anything that can be *meaningfully* converted to numbers

What is a neural network?




Like... pictures of handwritten numbers

What is a neural network?



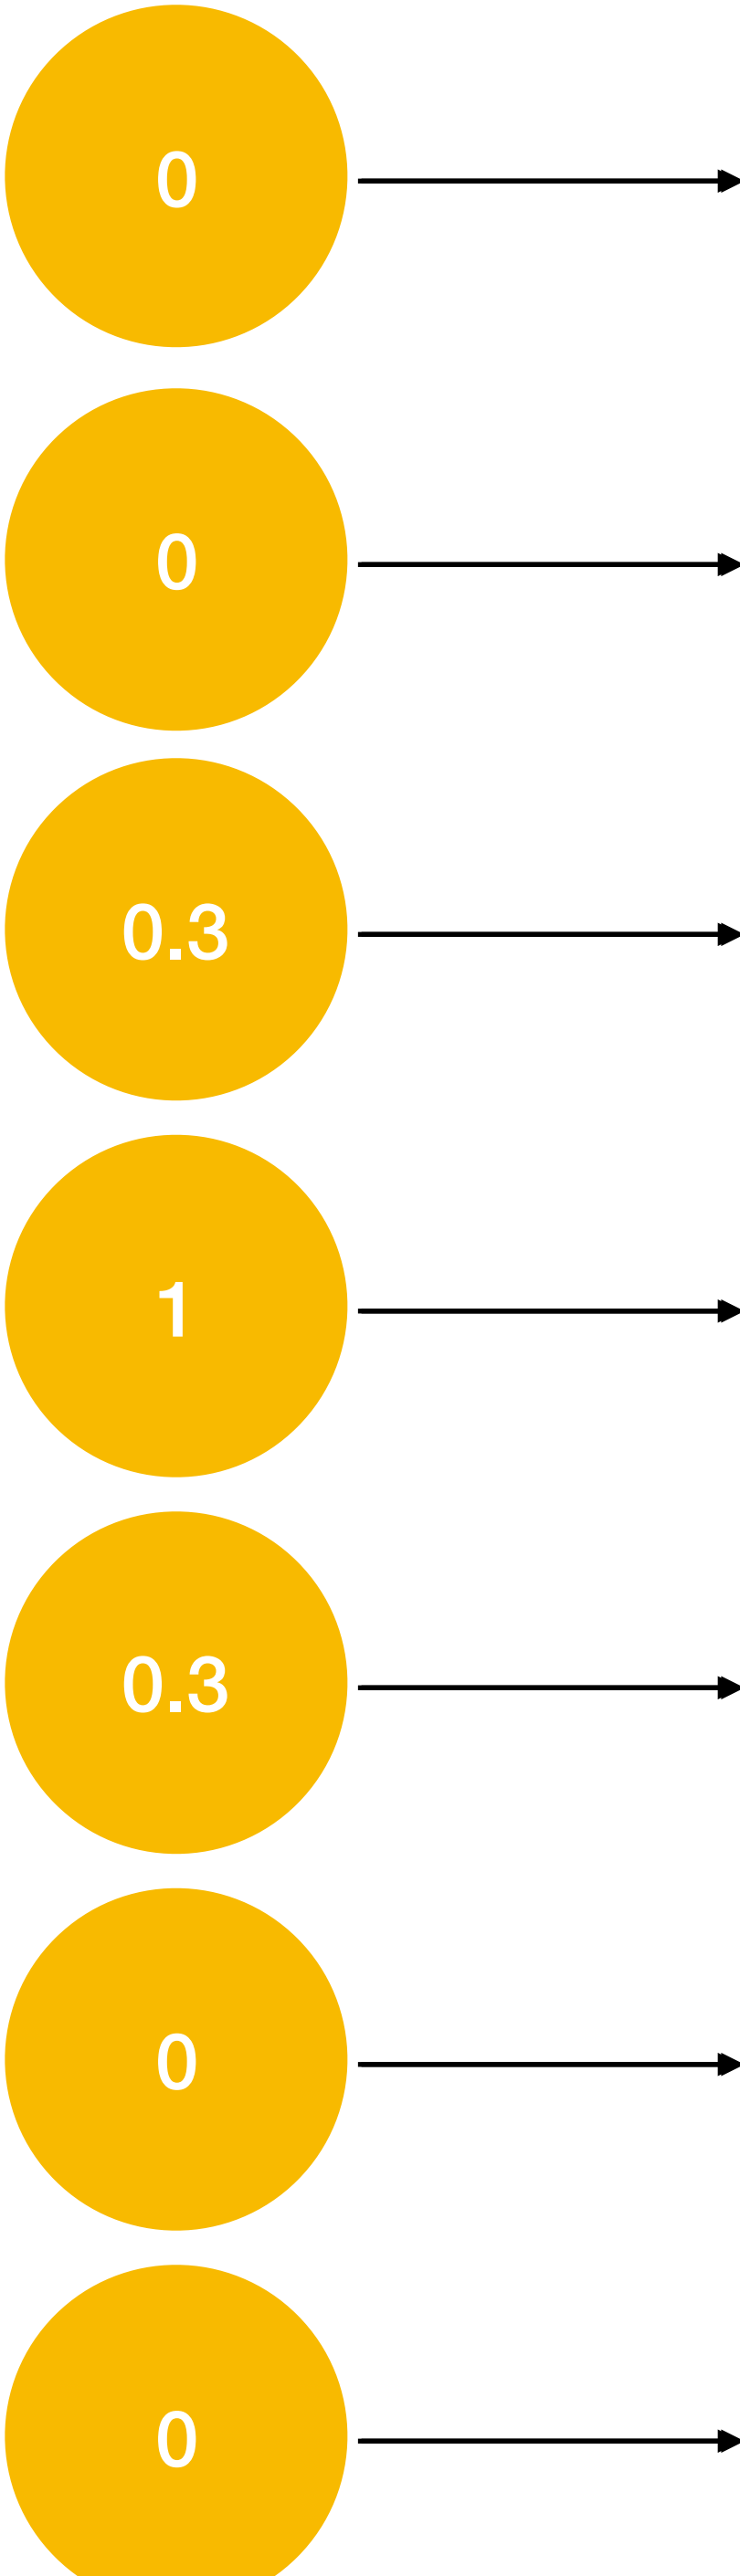
We can turn each one into values representing how dark each of the pixels are

What is a neural network?



| | | | | | | | |
|---|---|-----|-----|-----|-----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0.5 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0.7 | 0.7 | 0 | 0 |
| 0 | 0 | 0 | 0.3 | 1 | 0.3 | 0 | 0 |
| 0 | 0 | 0 | 1 | 1 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0.5 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

And then each of those values can be an input
to our network...



In this case, the goal
might be to work out
what digit we are
looking at

It's a 1

It's a 2

It's a 3

It's a 4

It's a 5

It's a 6

It's a 7...

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 |
| 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 |
| 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 |

Each output neuron can represent
one digit, and so we have 10 outputs
for 0, 1, 2, 3, all the way to 9

In this case, the goal
might be to work out
what digit we are
looking at

It's a 1

It's a 2

It's a 3

It's a 4

It's a 5

It's a 6

It's a 7...

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 |
| 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 |
| 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 |

If an output value is 1, the network
thinks it's the corresponding digit. If
it's 0, it does not.

To begin with, the network won't really have any idea which digits it's looking at, so the outputs will be all over the place



Err, it might be a "1"



Nah, maybe it's a 2.



Erm... I know this... is it a "3" maybe??



oh, oh, oh! I think it might be a 4?!



0.4



0.6



0.2



you have to train it by
telling it what's
what...

0.3

0.5

0.3

0.7


0.4

0.6

0.2

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 |
| 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 |
| 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 |

It's a "1"!!!



And by doing so, it
gradually learns..

0.9

I am pretty damn
sure it's a "1".

0.2

There's a wee chance that
it might be a "2" though..

0

OK, it's DEFINITELY not a
"3"

0

0


It's a "1"!!!

0

Umm, but it could be a "7".
Damn these digits for looking so
similar..silly humans

0.5

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 |
| 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 |
| 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 |



...until eventually, it
has learned what the
digits are

0.9

0.01

0

0

0


0

0.02


Defo a "1"

| | | | | | | | | | |
|---|---|---|---|---|---|---|---|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 |
| 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 | 2 |
| 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 | 3 |
| 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 | 4 |
| 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 | 5 |
| 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 | 6 |
| 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 | 7 |
| 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 | 8 |
| 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 | 9 |

Woohoo! Good job!




What is a neural network?



So, a network takes in some inputs that represent something, and can be trained to output appropriate values as a result.

What is a neural network?



In fact, a neural network can, in theory, represent any mapping from input to output values!

- "universal approximation theorem"

It can learn to map pixels representing handwritten numbers to their value...




“I think it’s a 1”

“I think it’s a 3”

“I think it’s a 7”

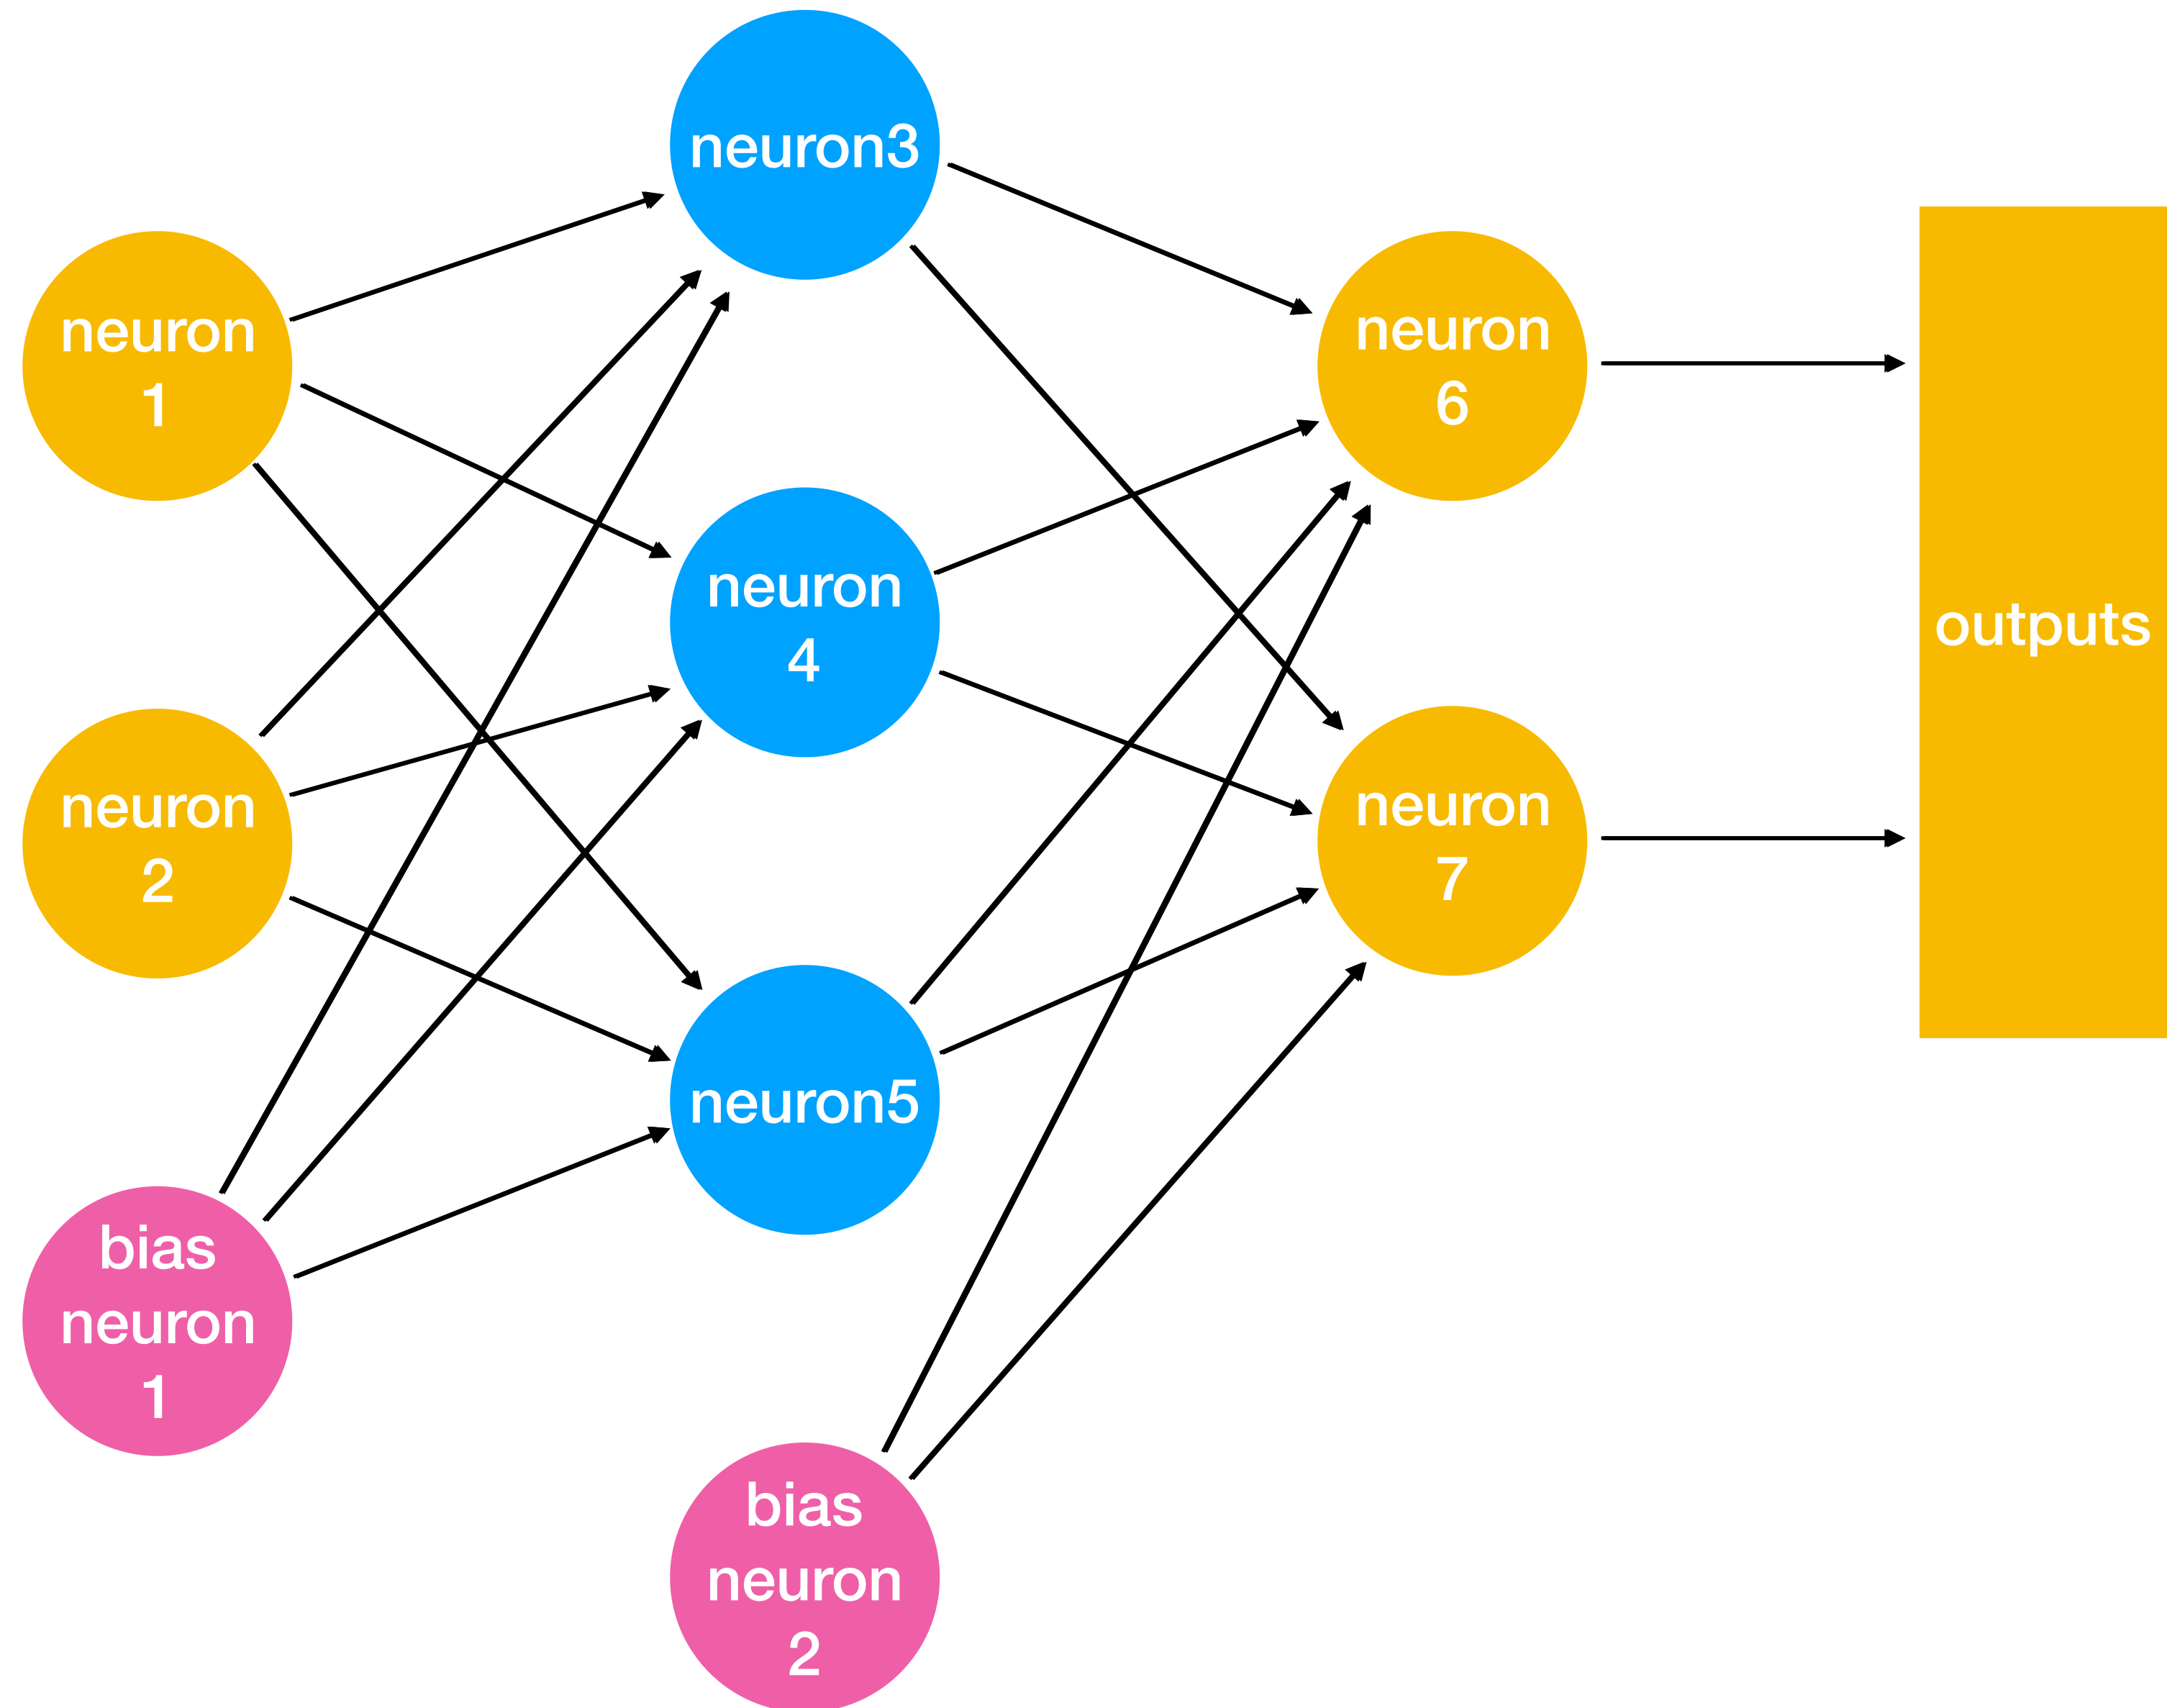
It can learn to map black and white pictures to colour pictures...



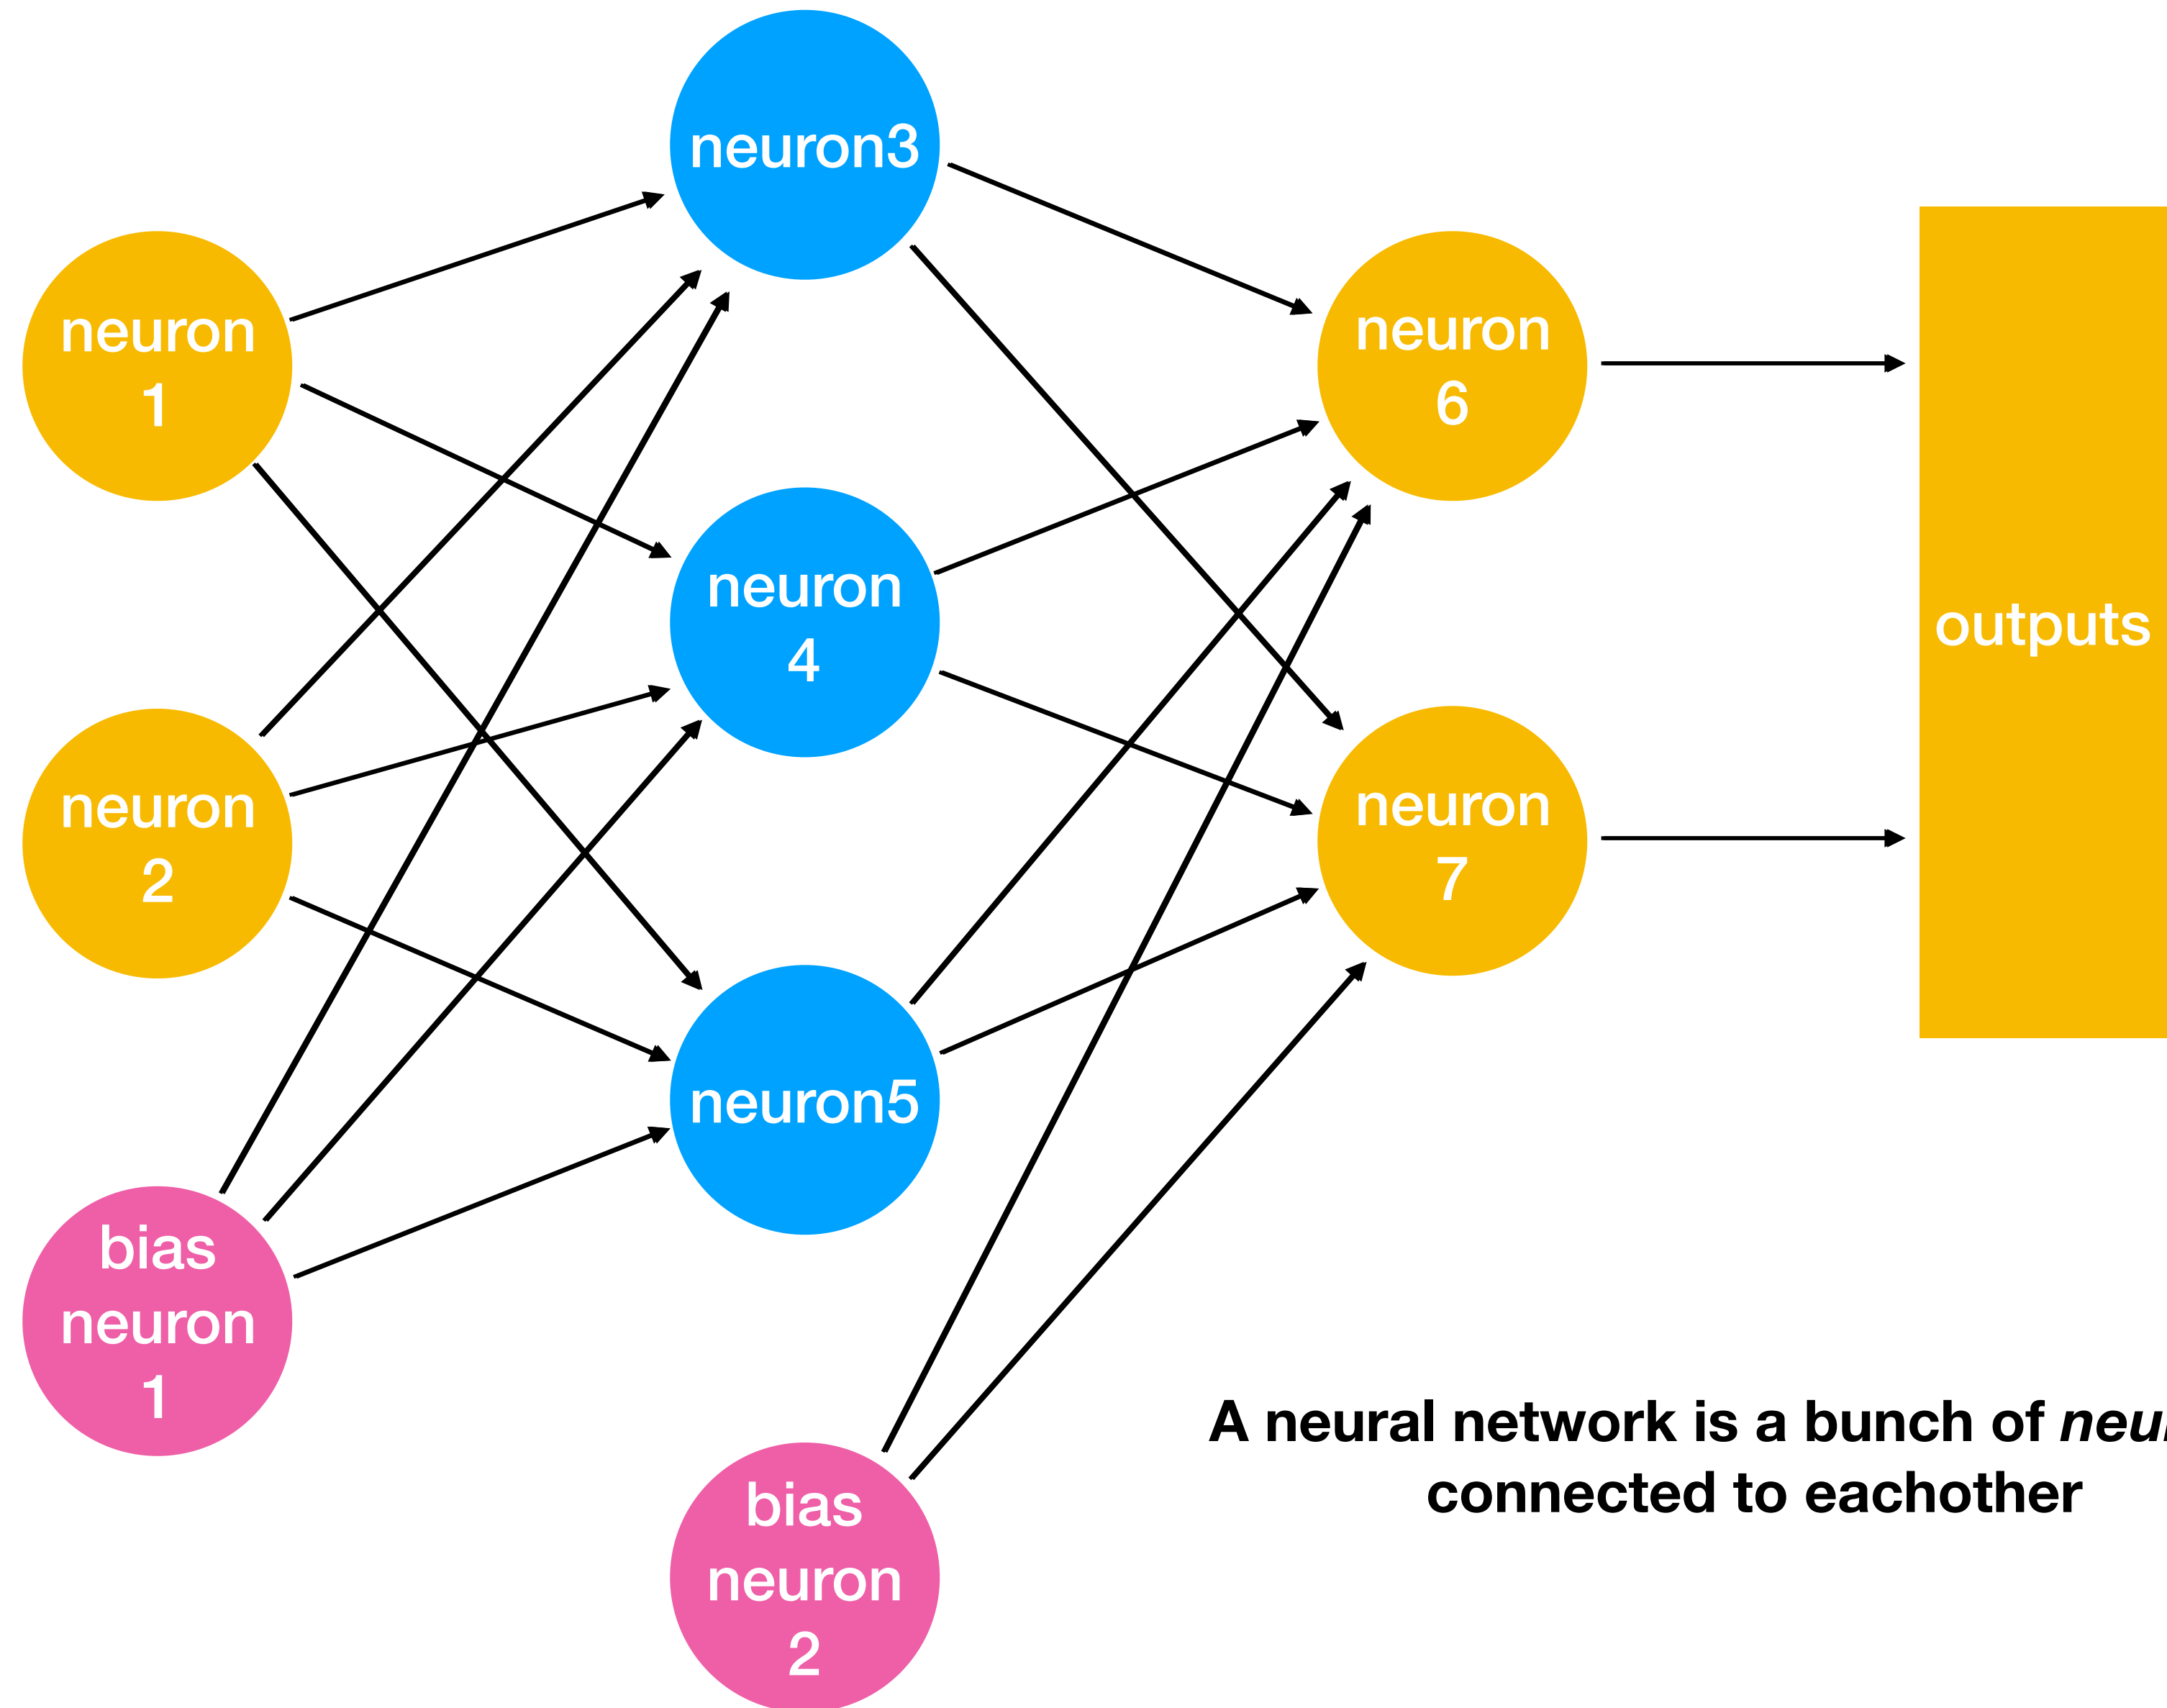
And loads of other things!

<http://www.yaronhadad.com/deep-learning-most-amazing-applications/>

Let's peel away the black box...

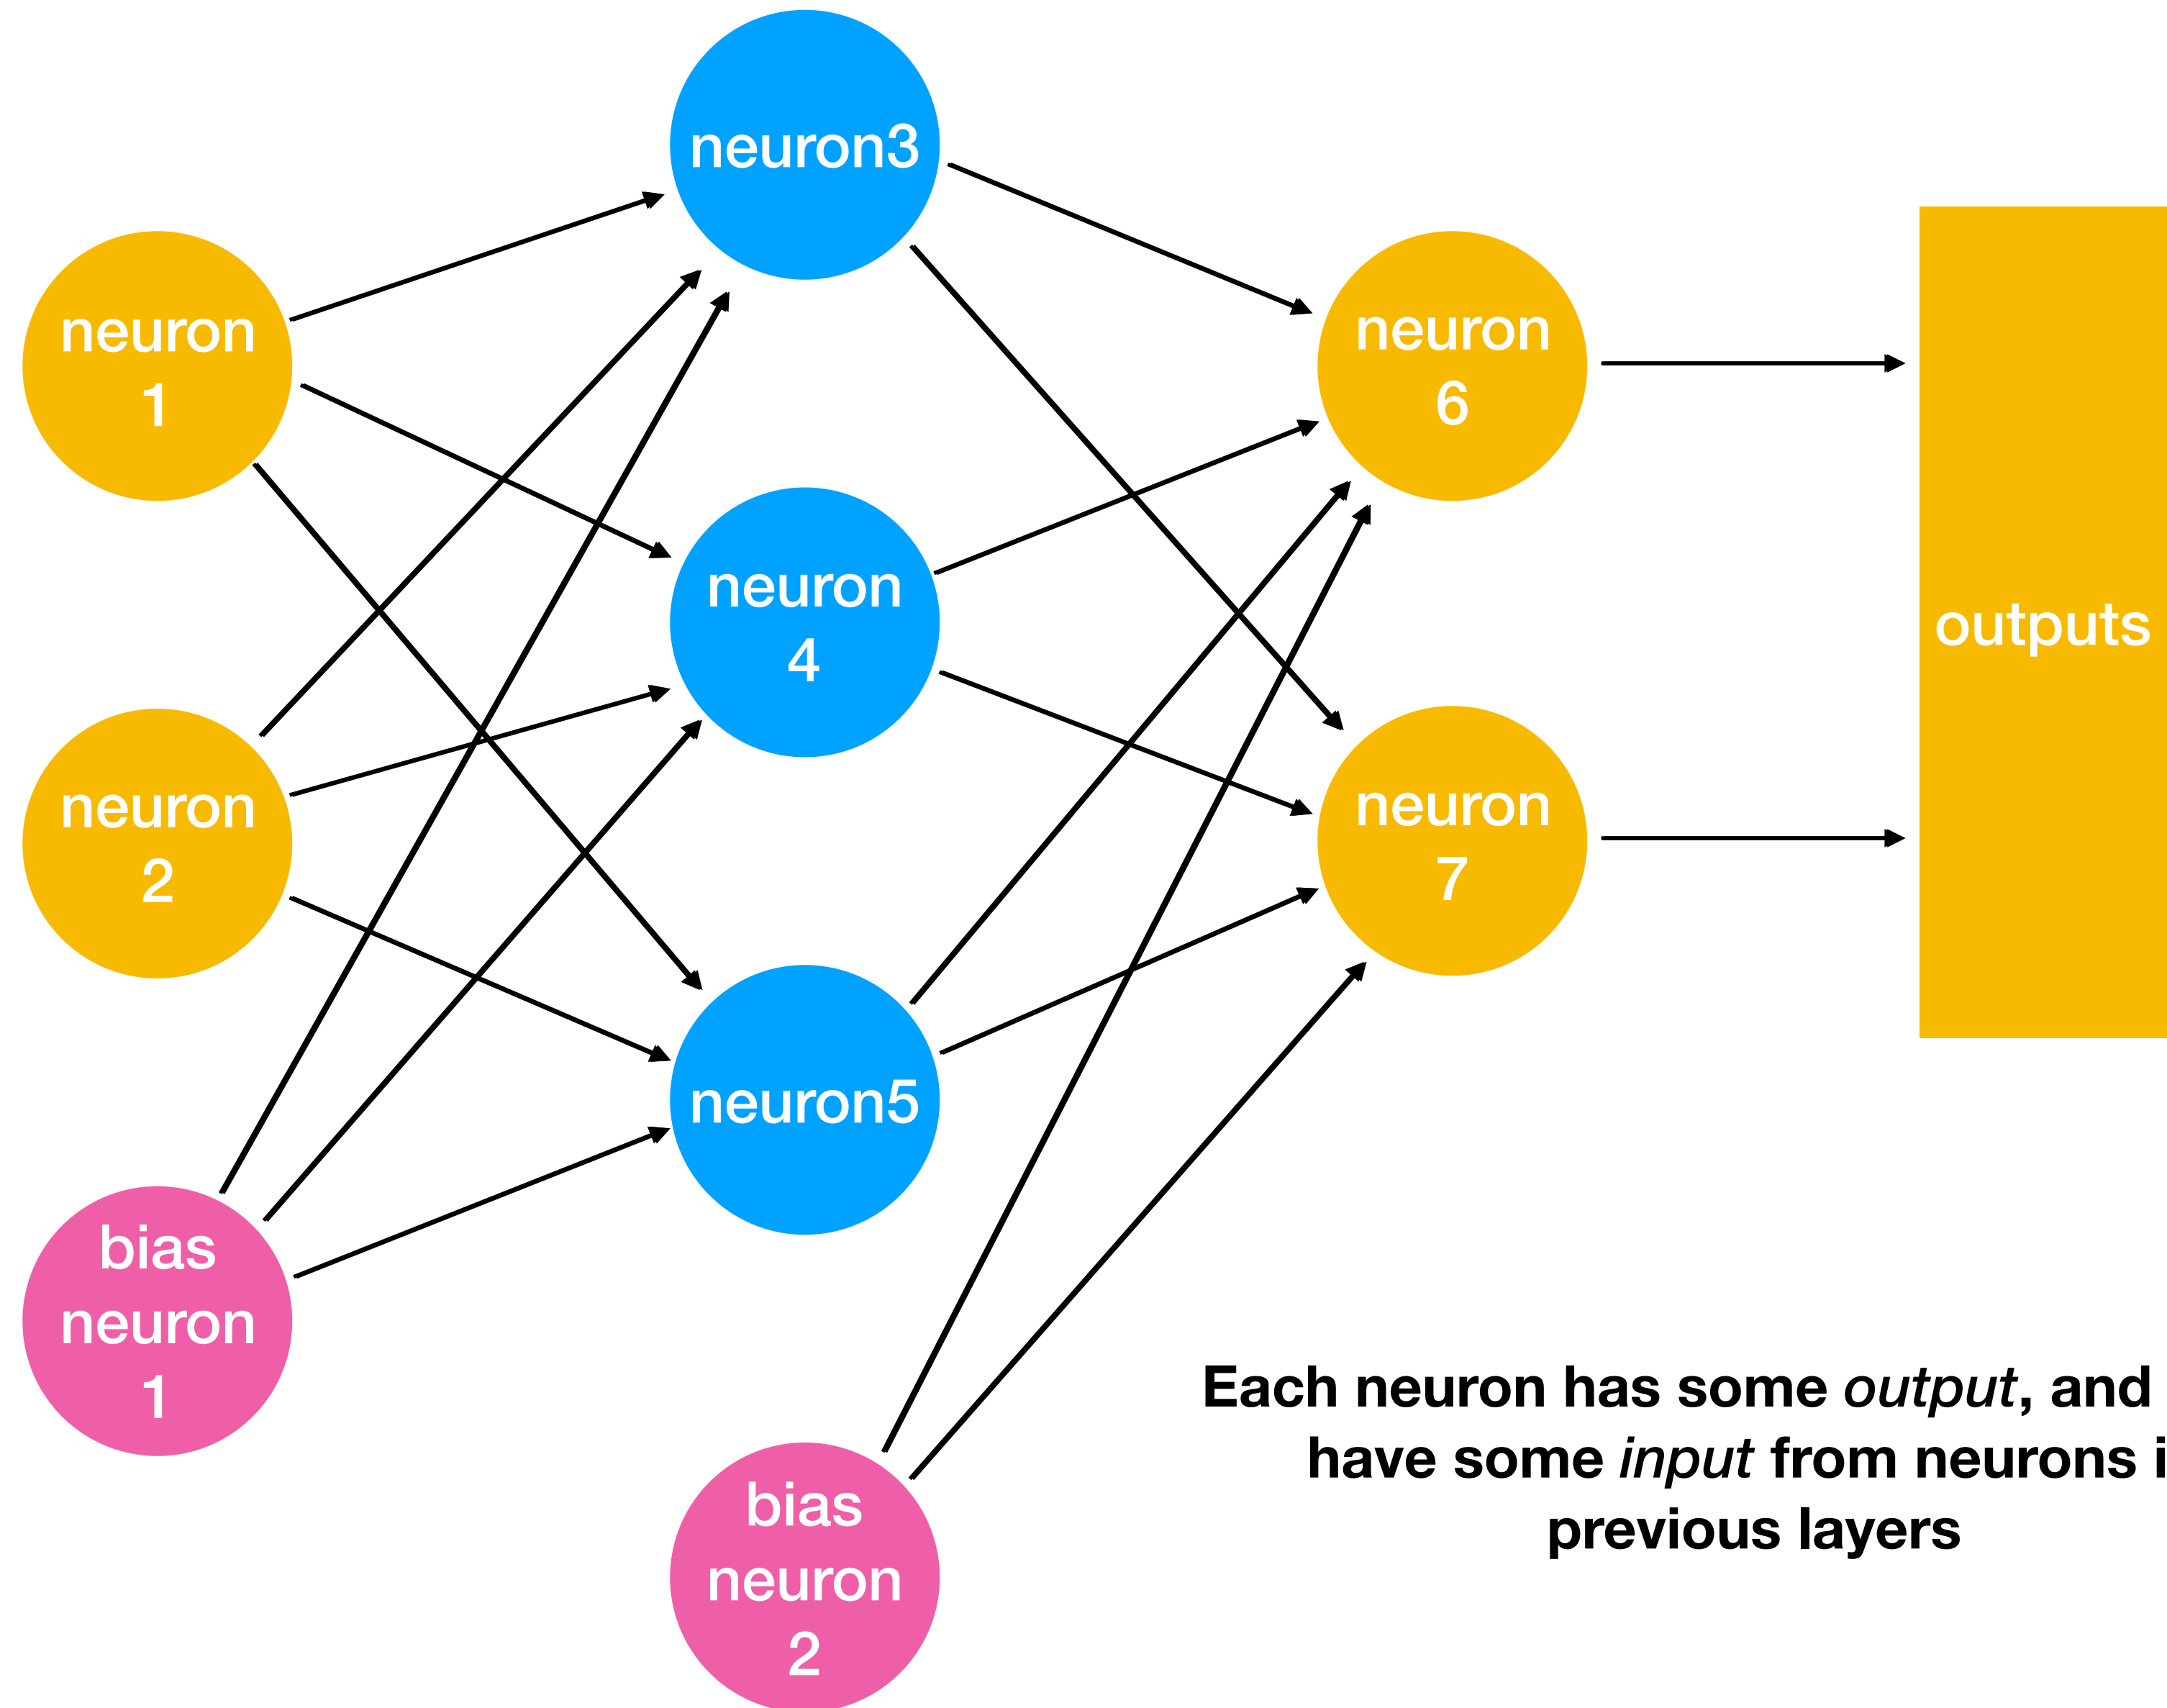


Let's peel away the black box...



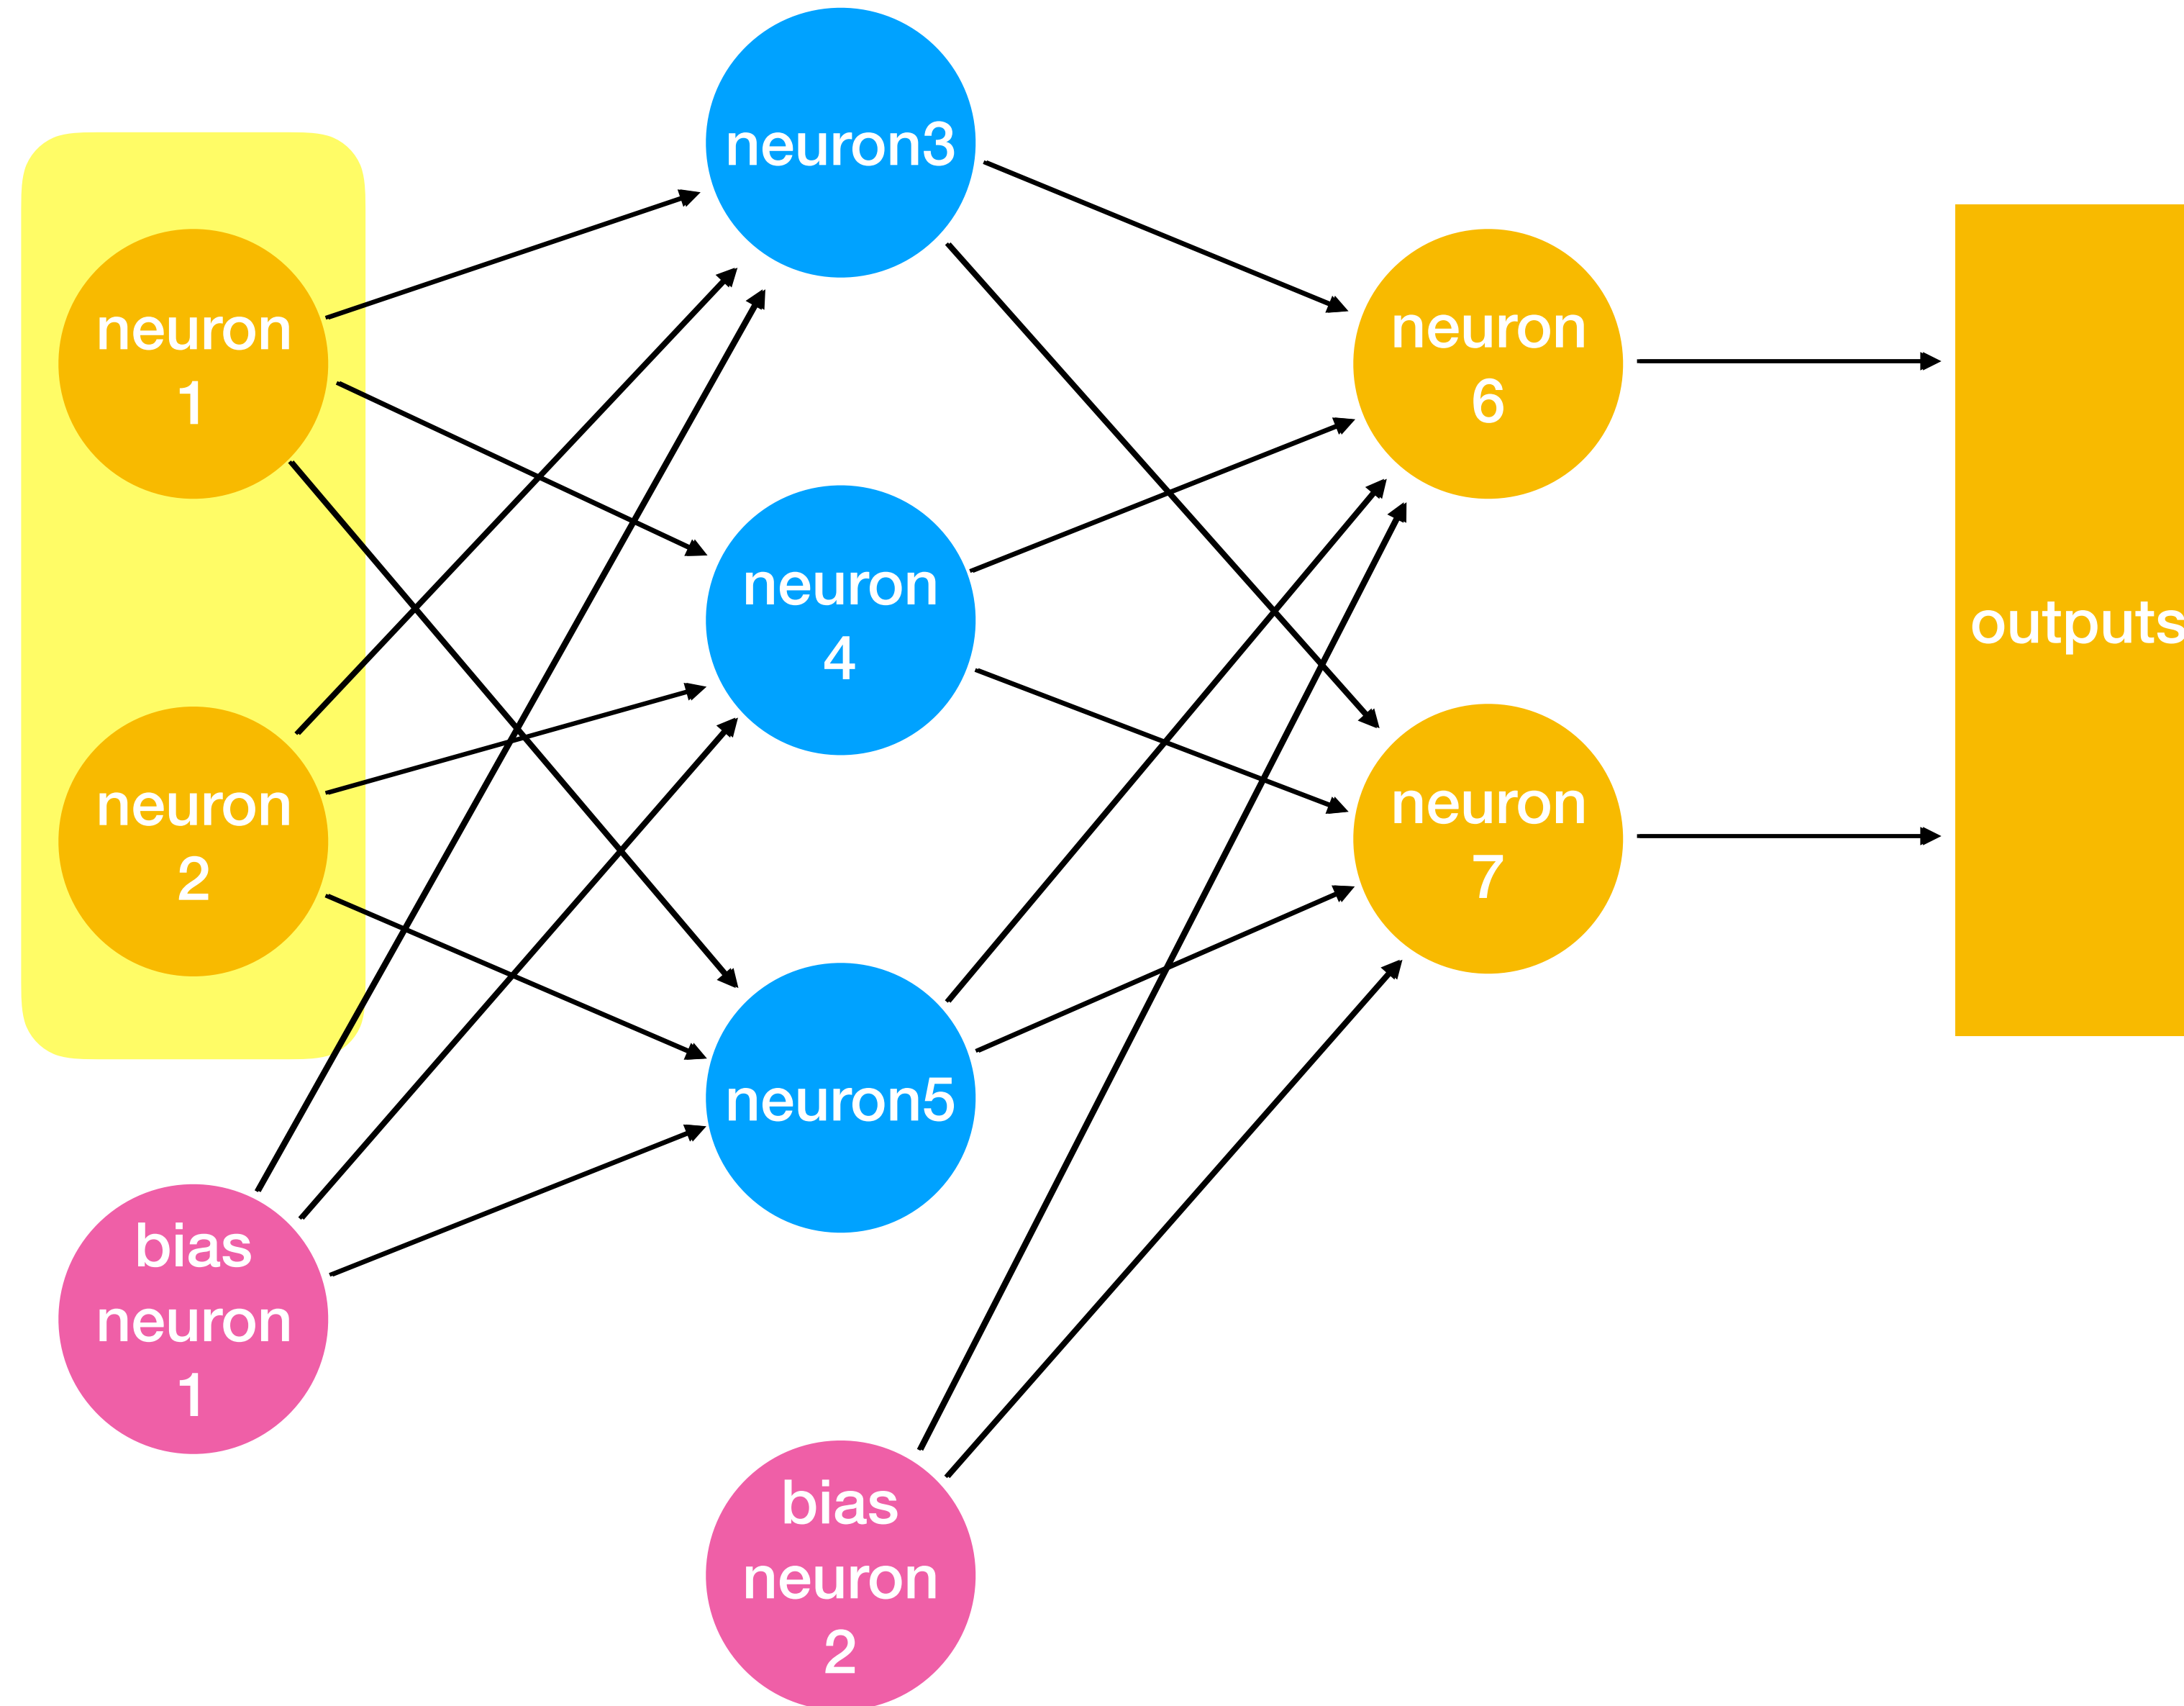
A neural network is a bunch of *neurons*, connected to each other

Let's peel away the black box...

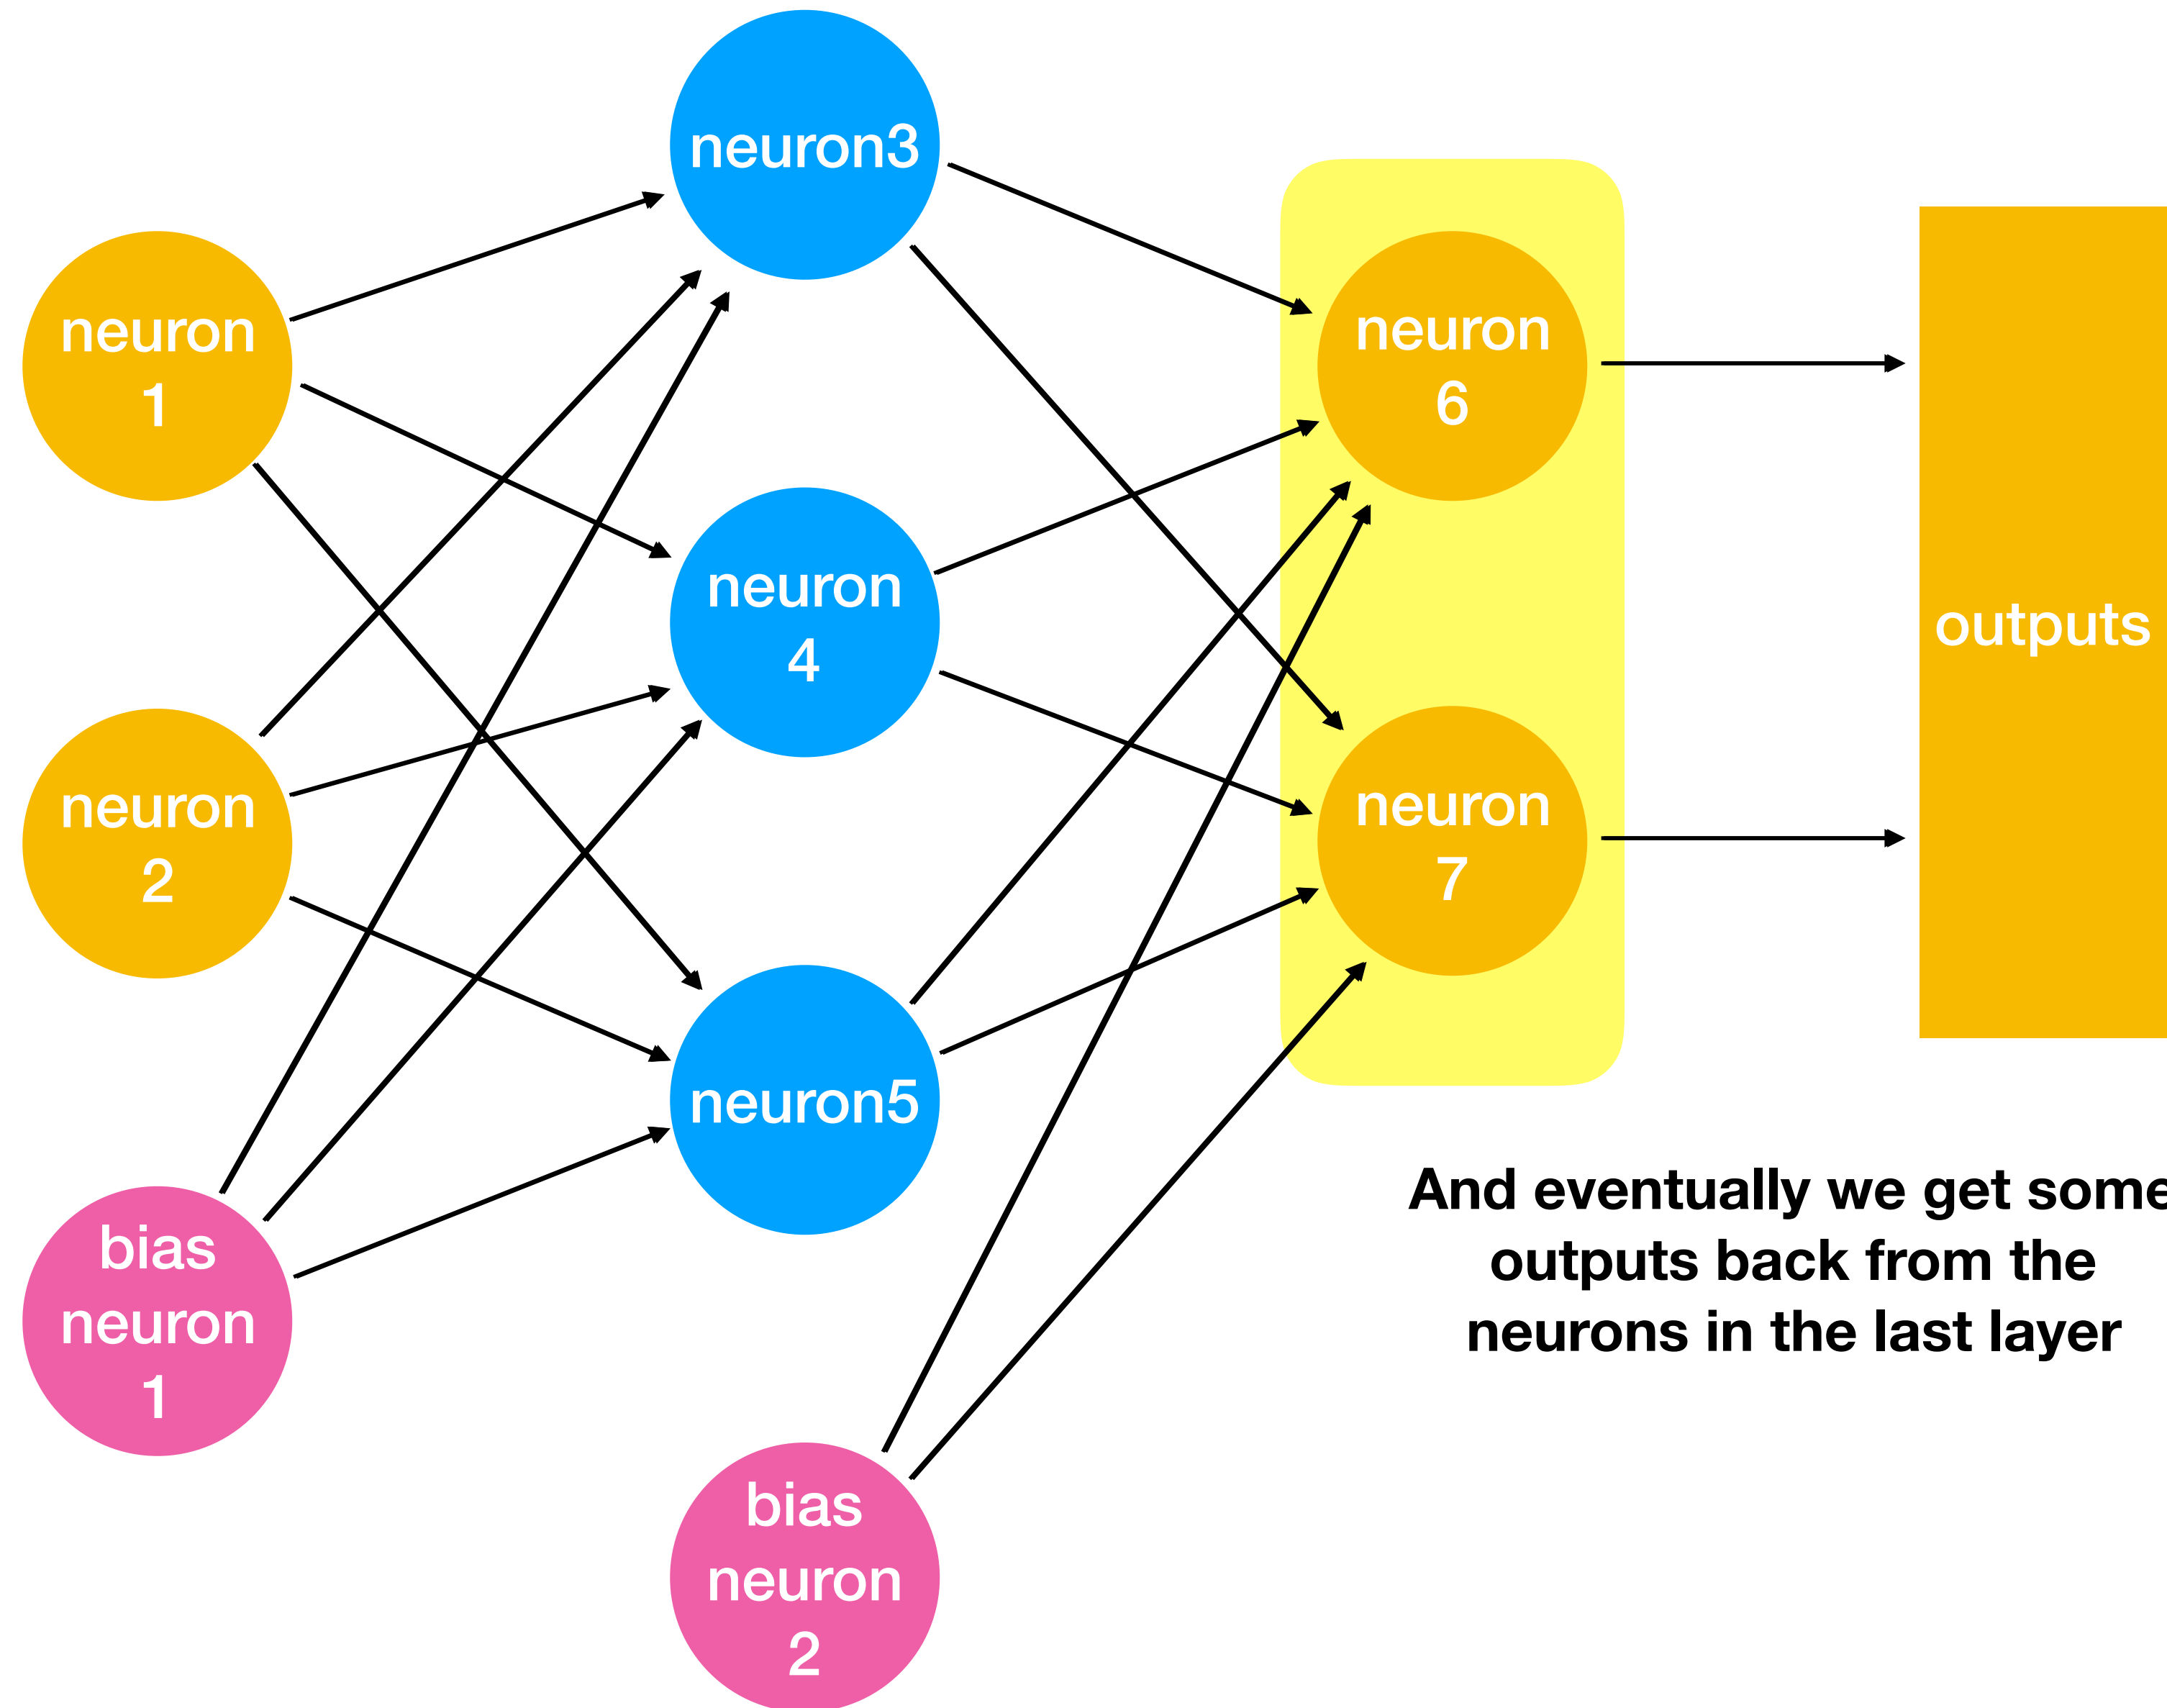


Let's peel away the black box...

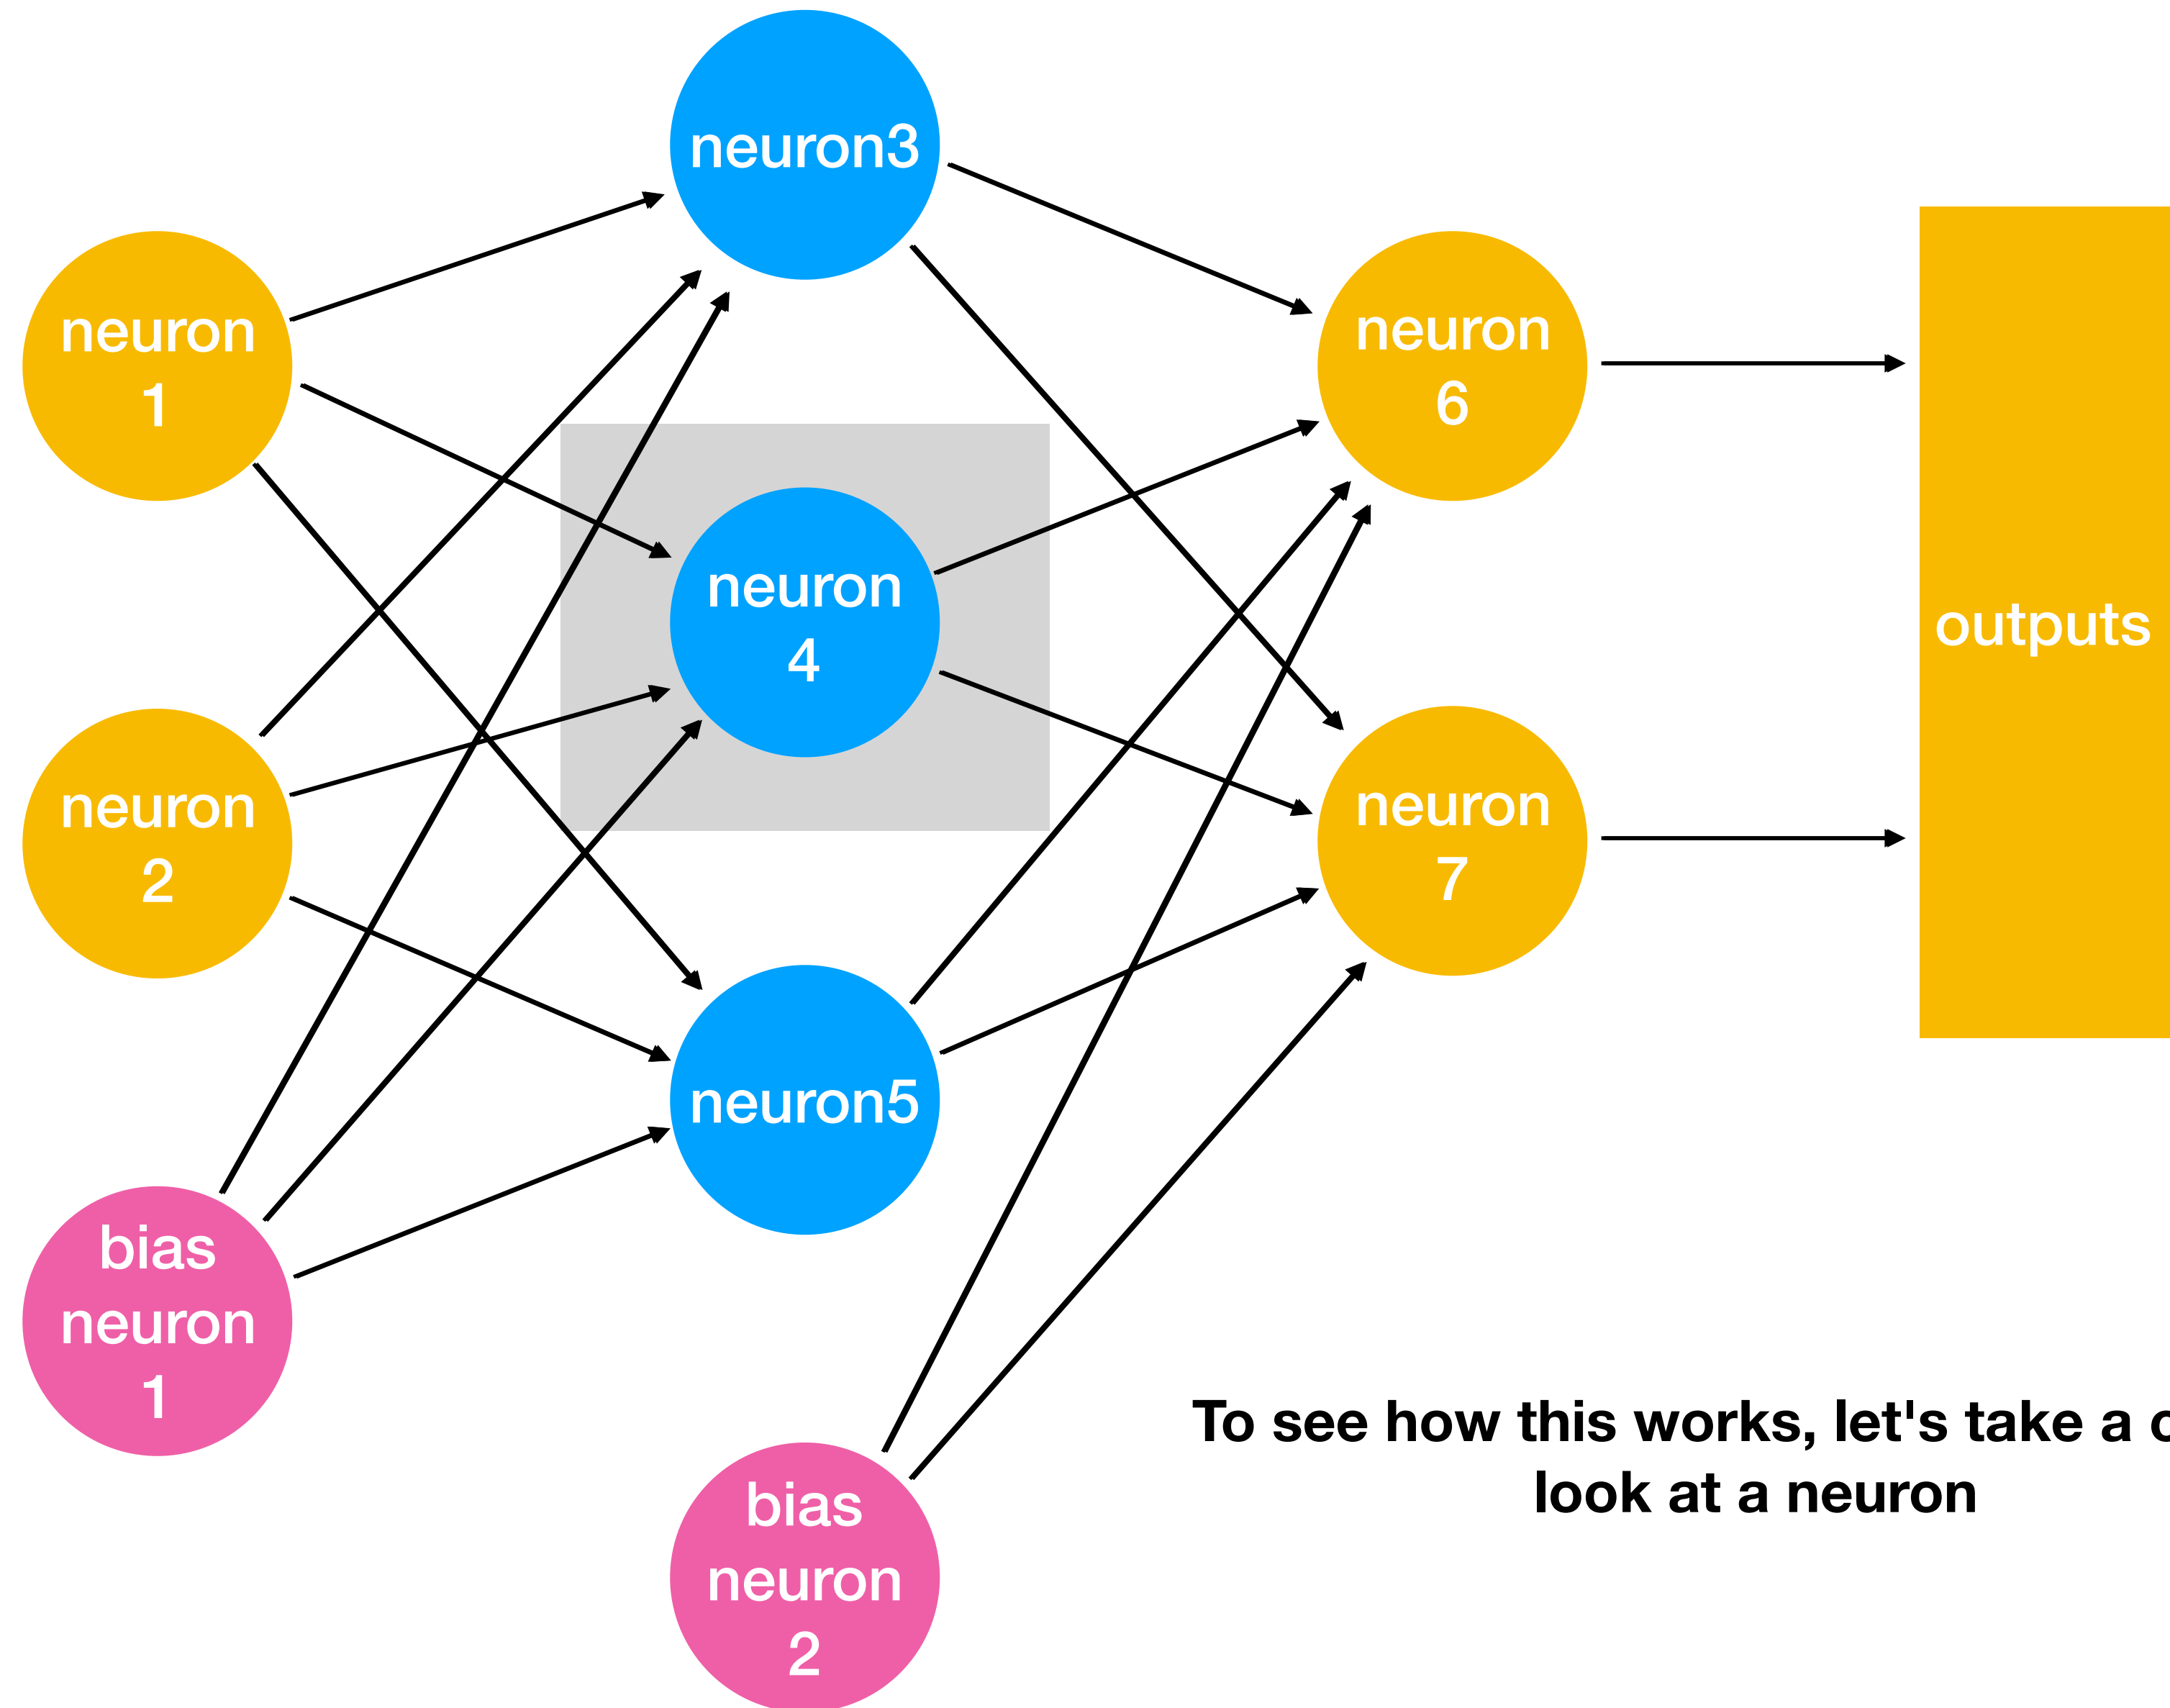
We set the neurons on the first (leftmost) layer to our input values



Let's peel away the black box...

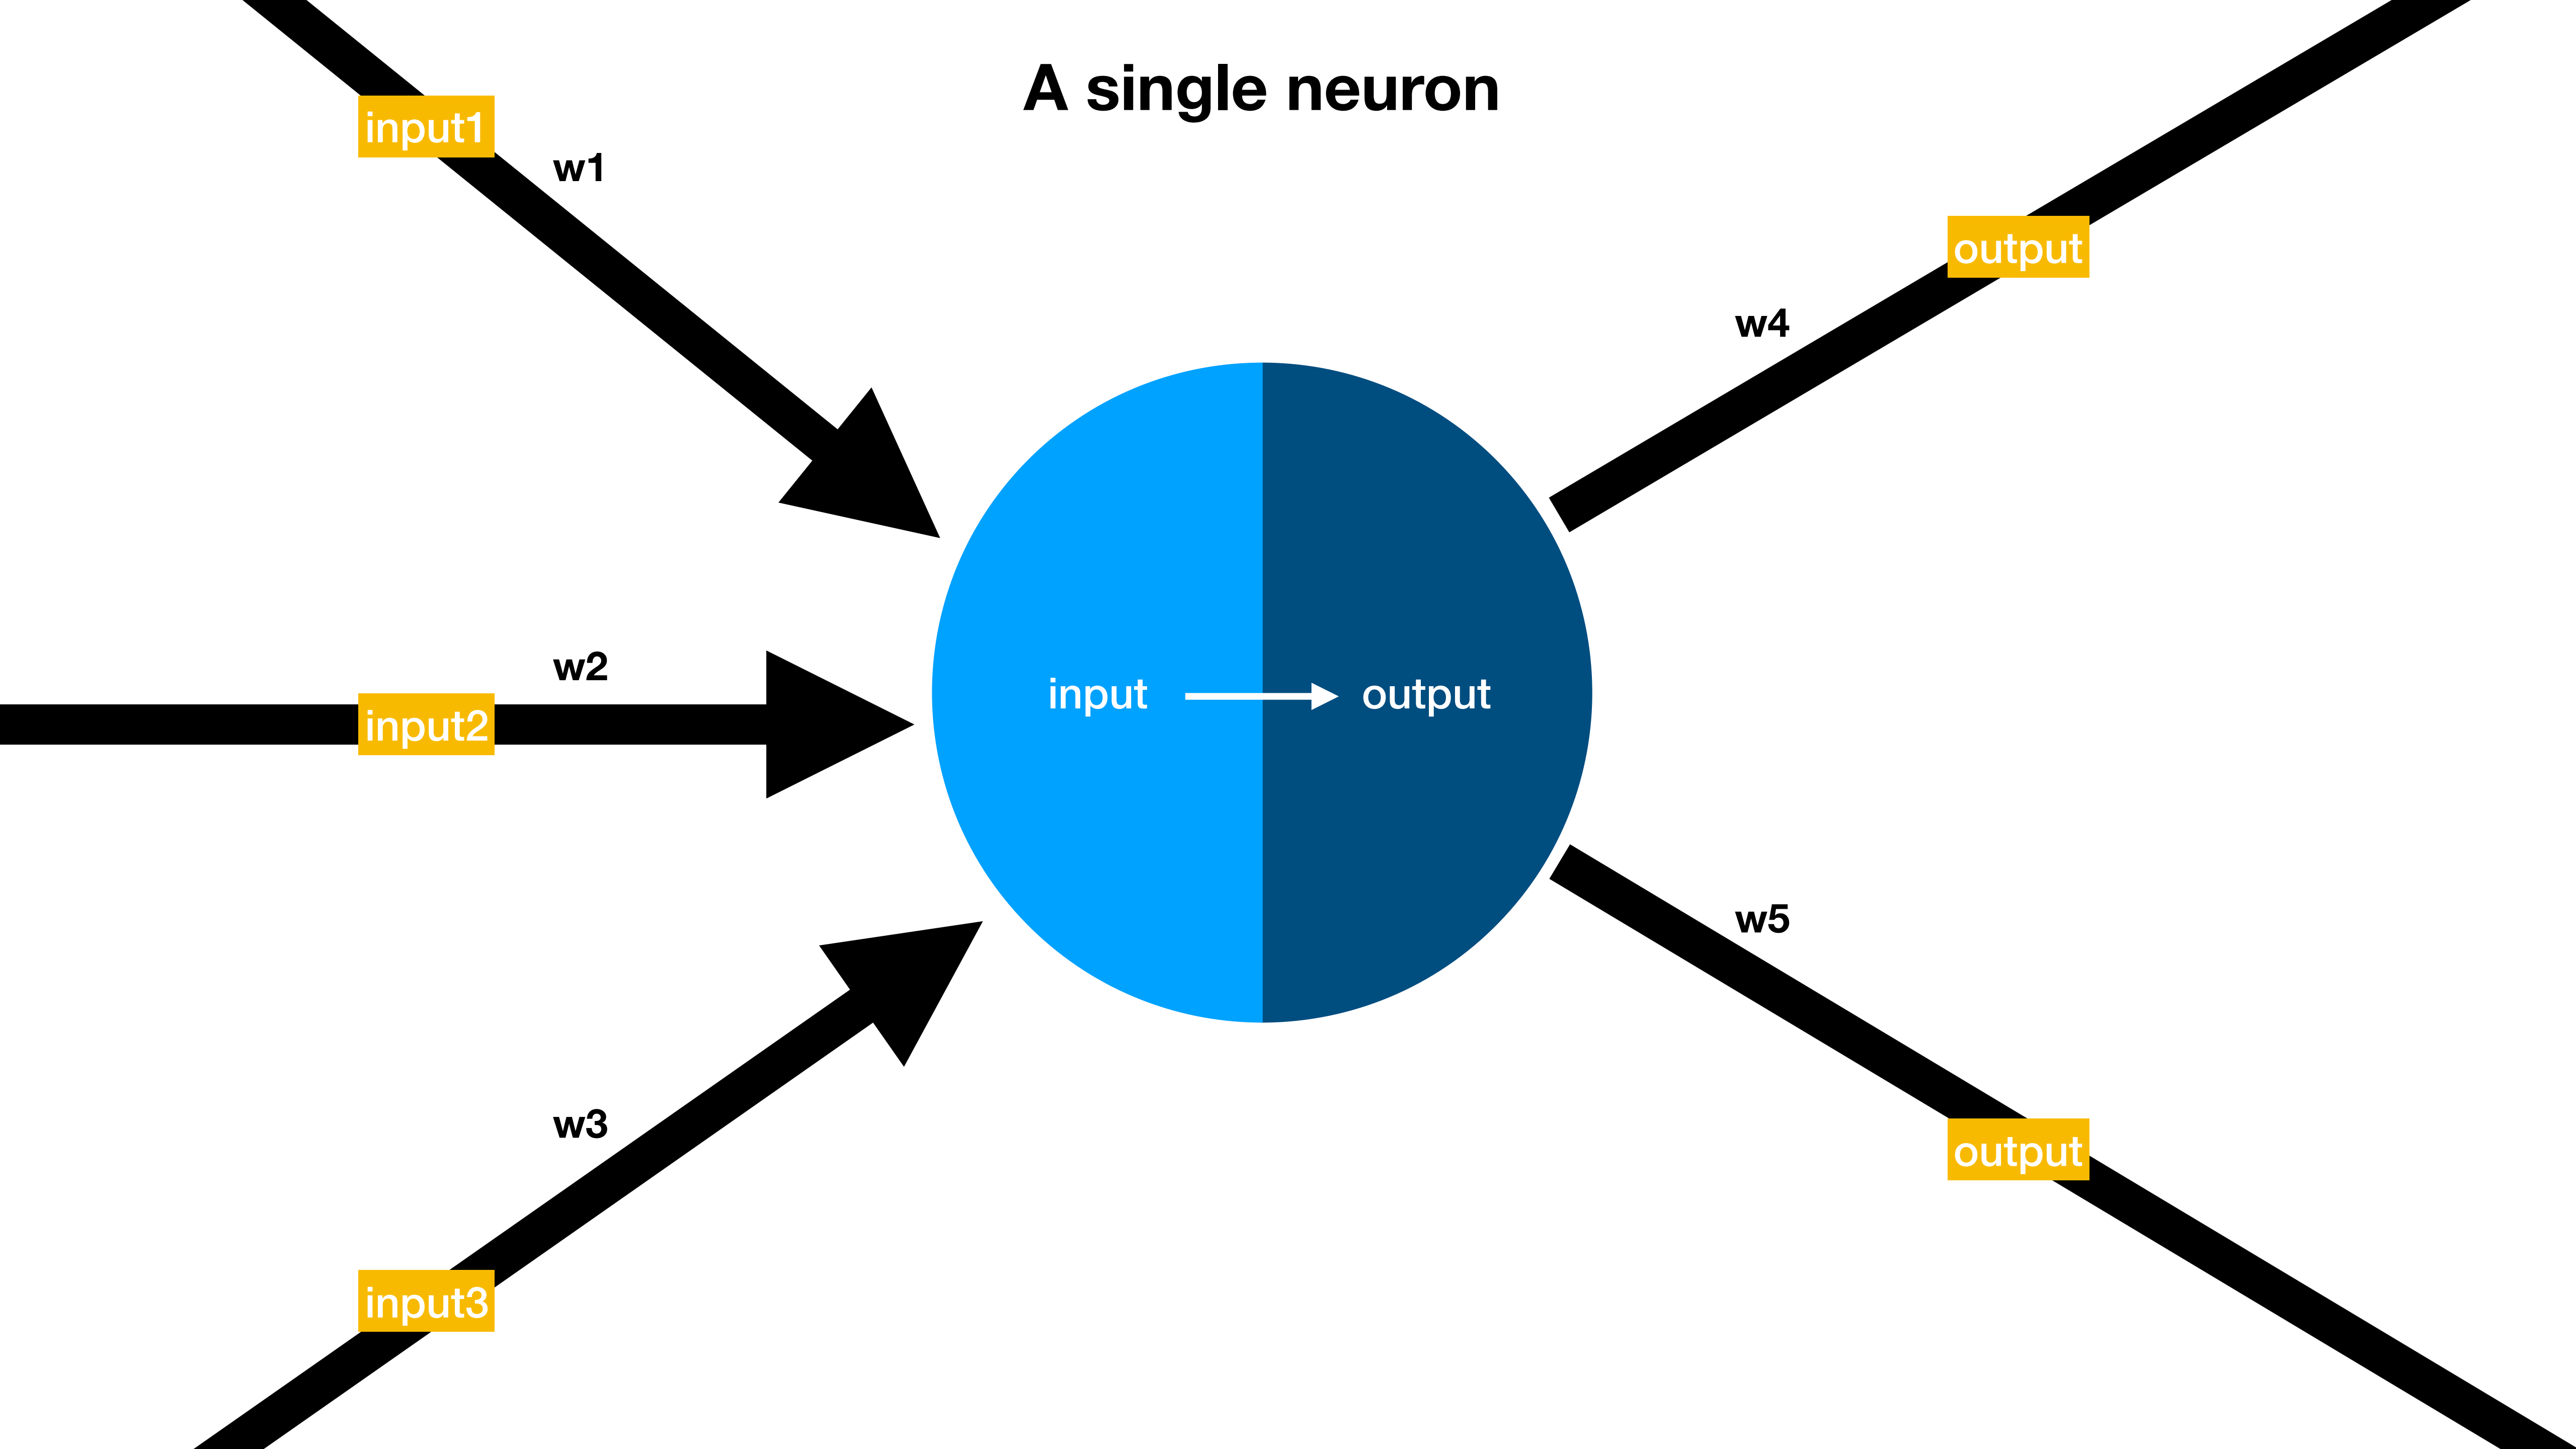


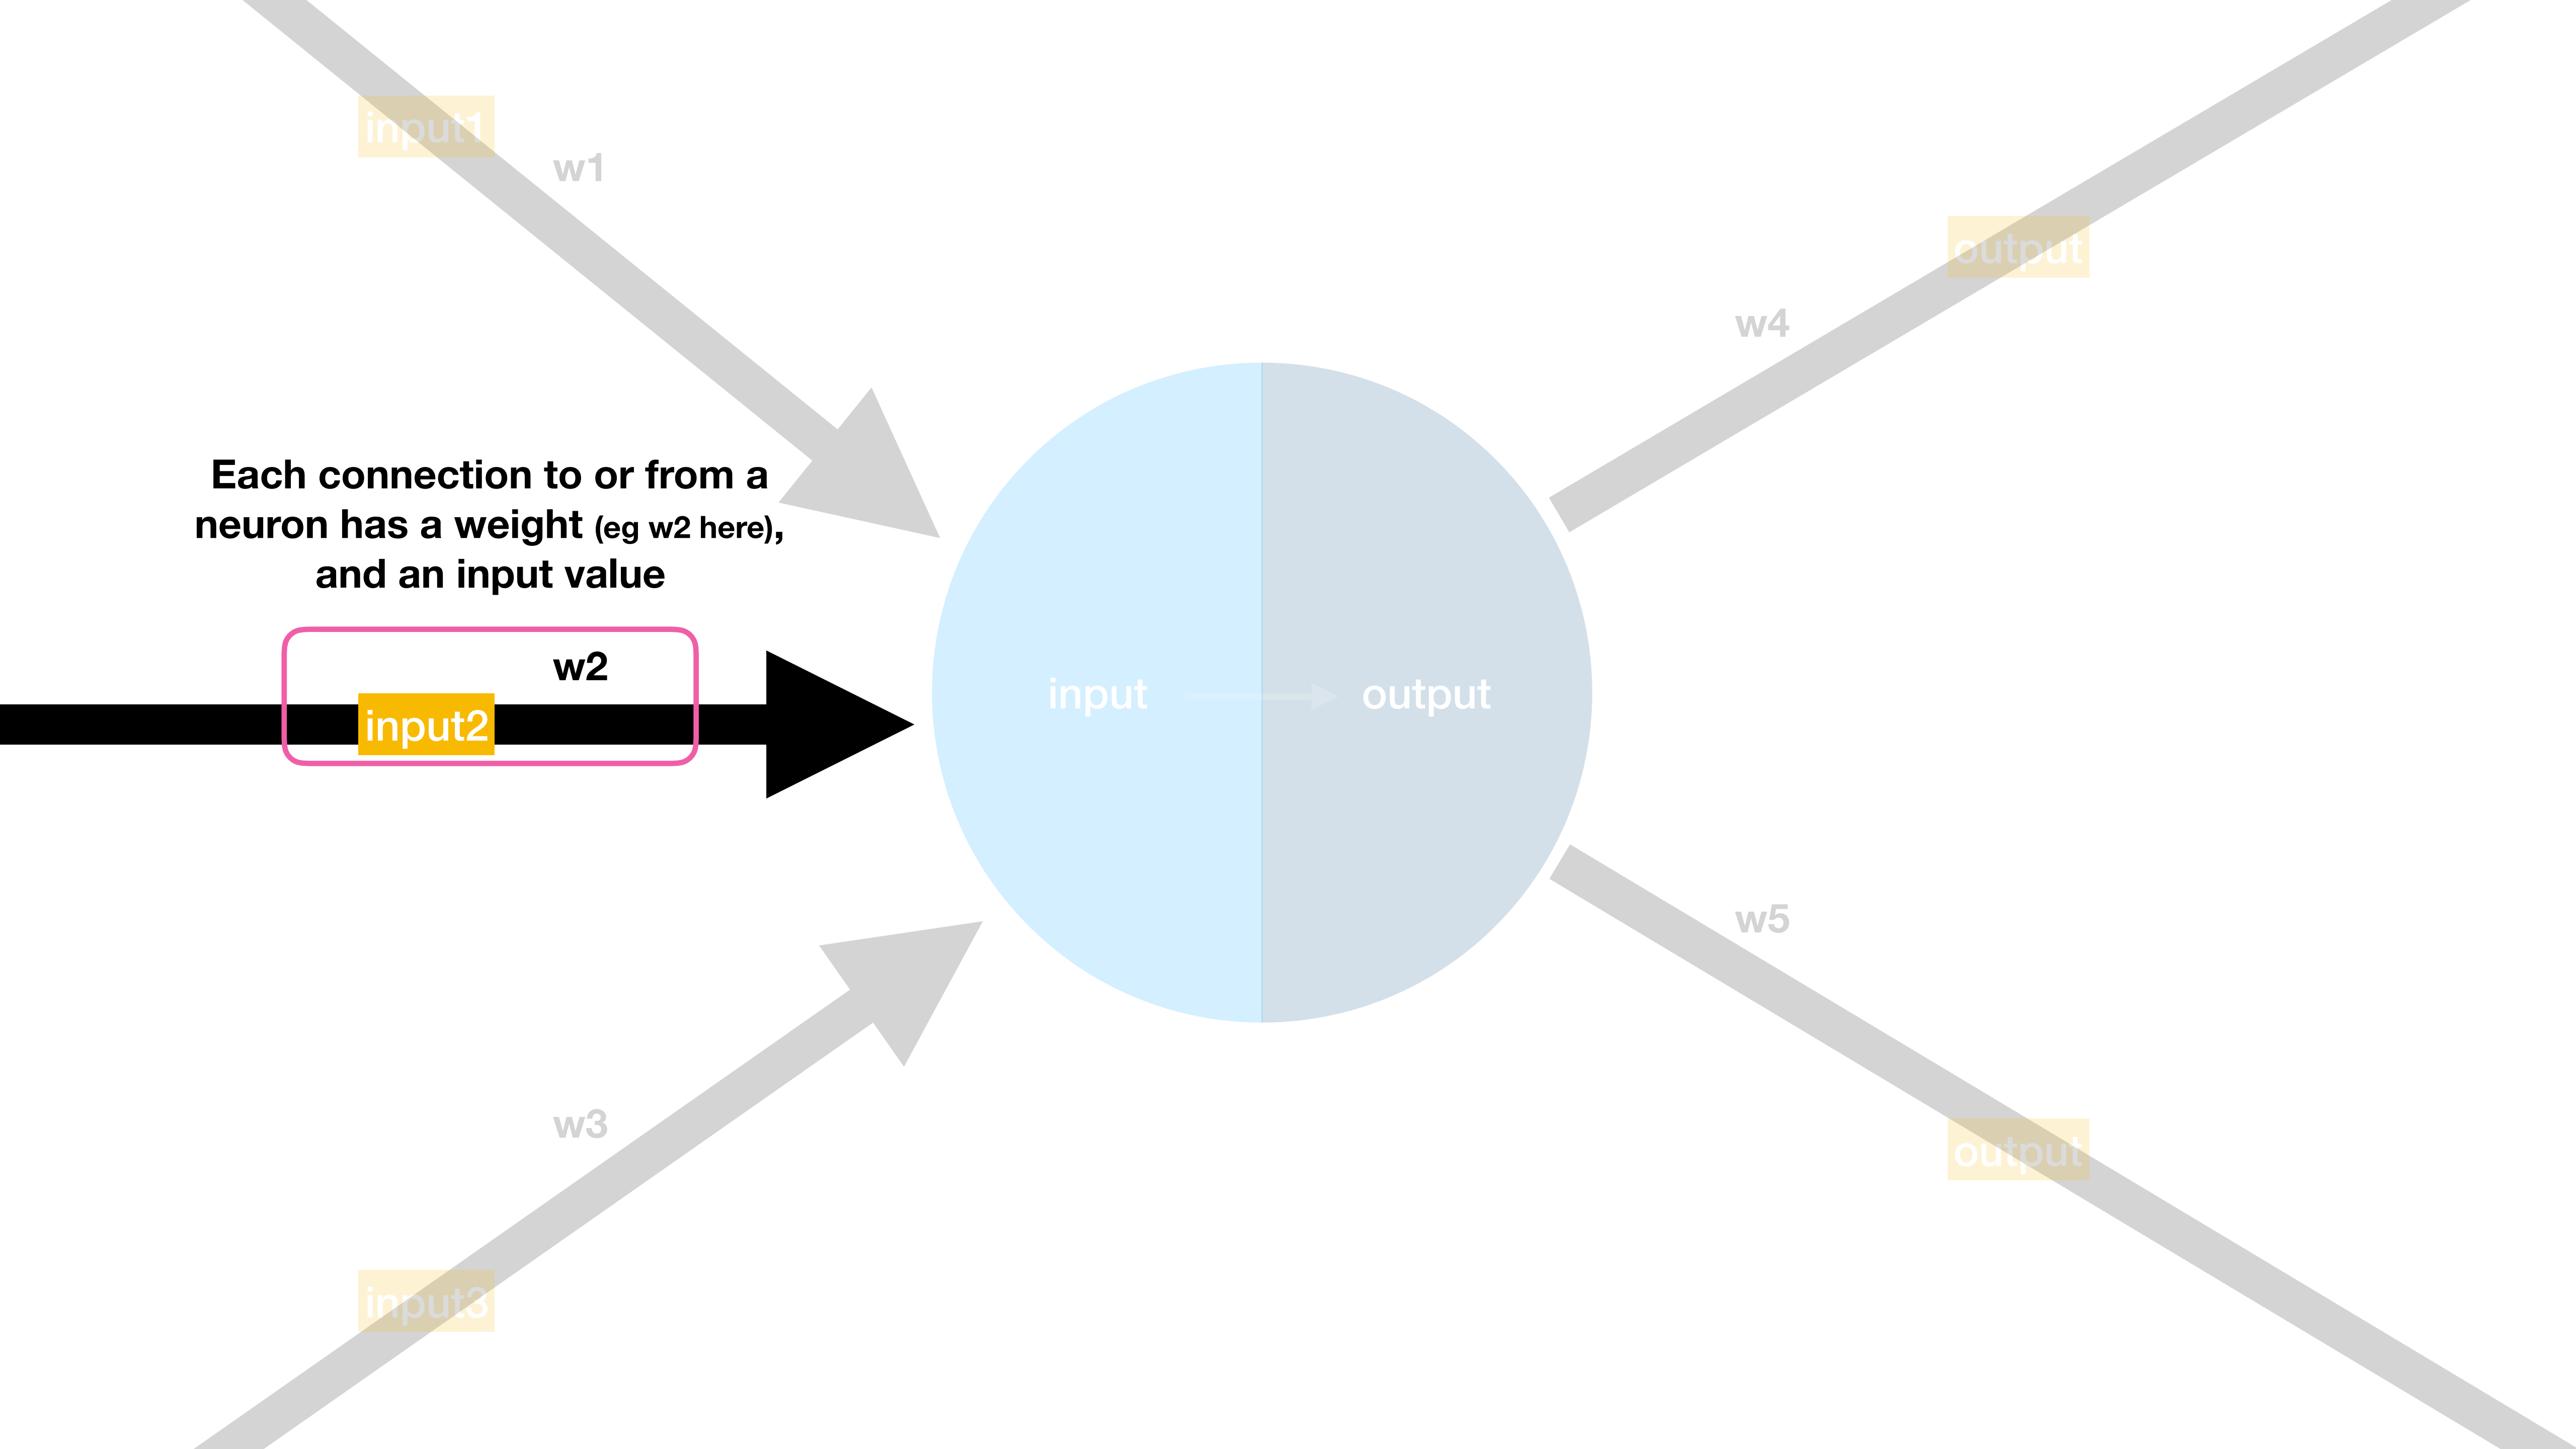
And eventually we get some
outputs back from the
neurons in the last layer



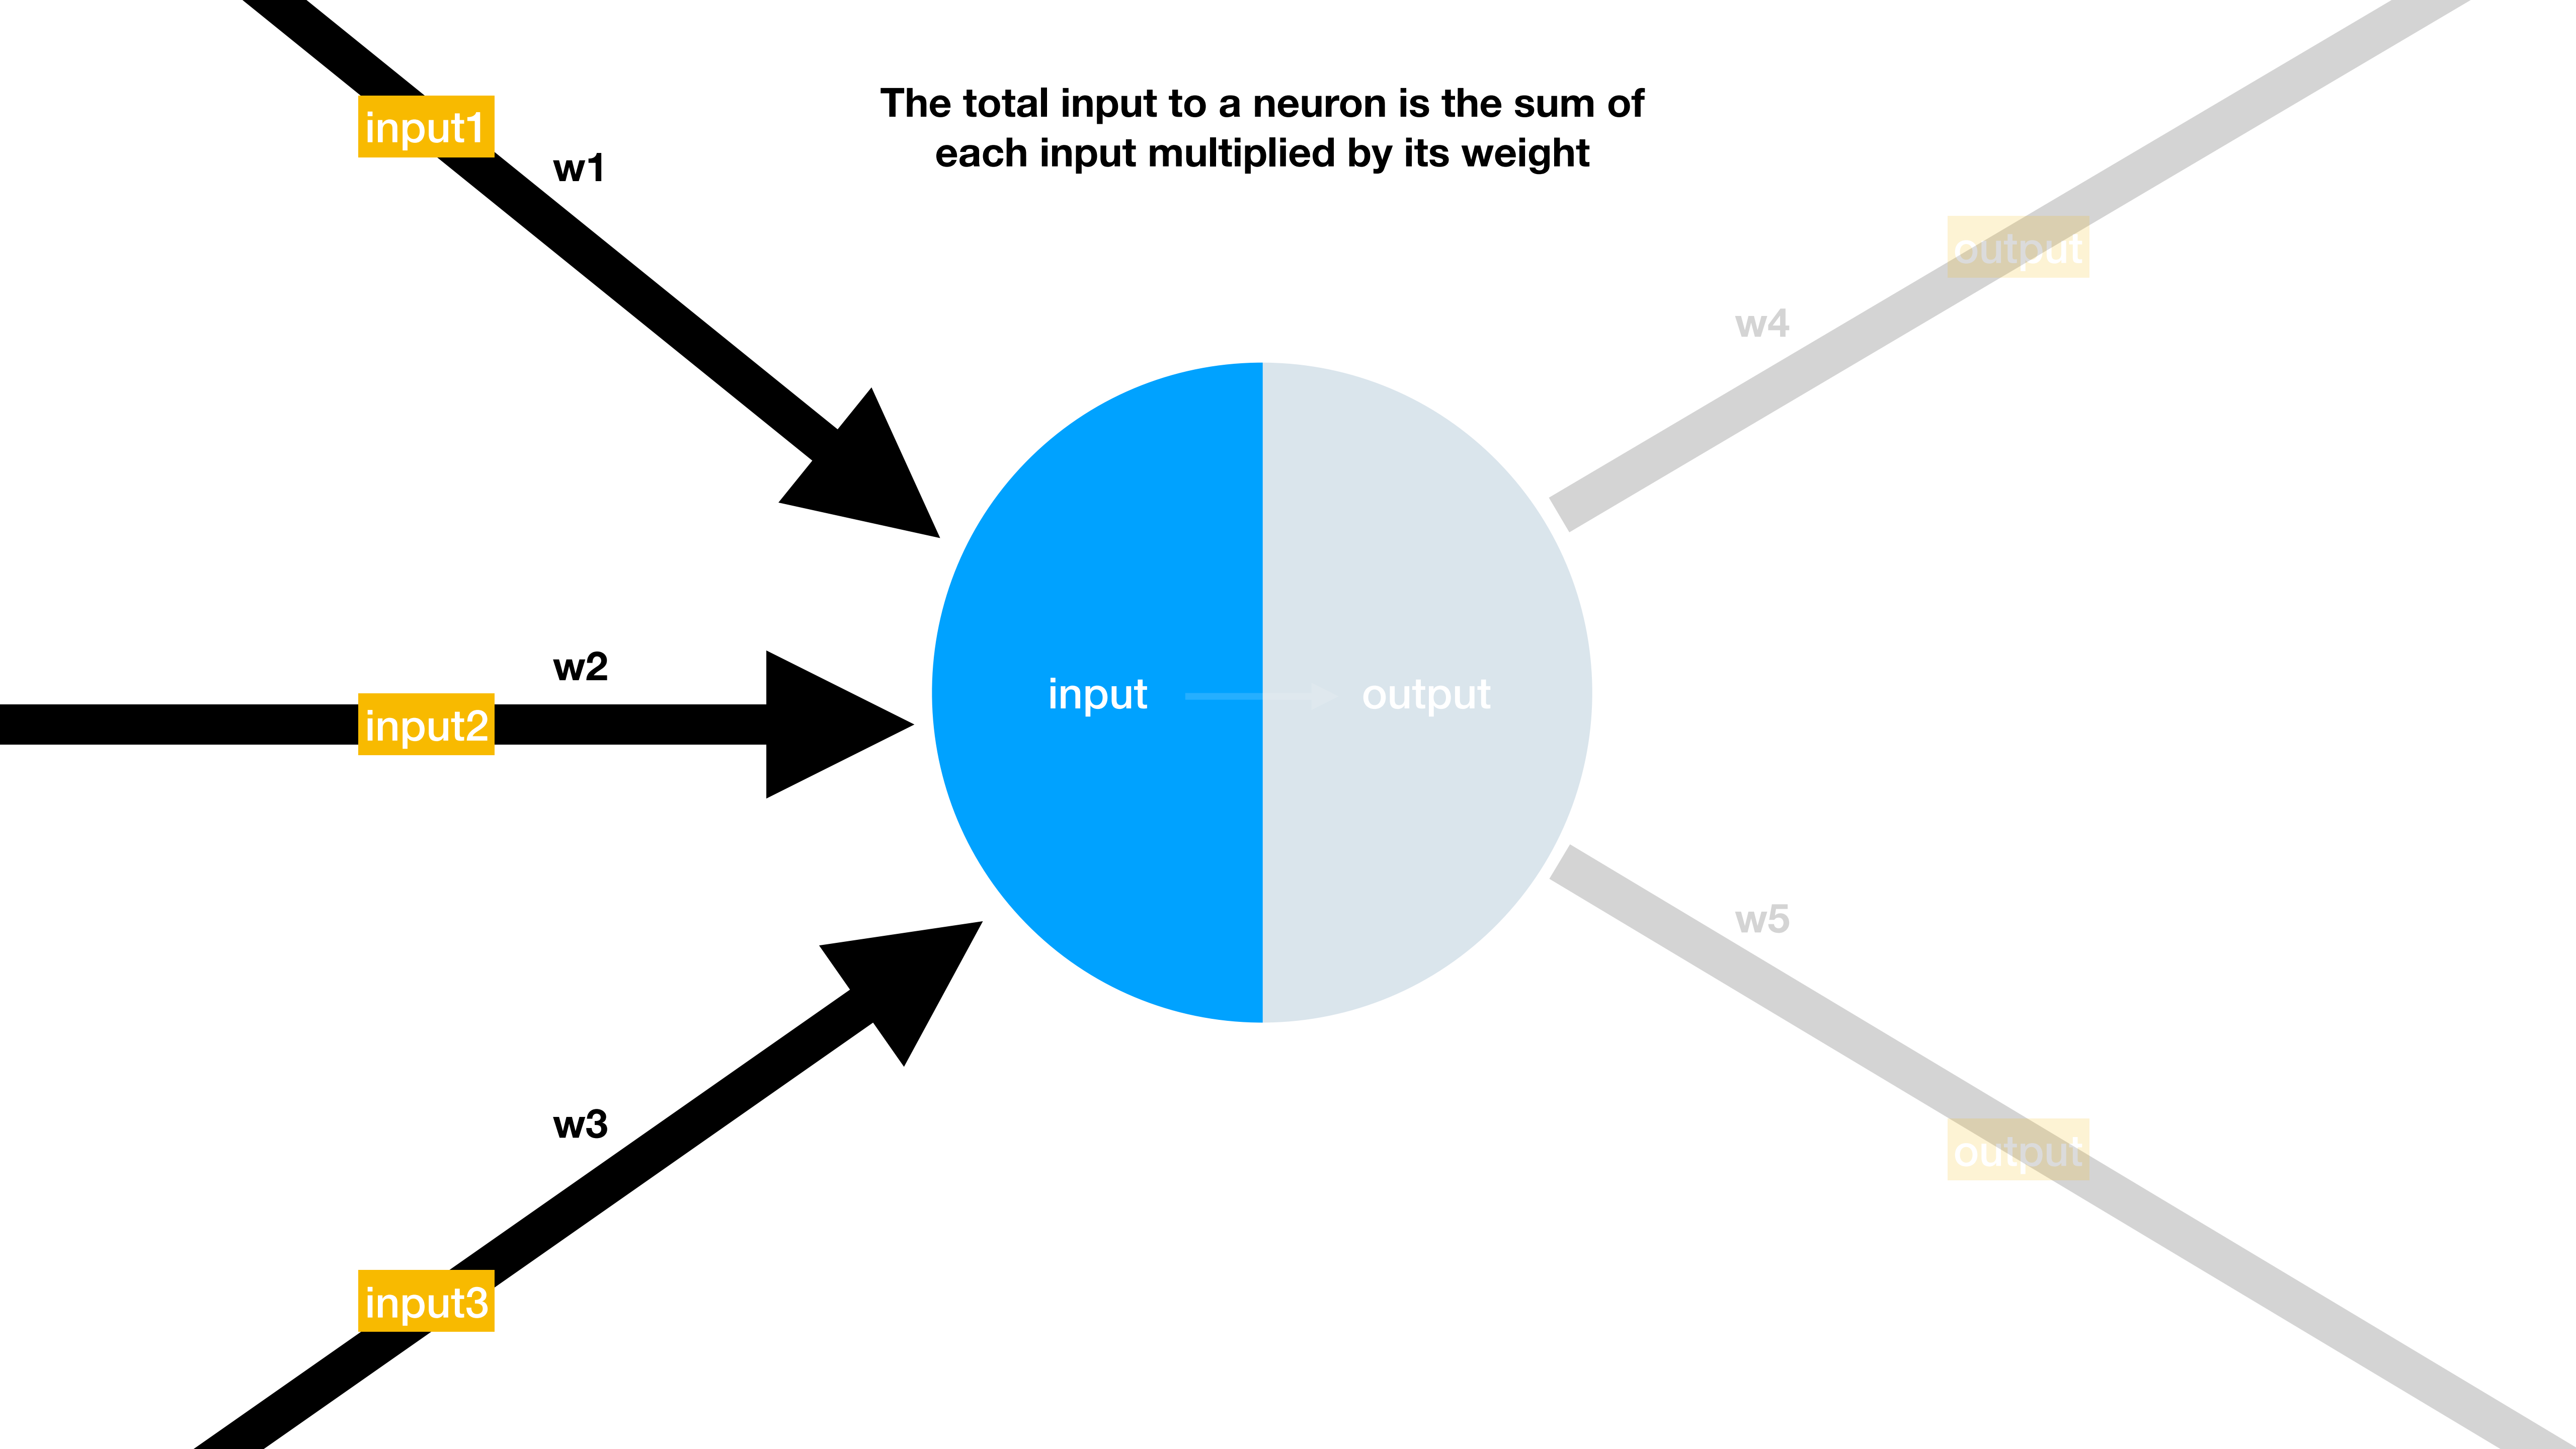
To see how this works, let's take a closer look at a neuron

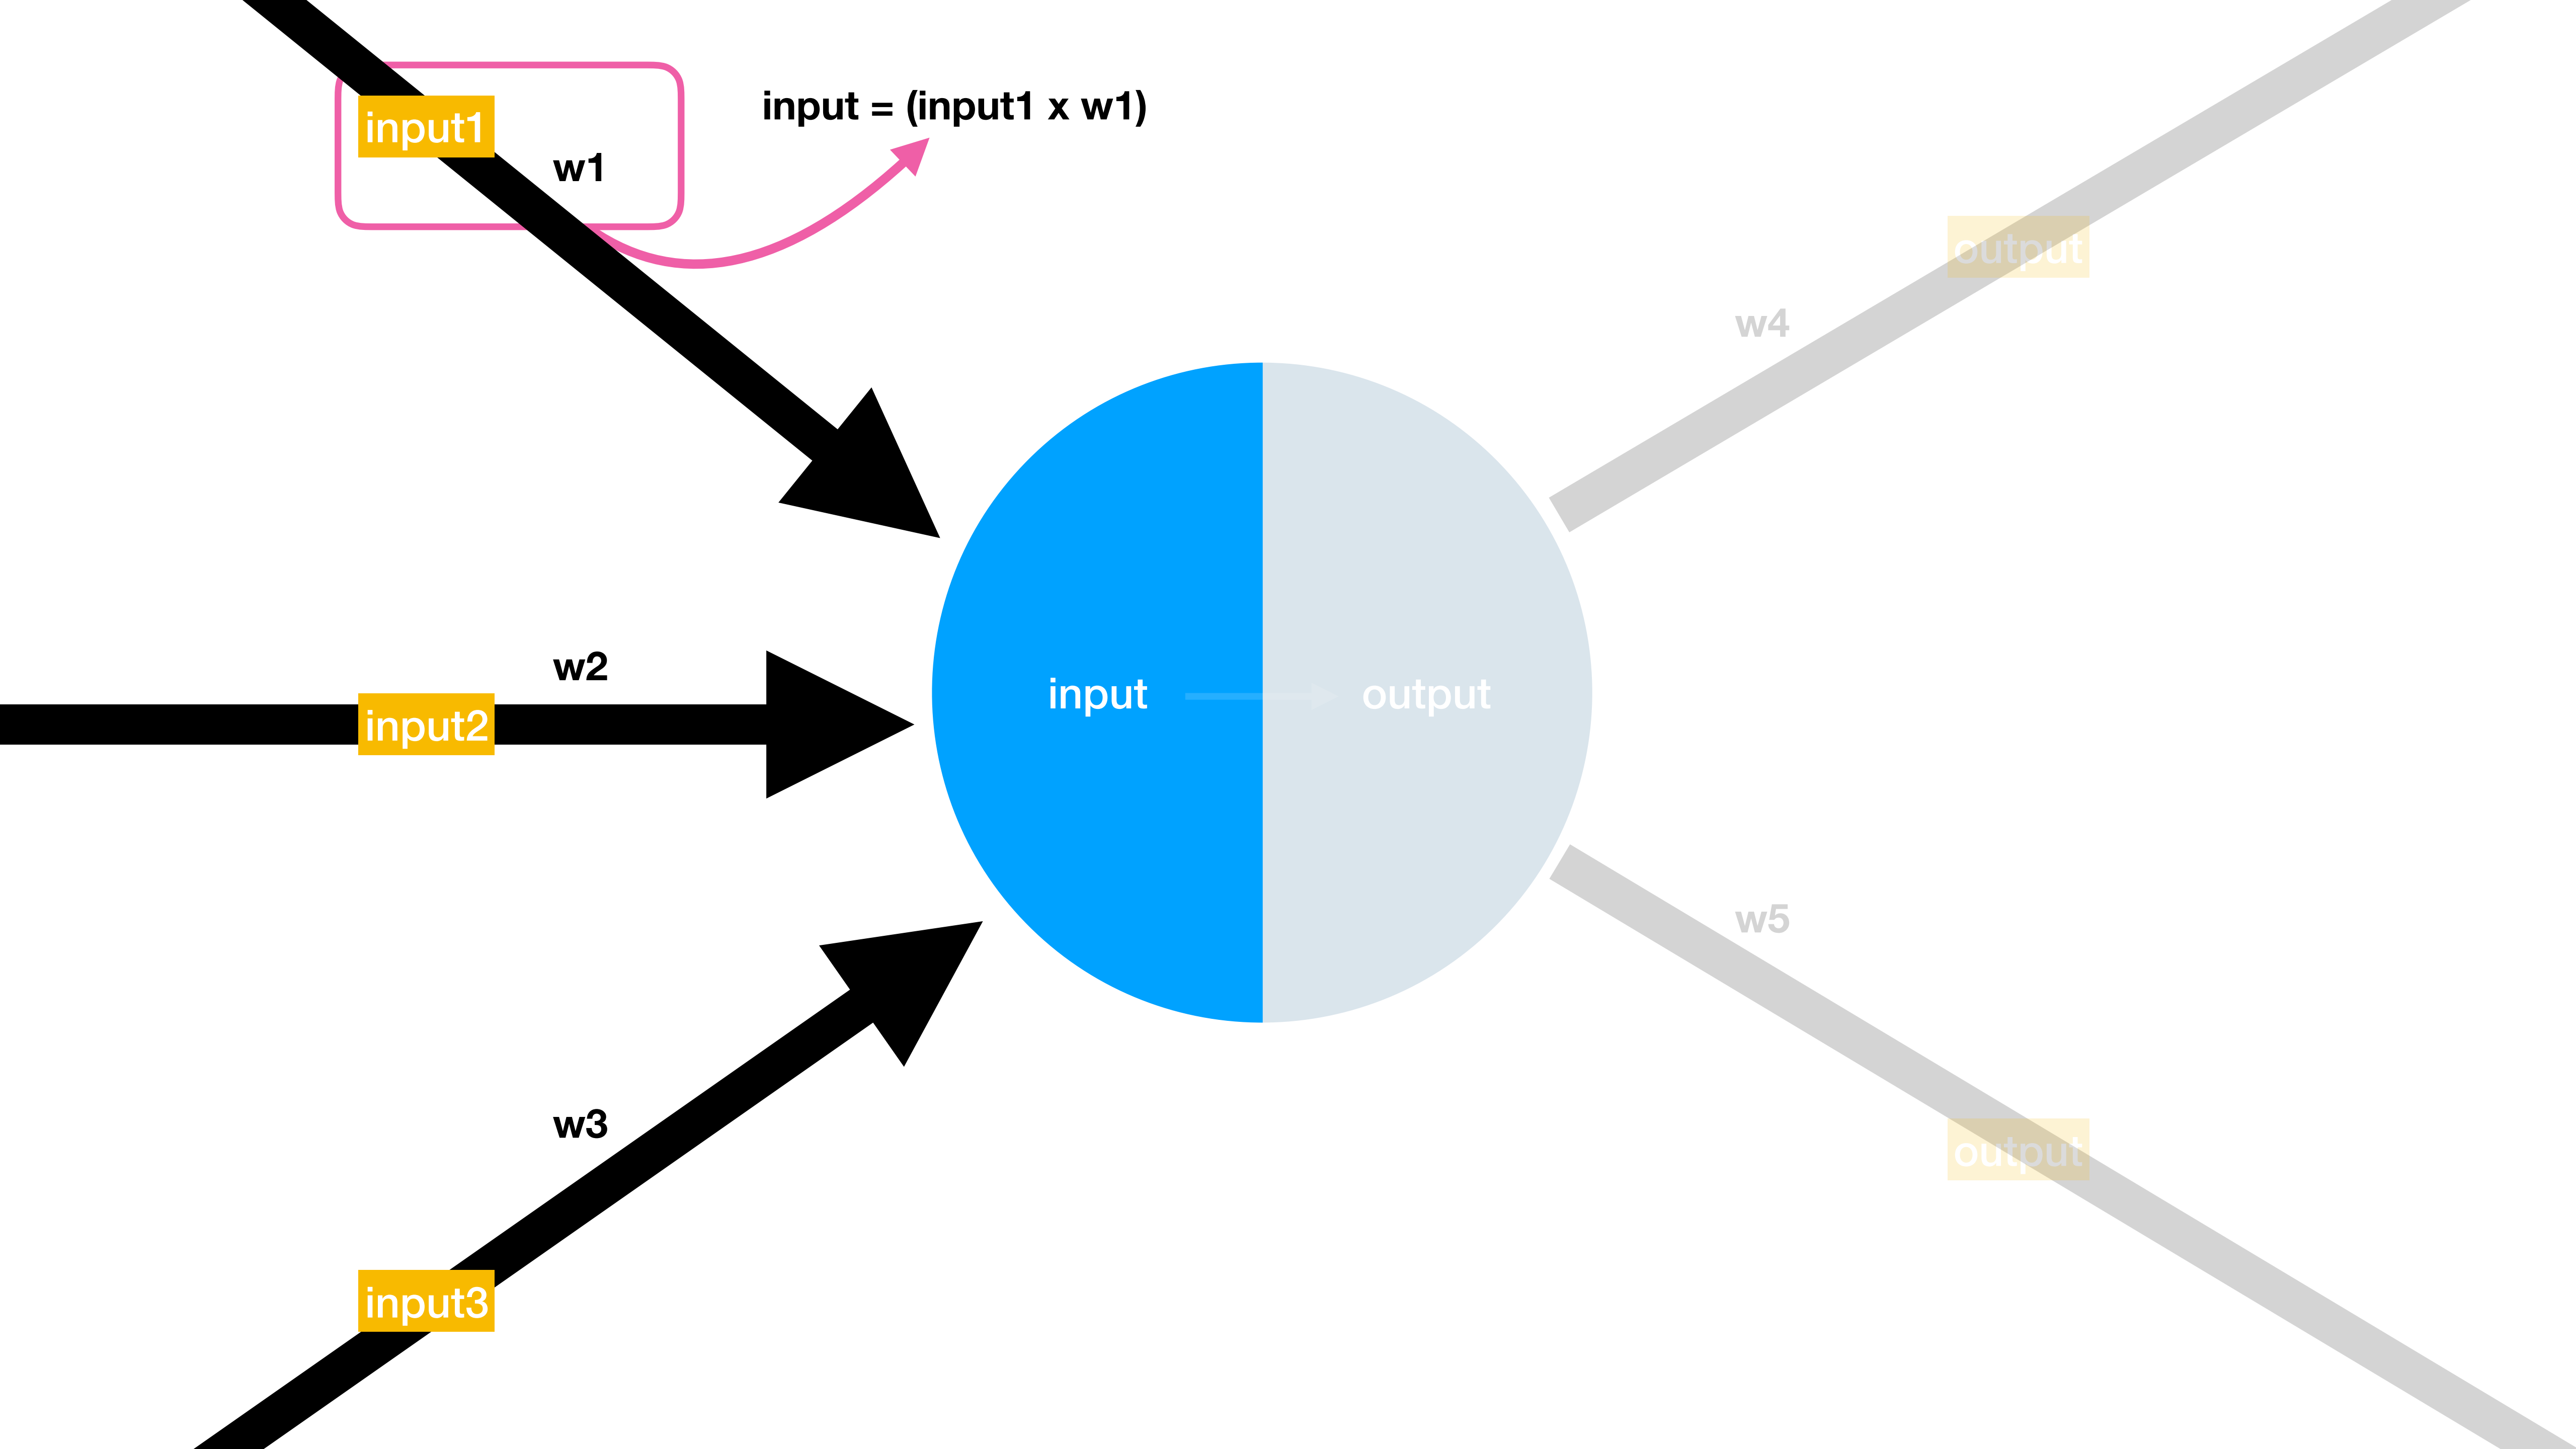
A single neuron




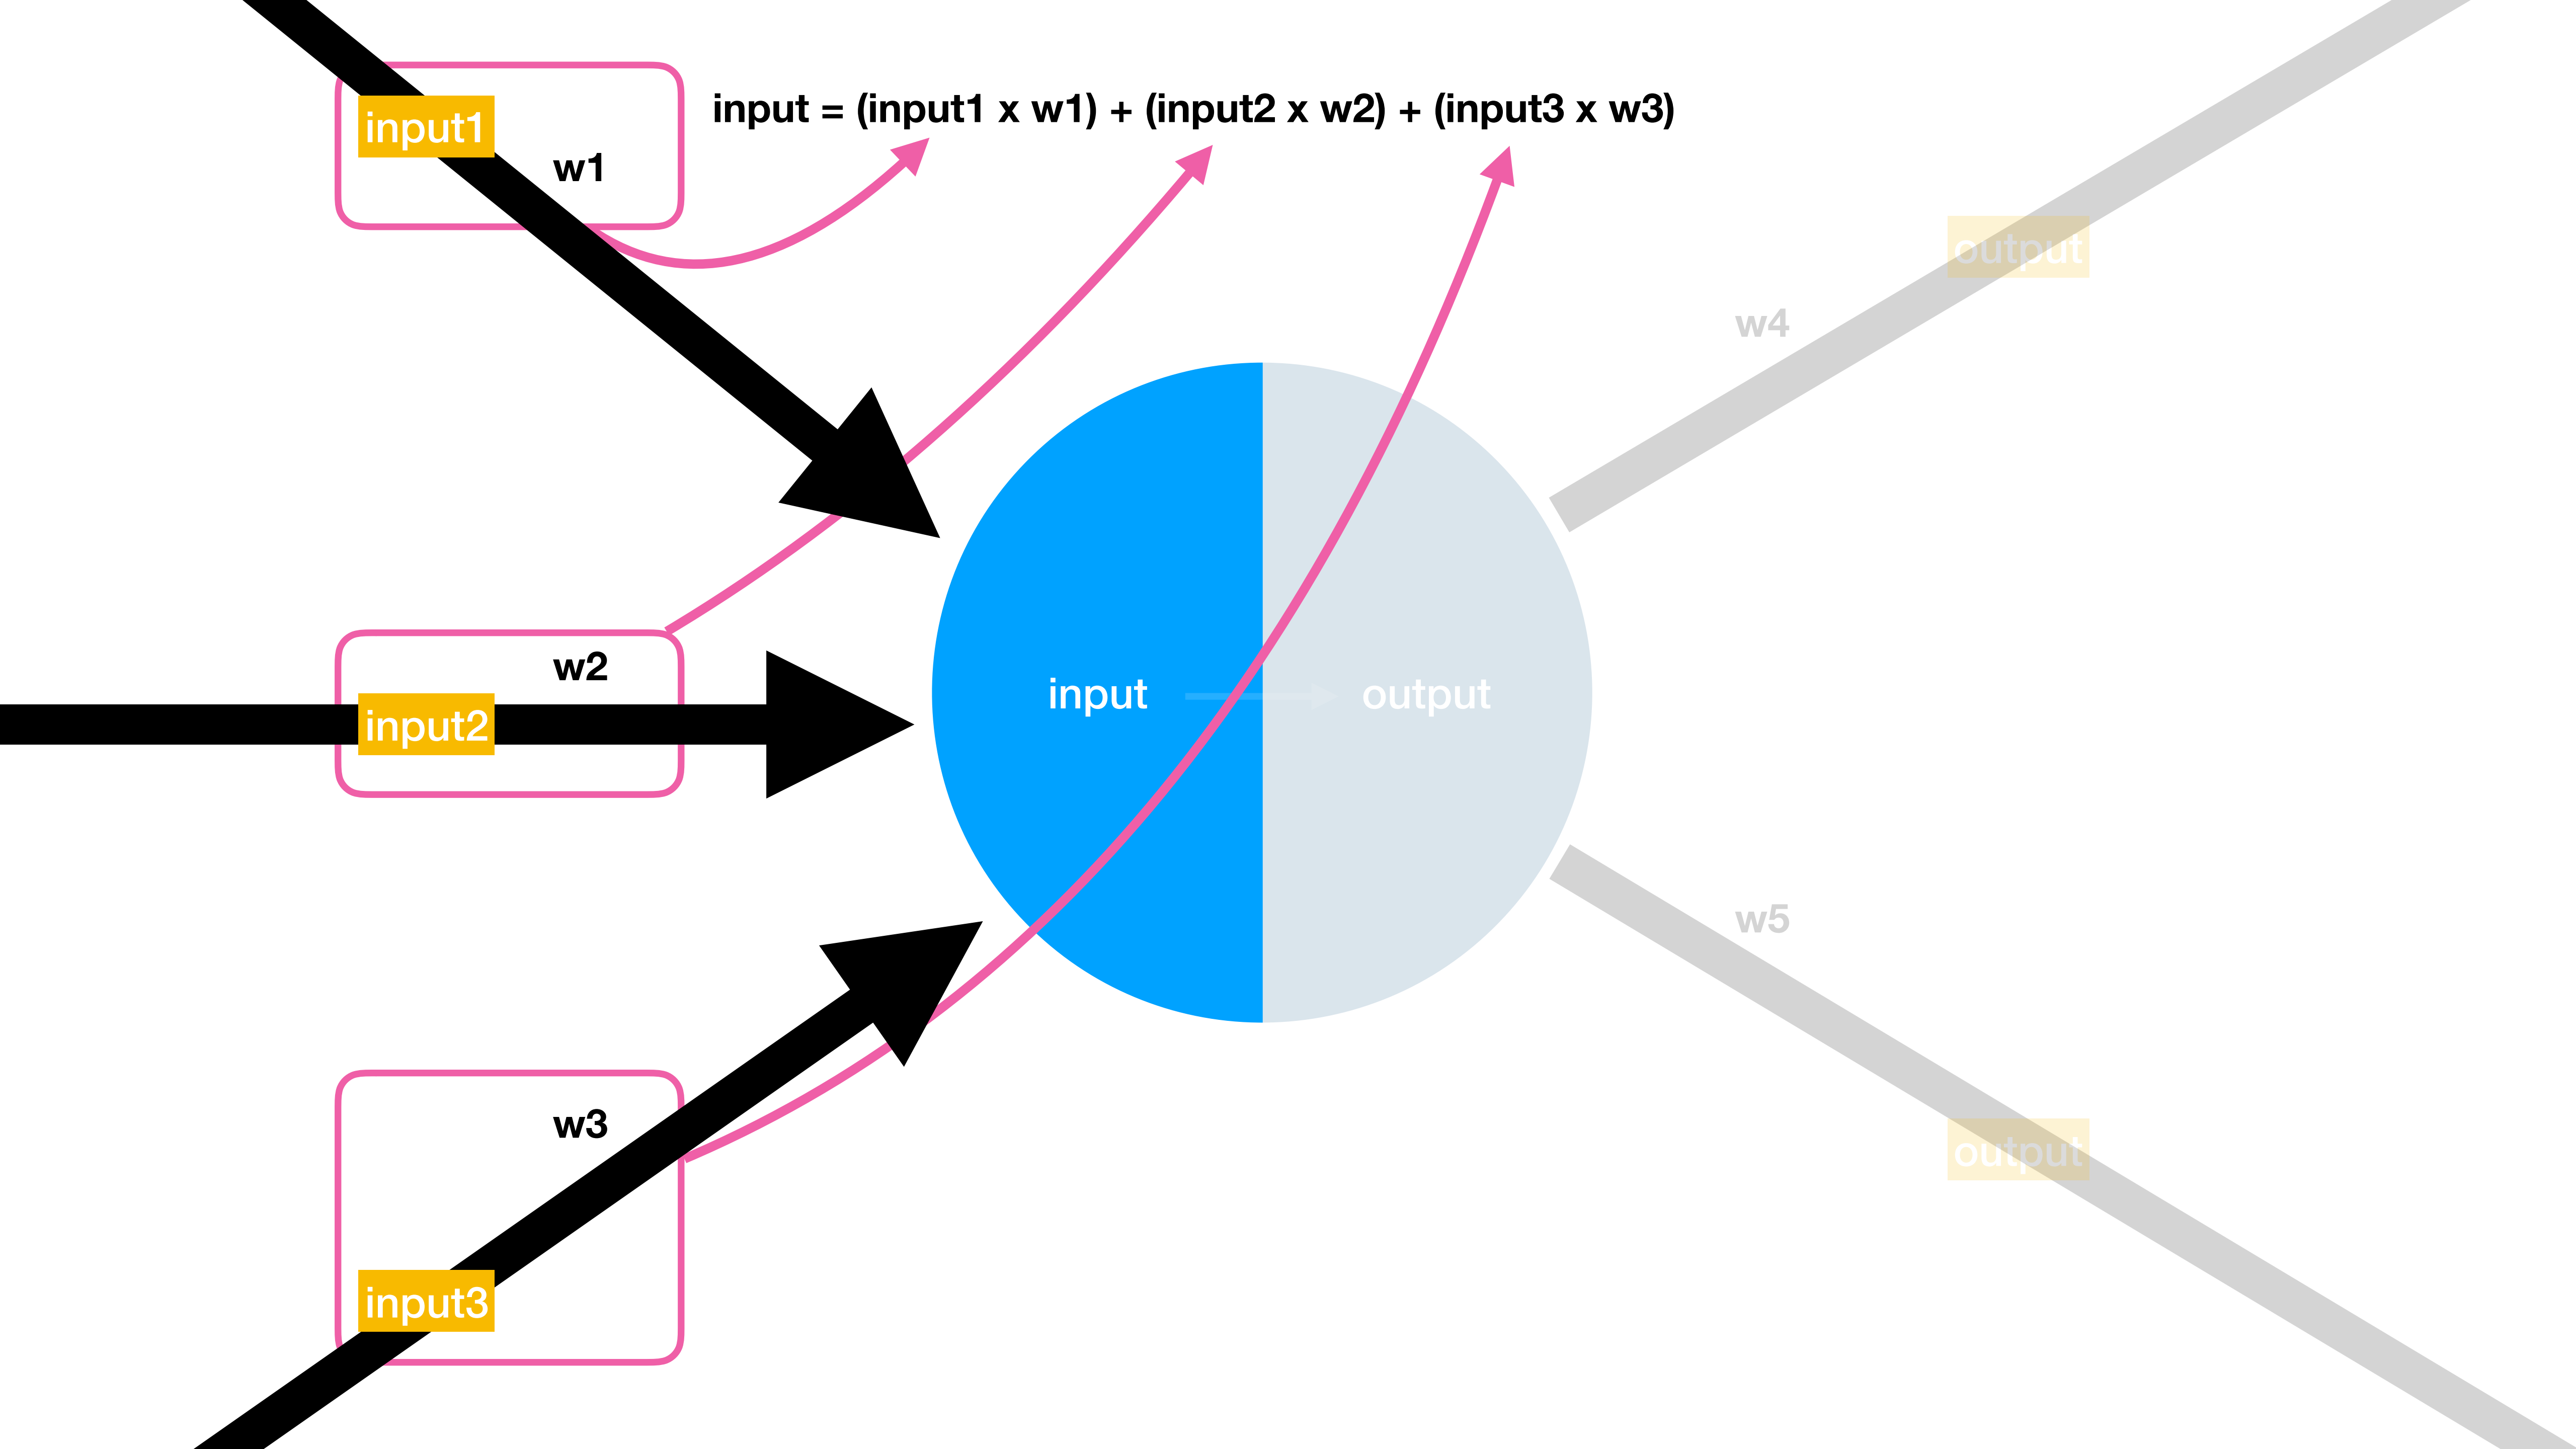


The total input to a neuron is the sum of each input multiplied by its weight

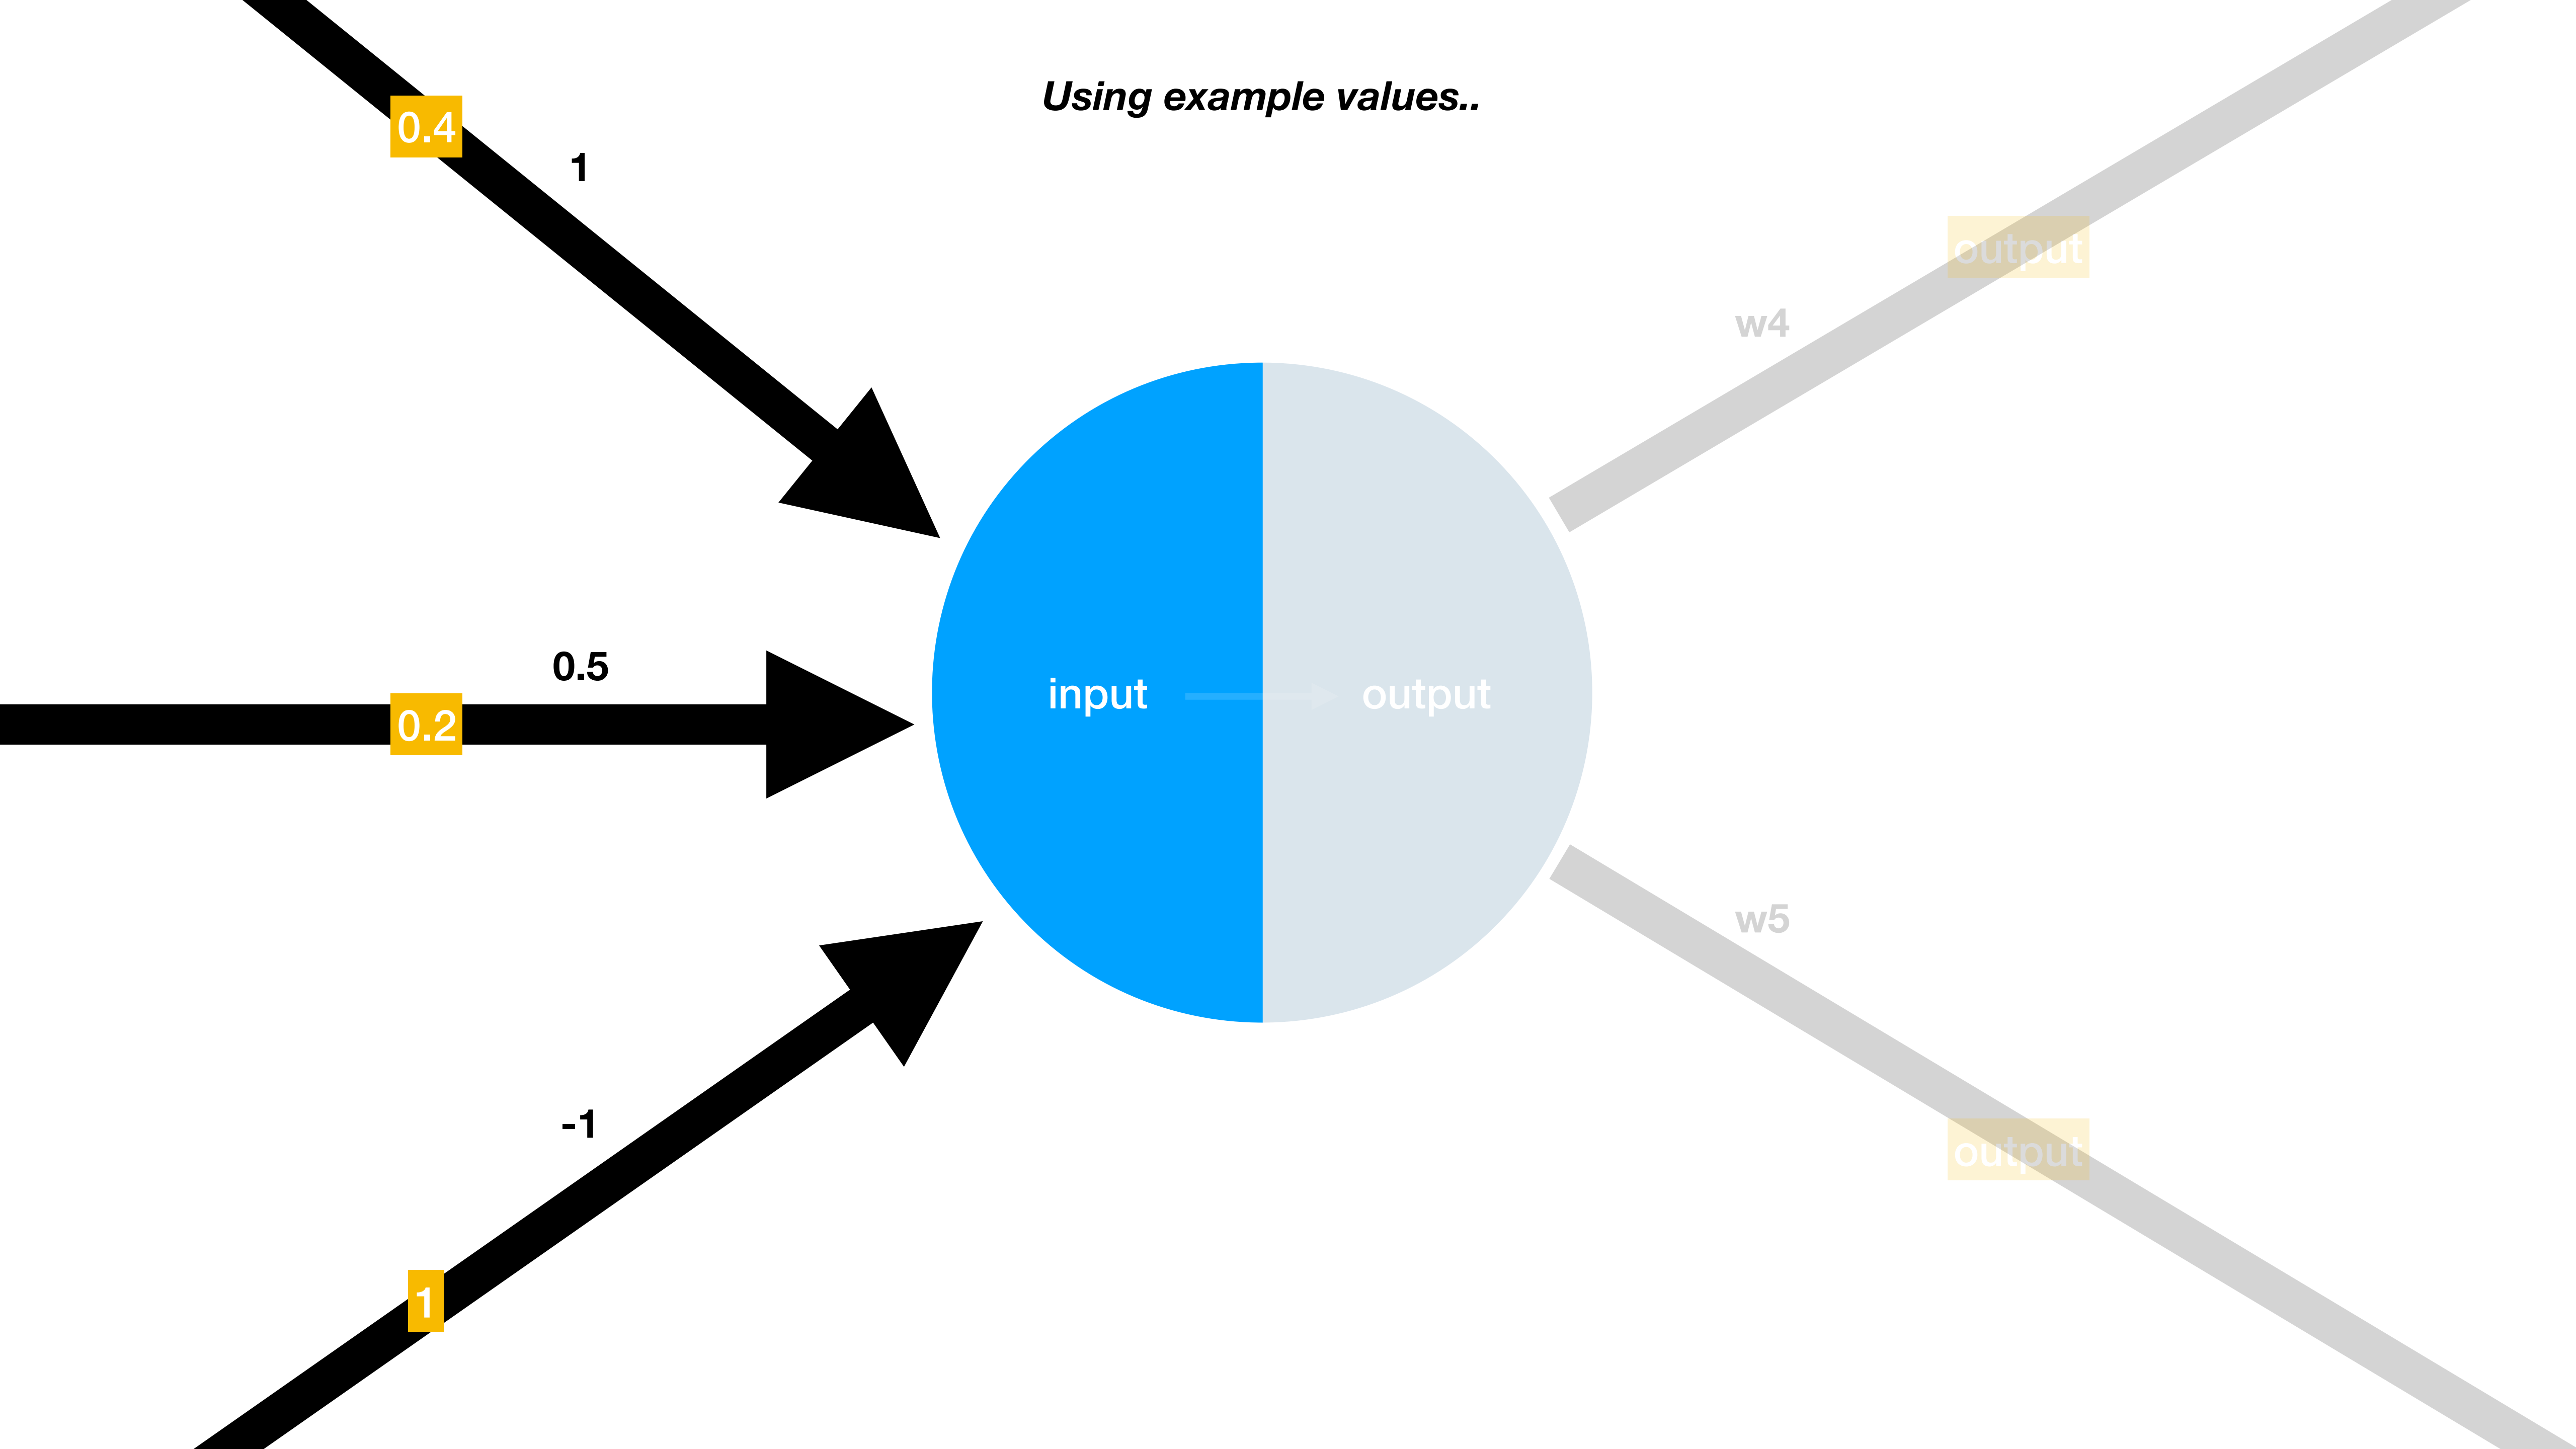


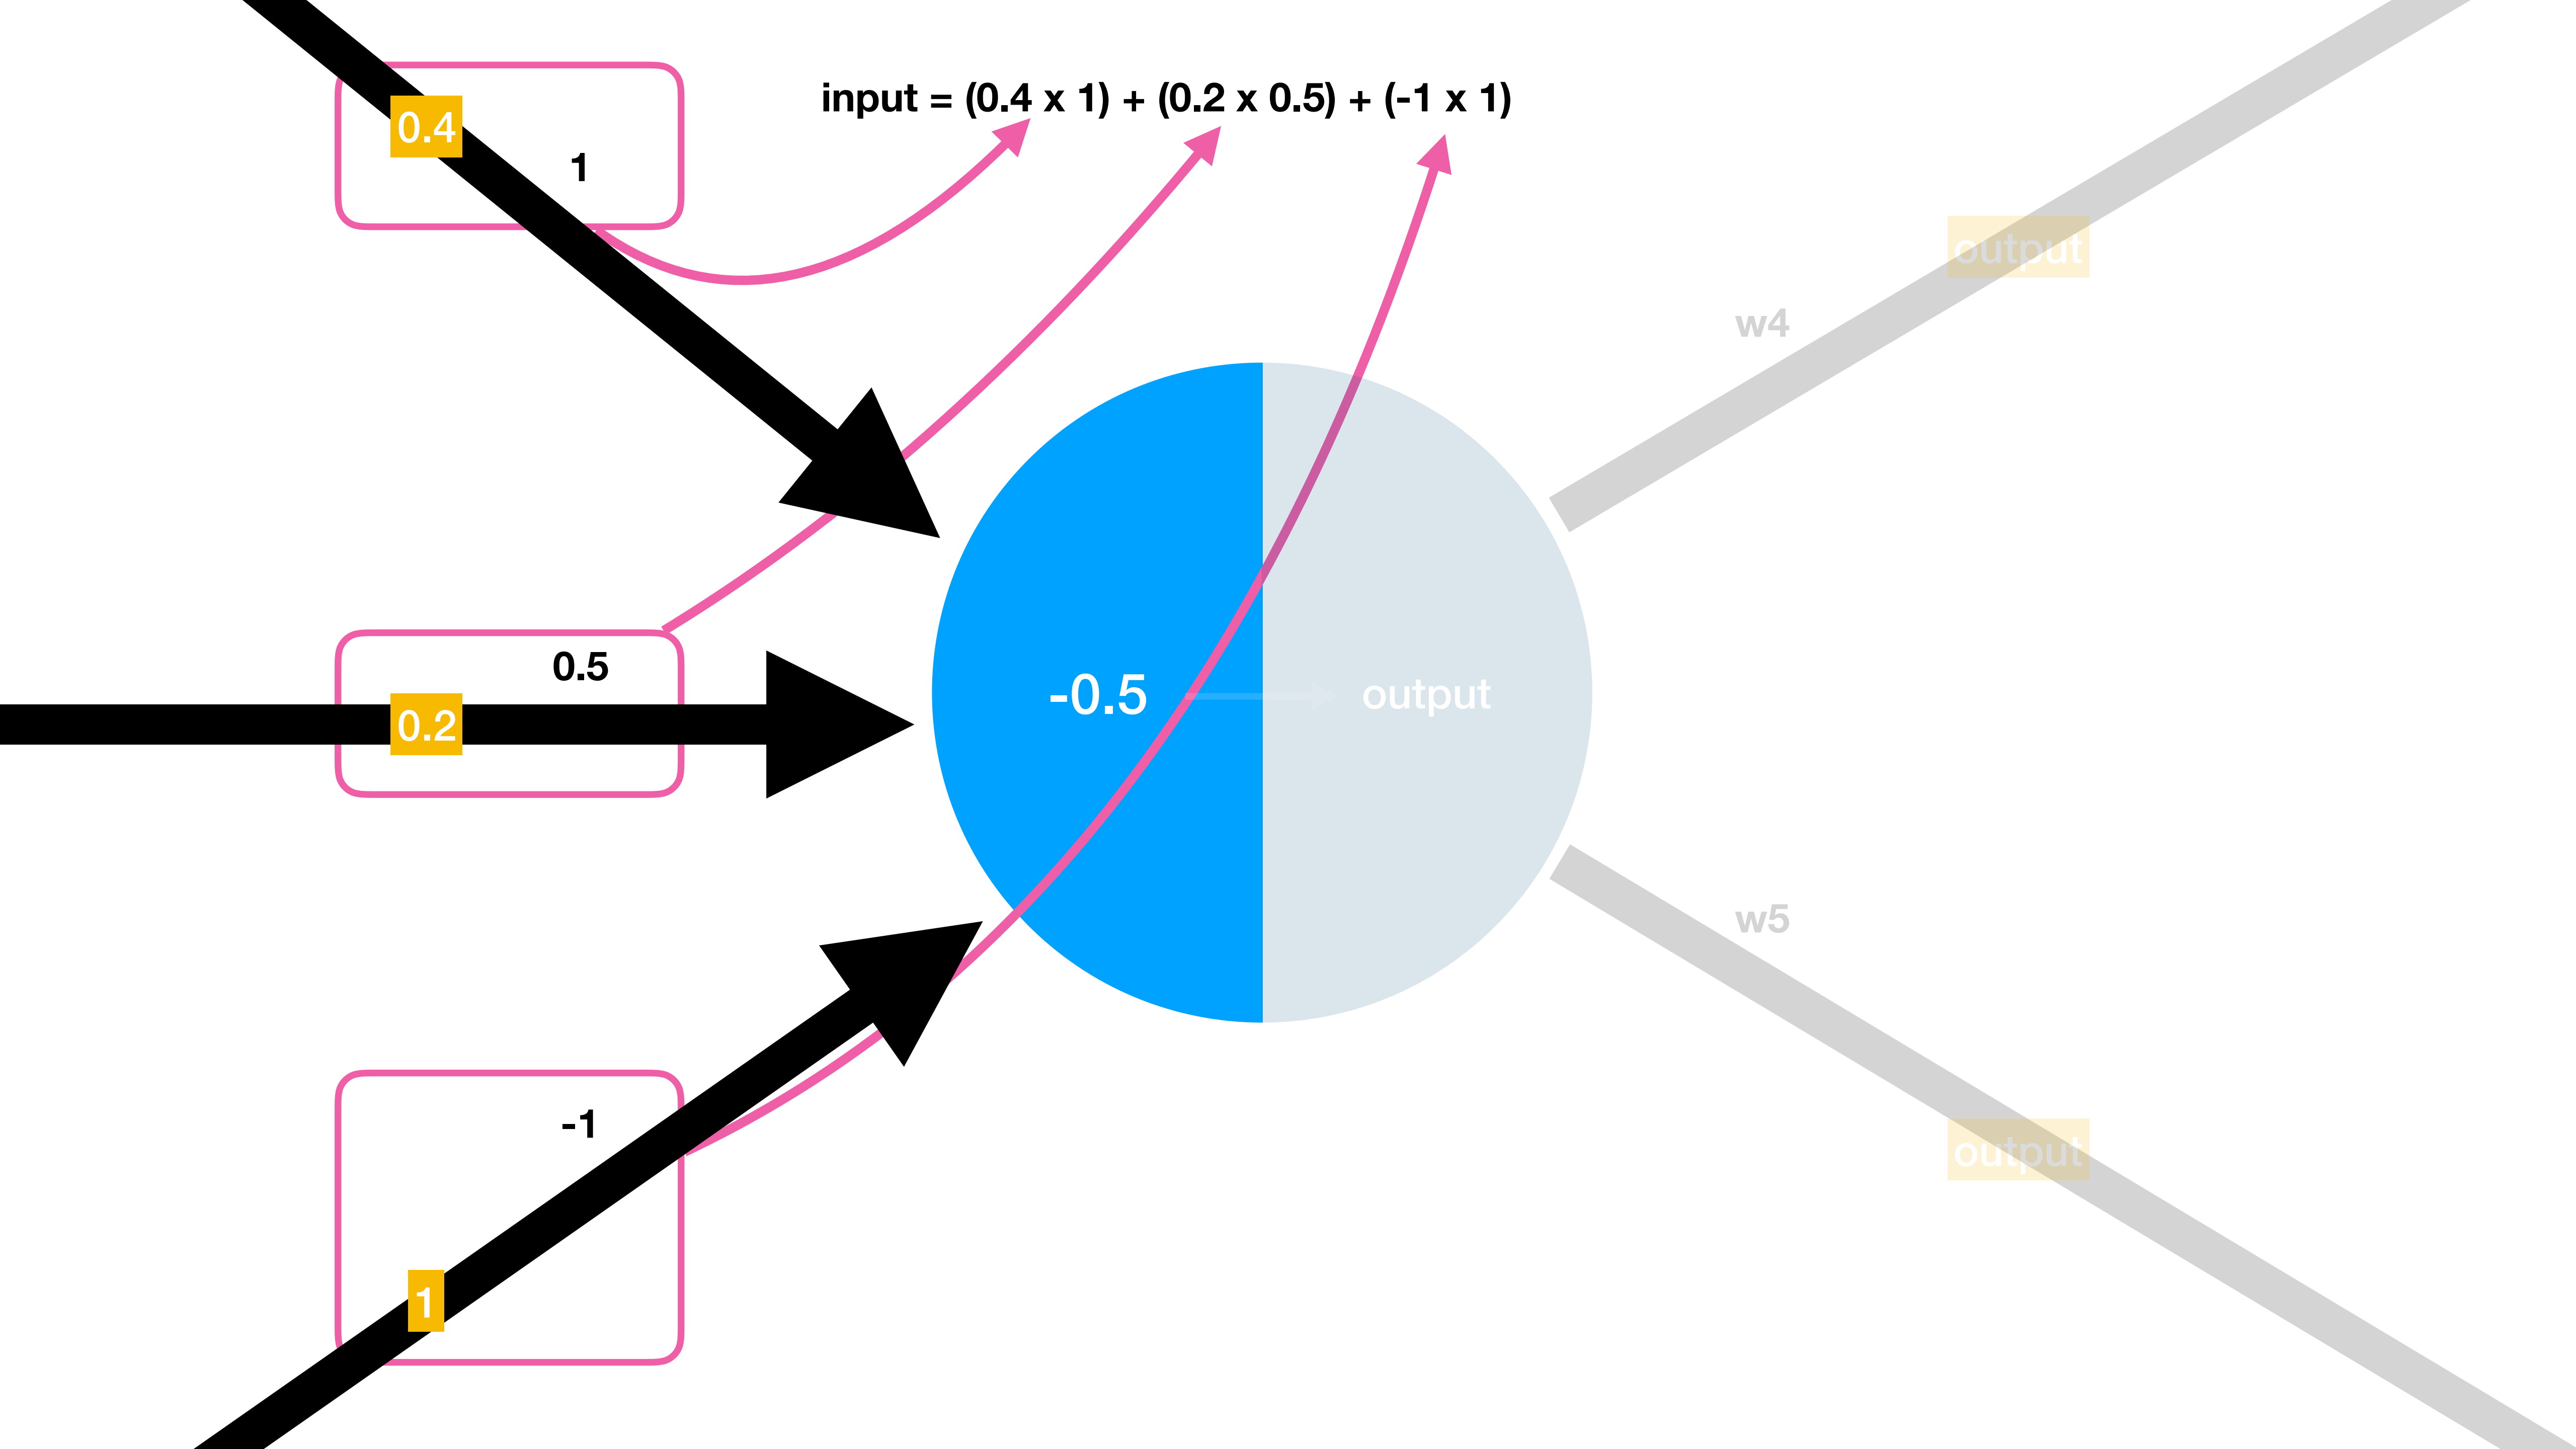


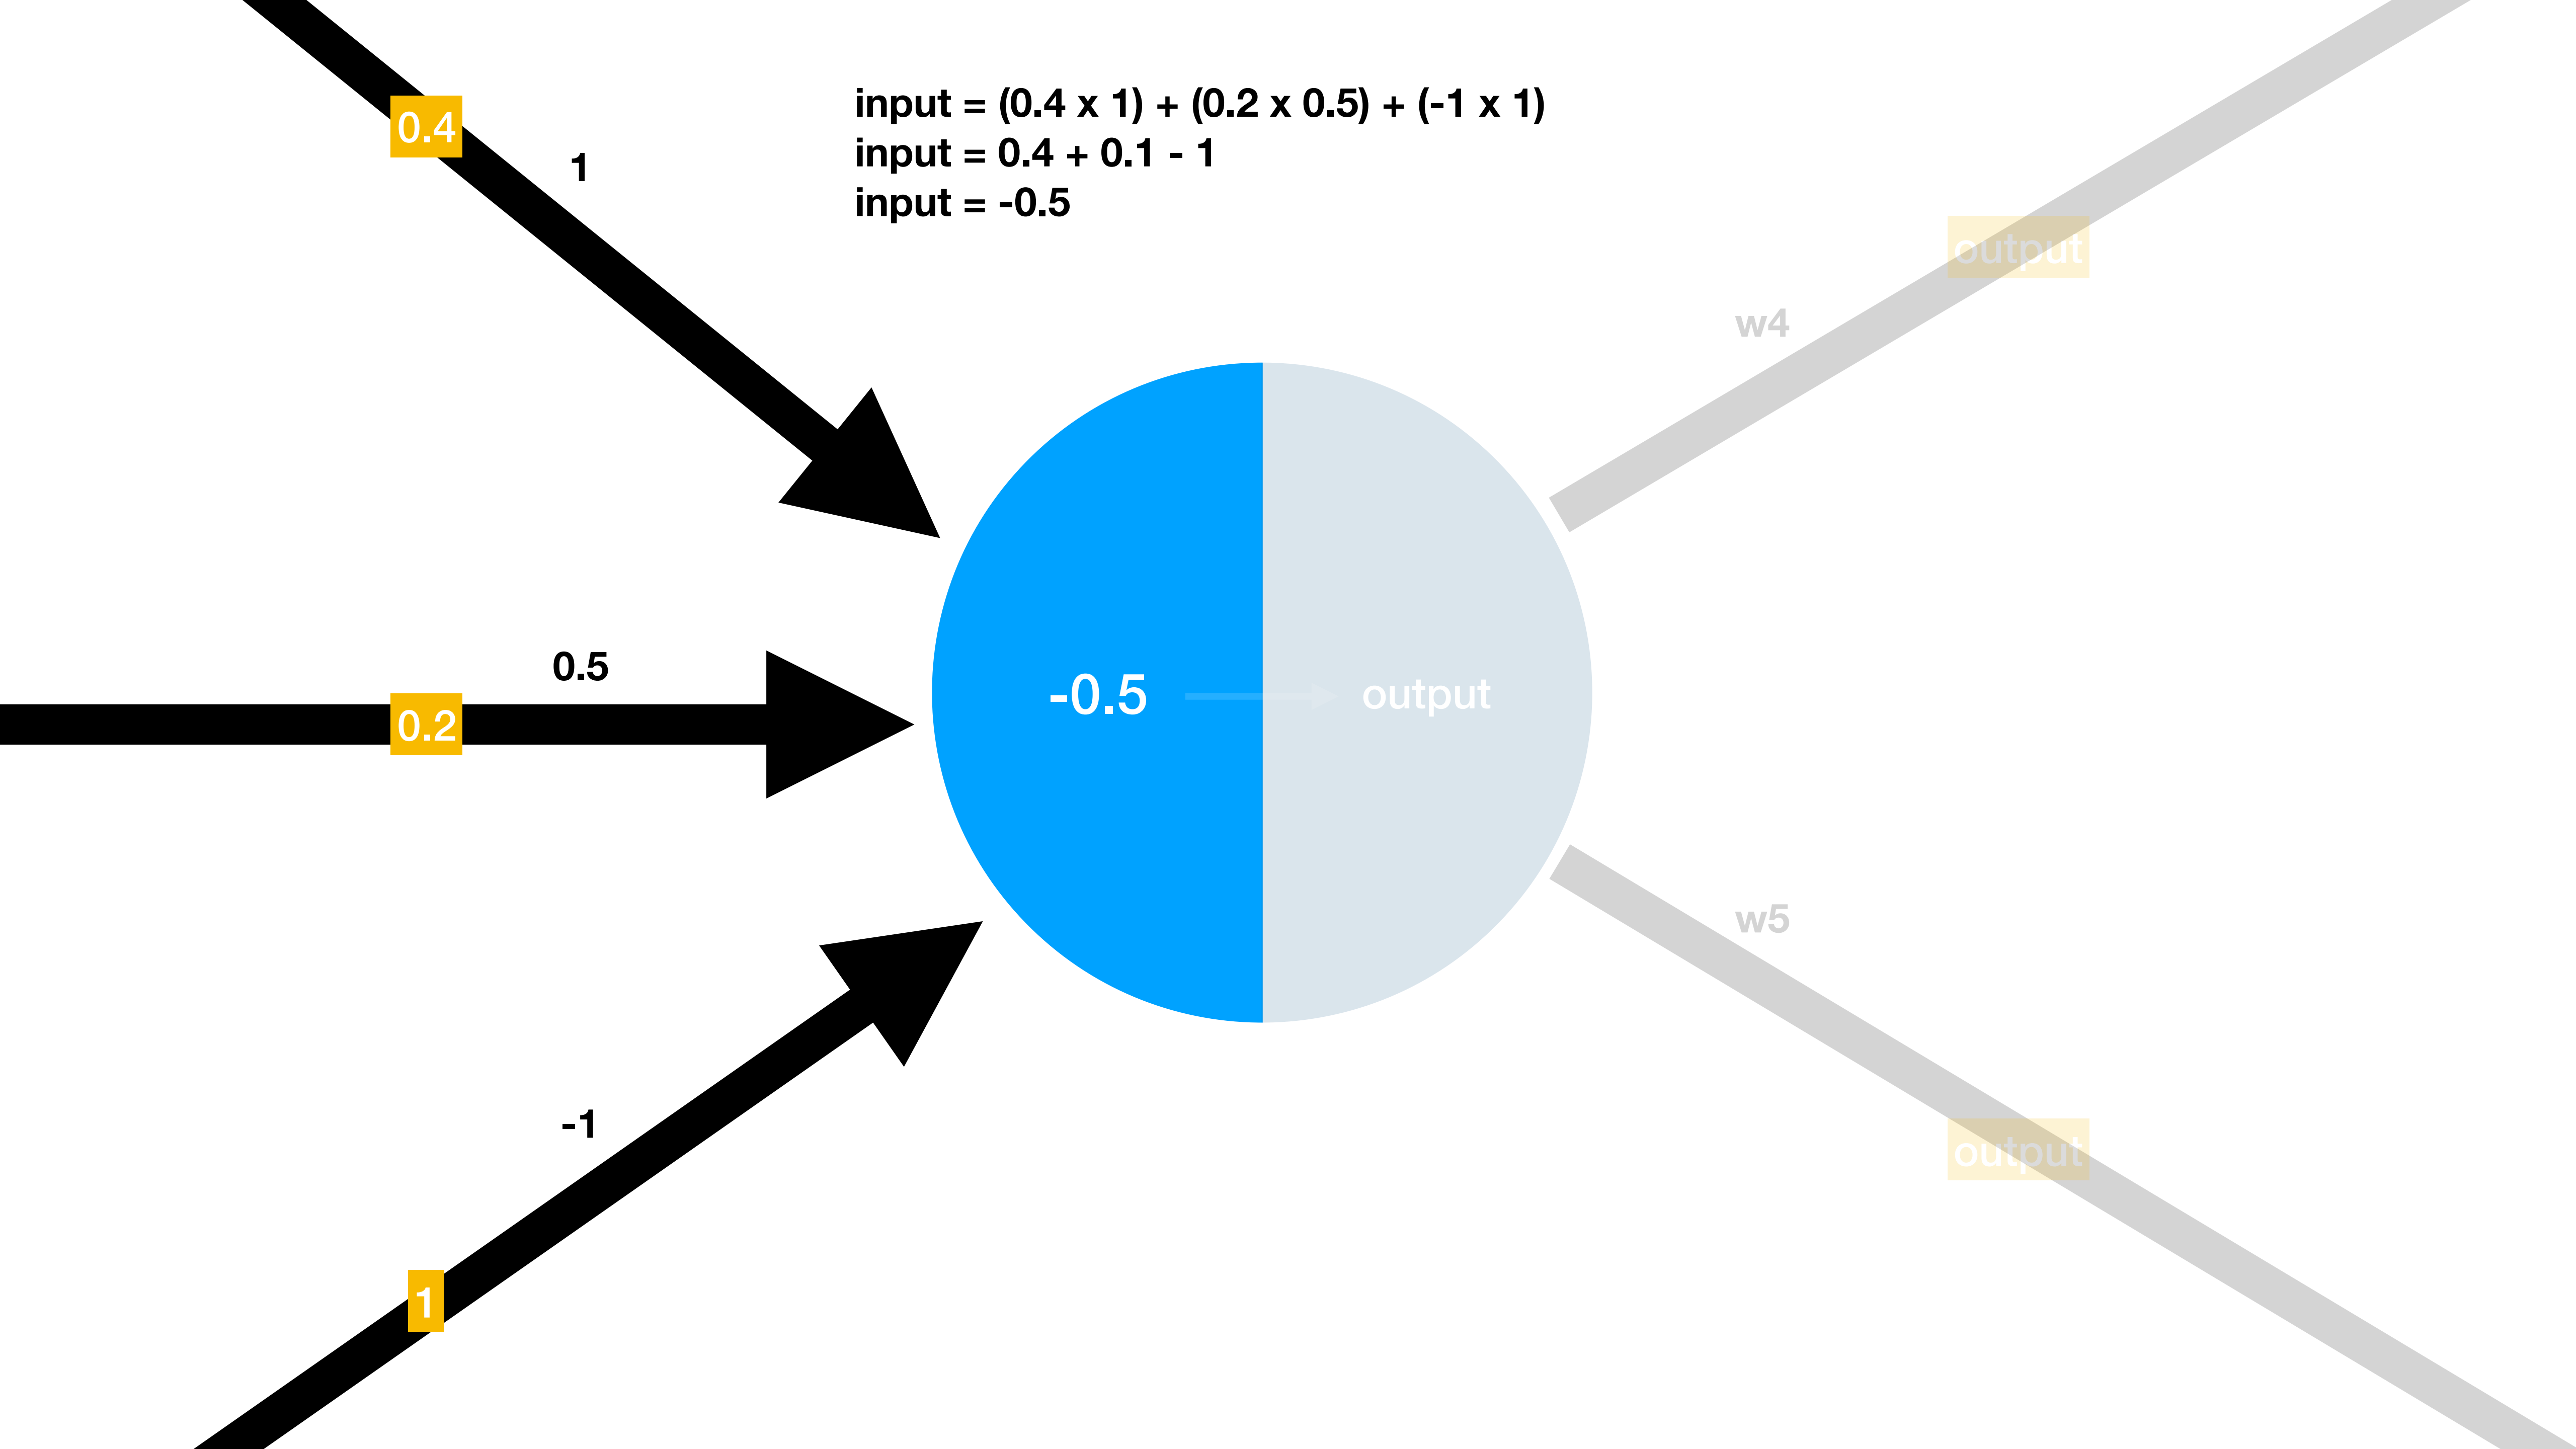


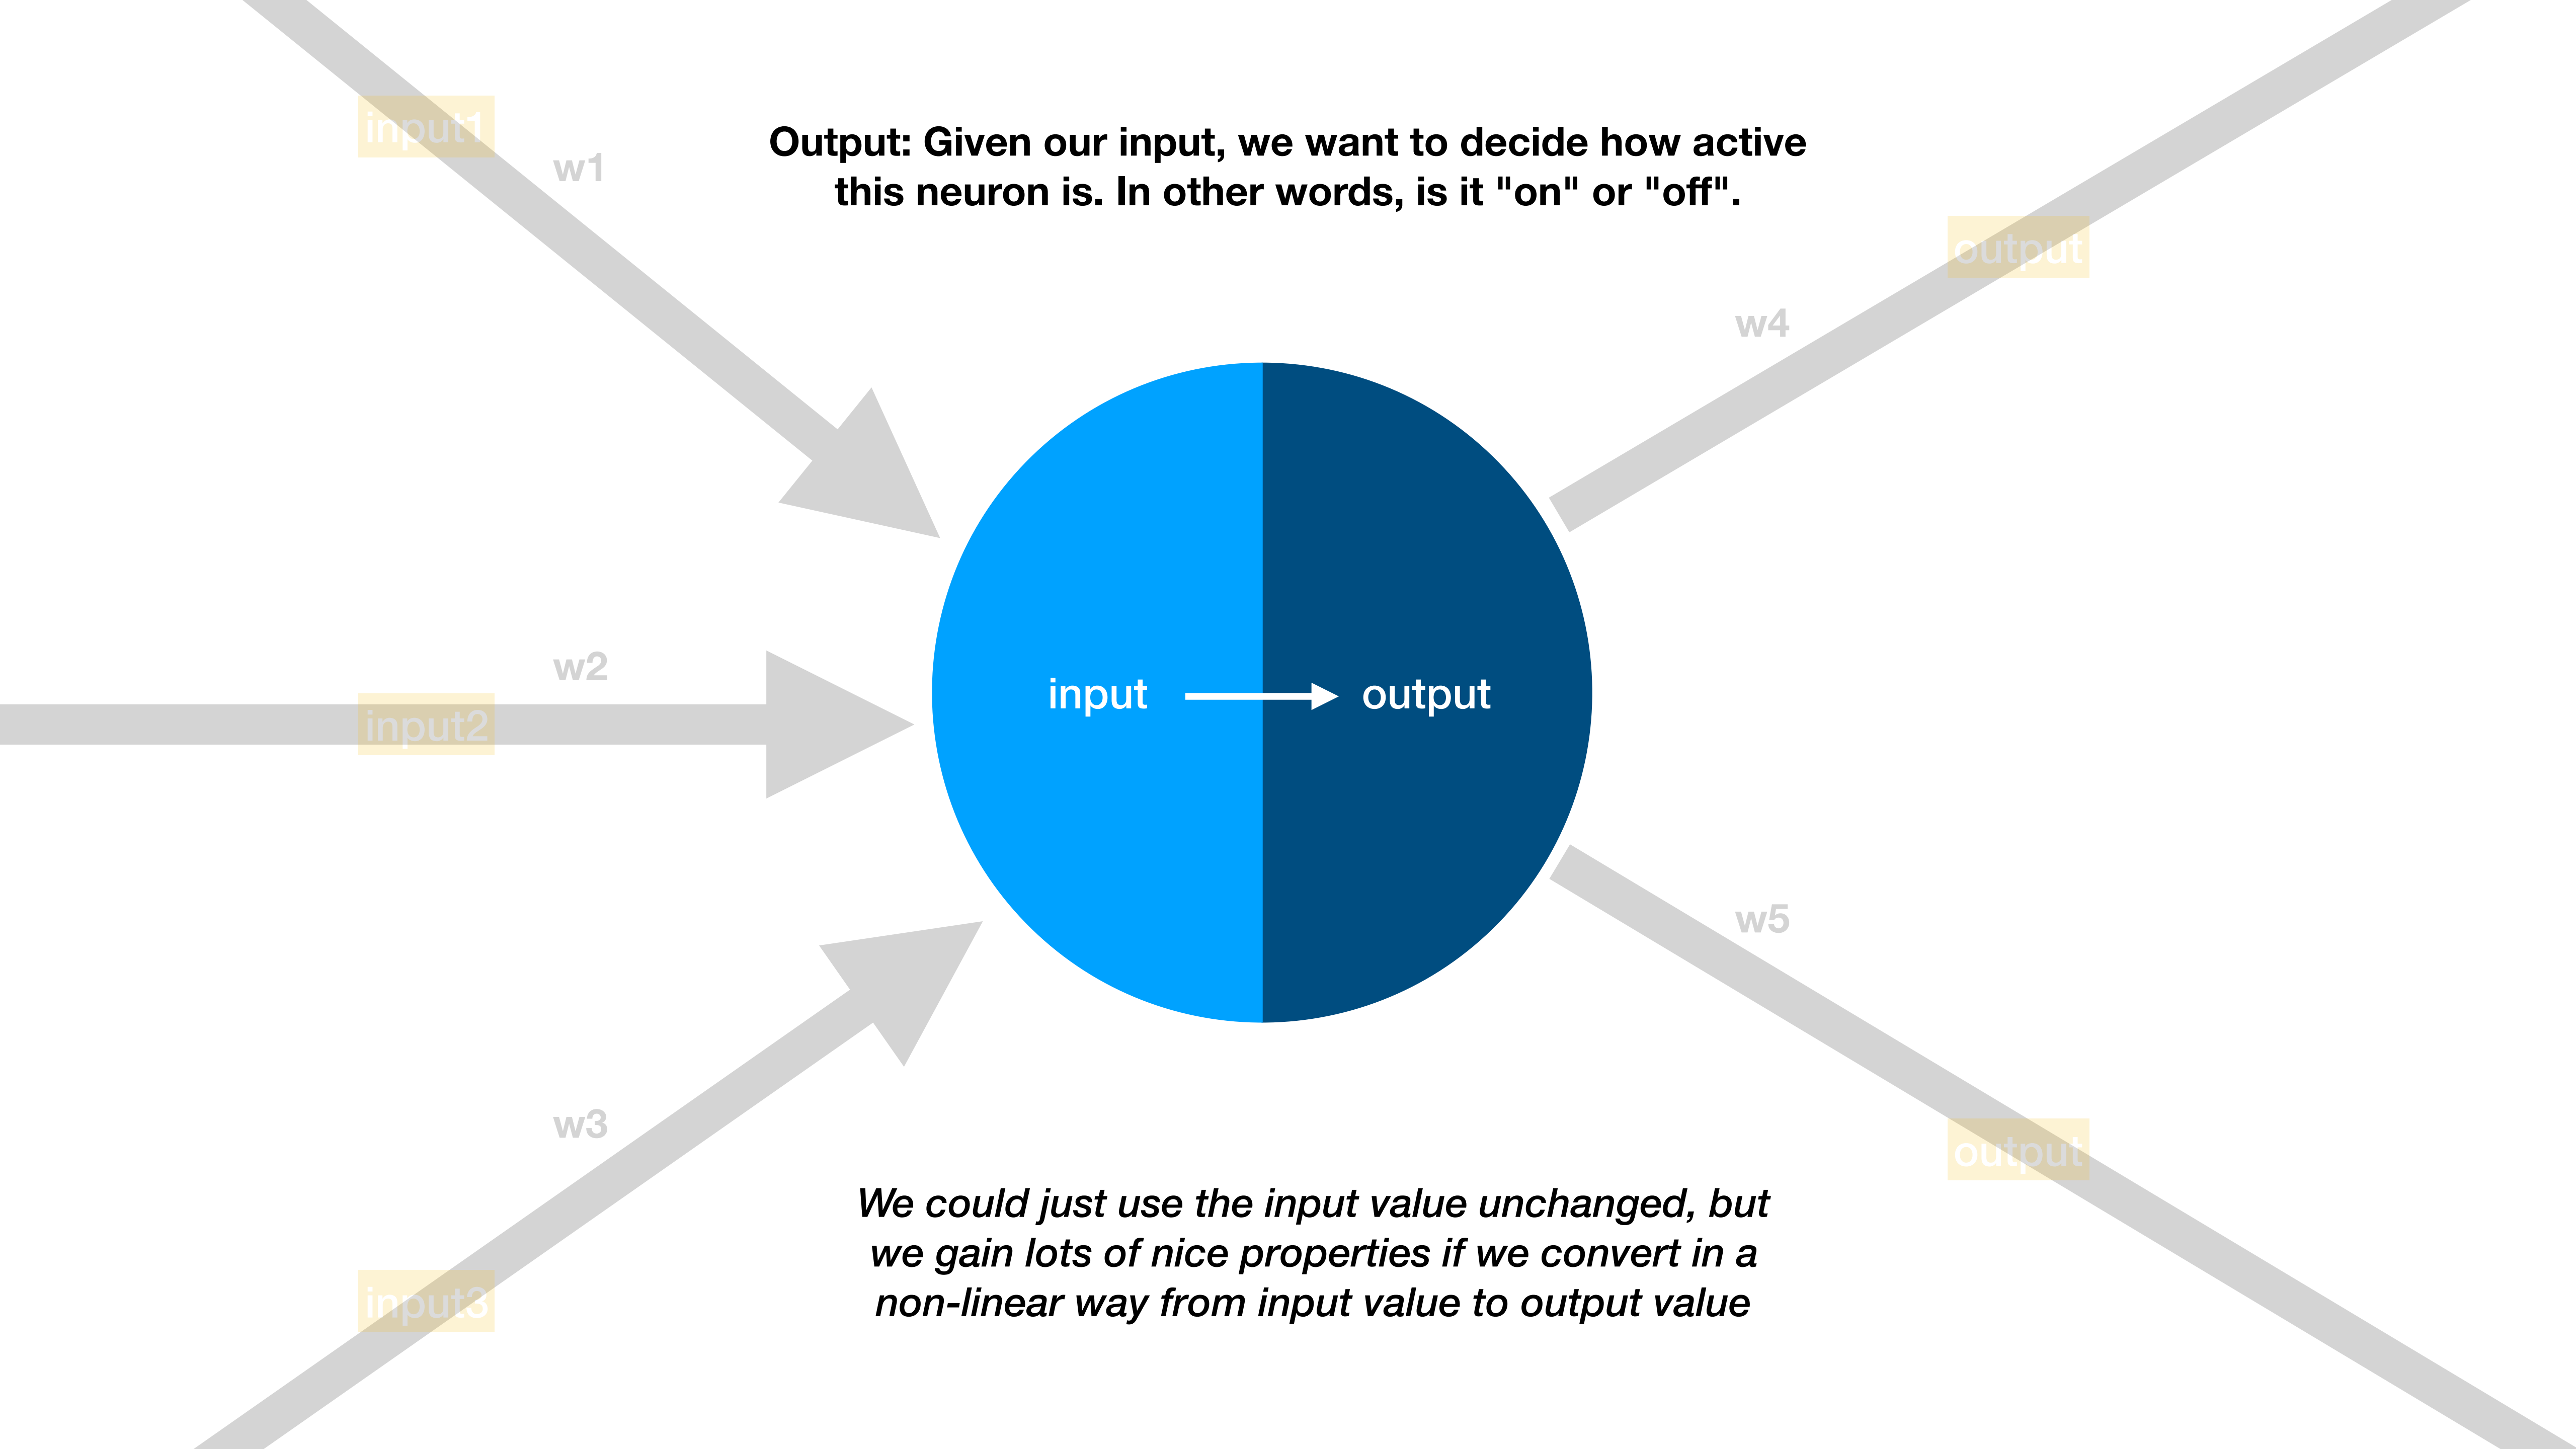


Using example values..







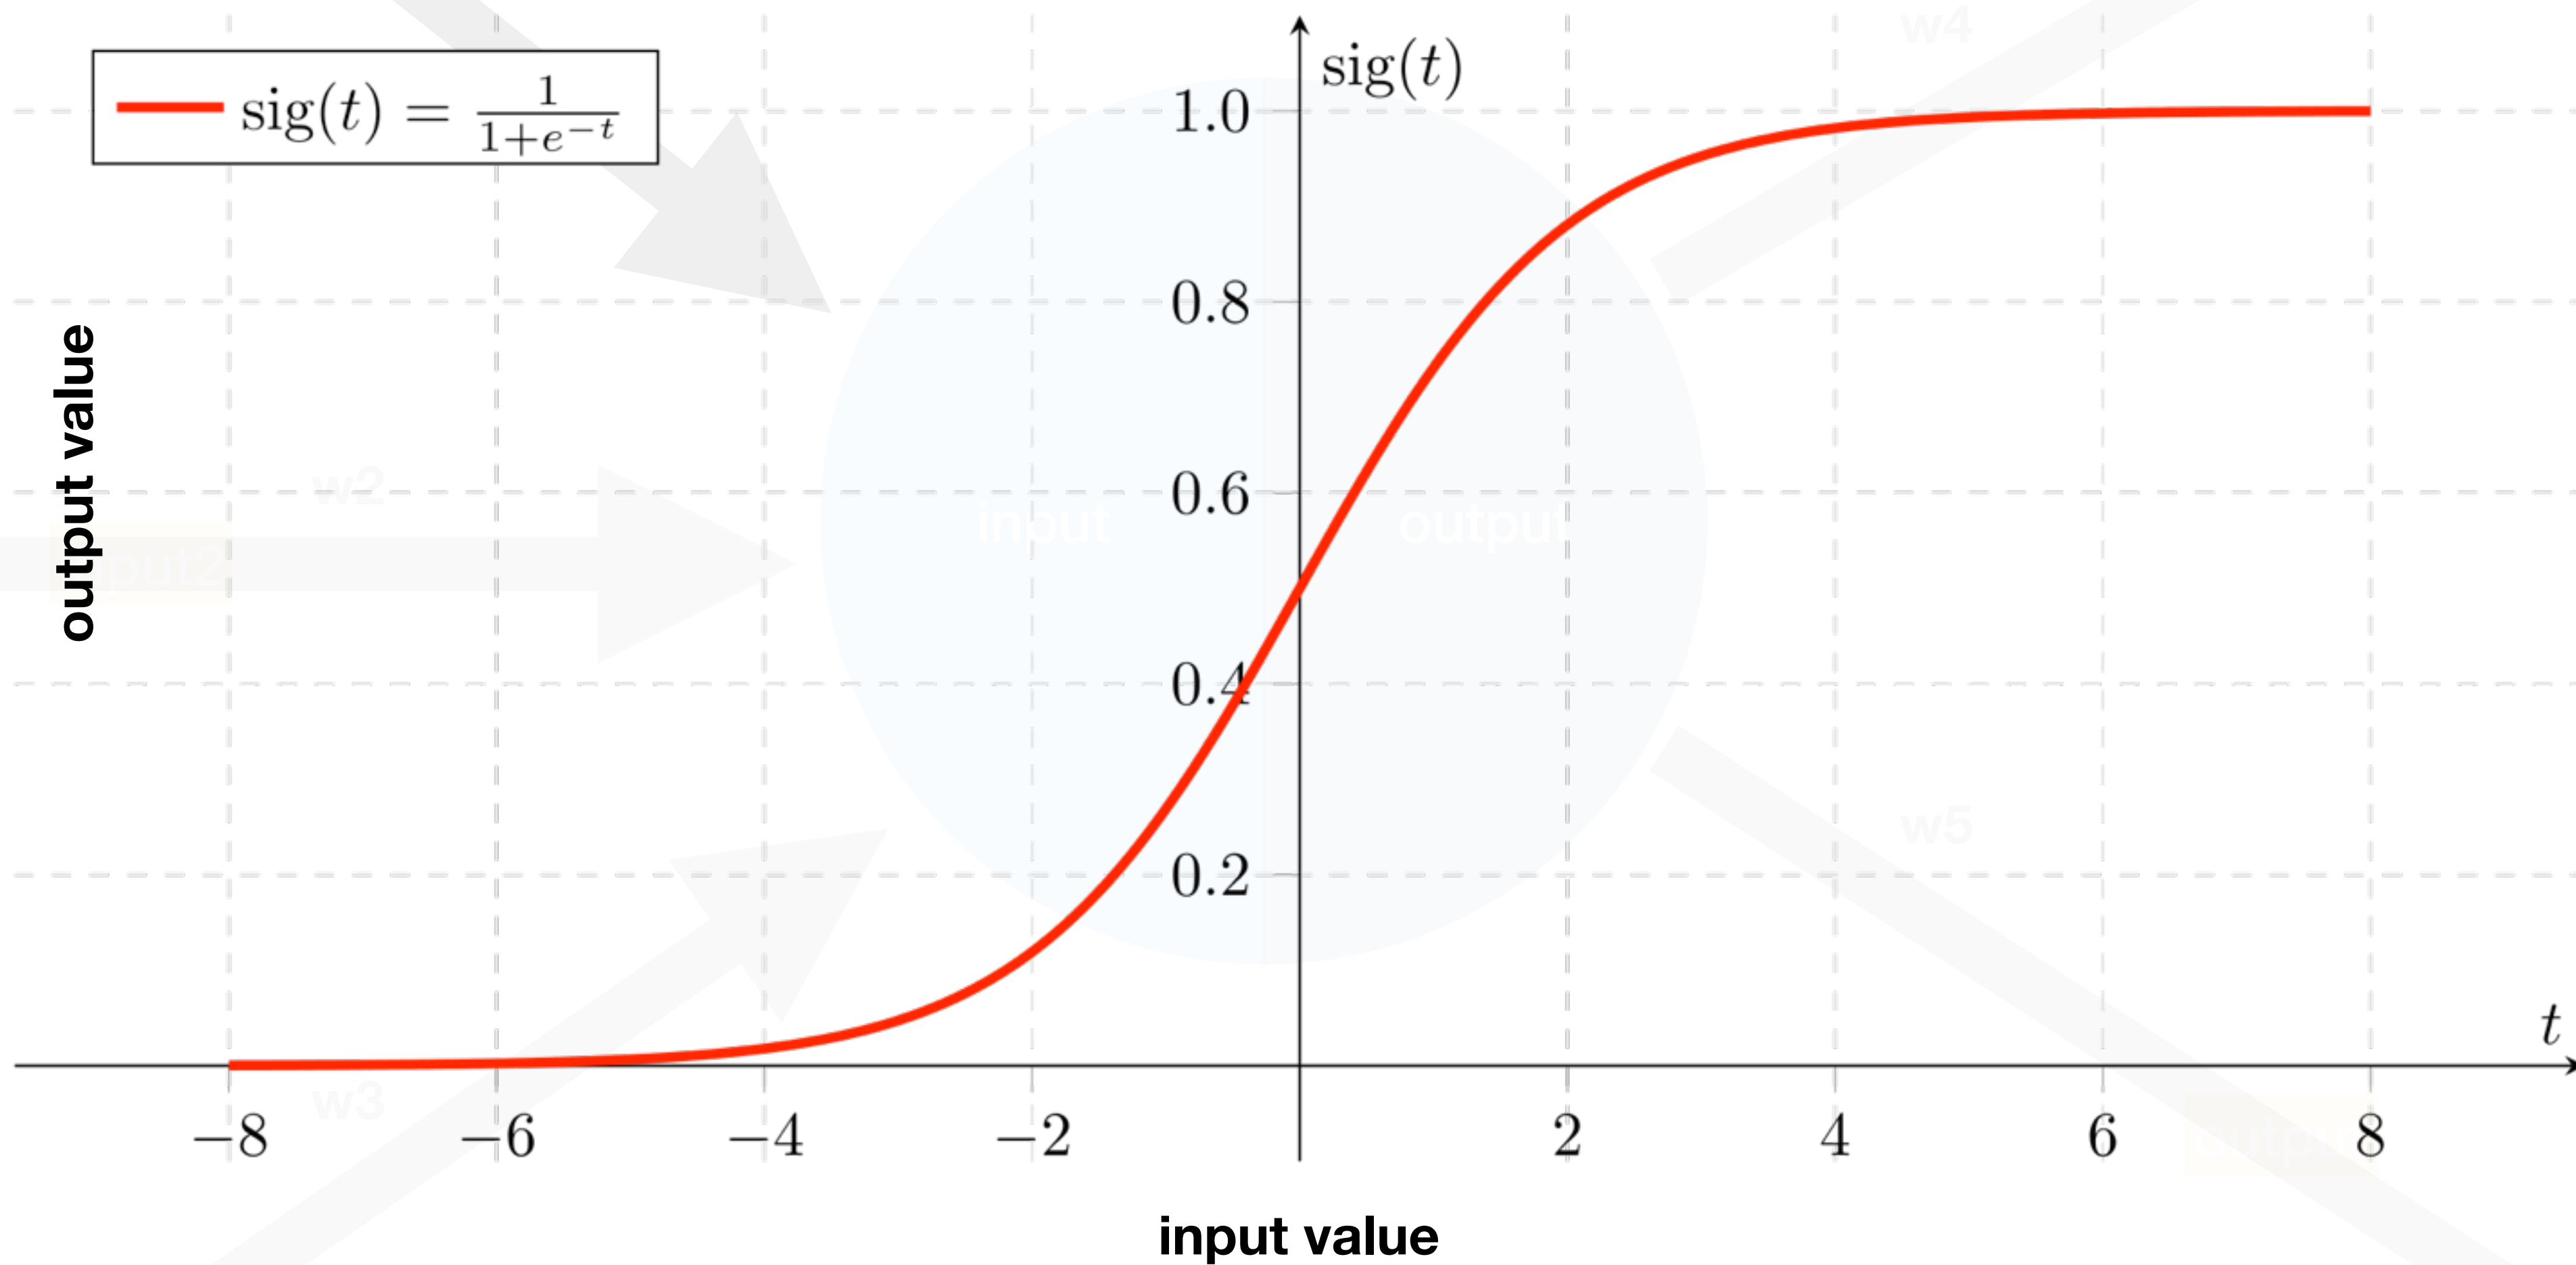


Output: Given our input, we want to decide how active this neuron is. In other words, is it "on" or "off".

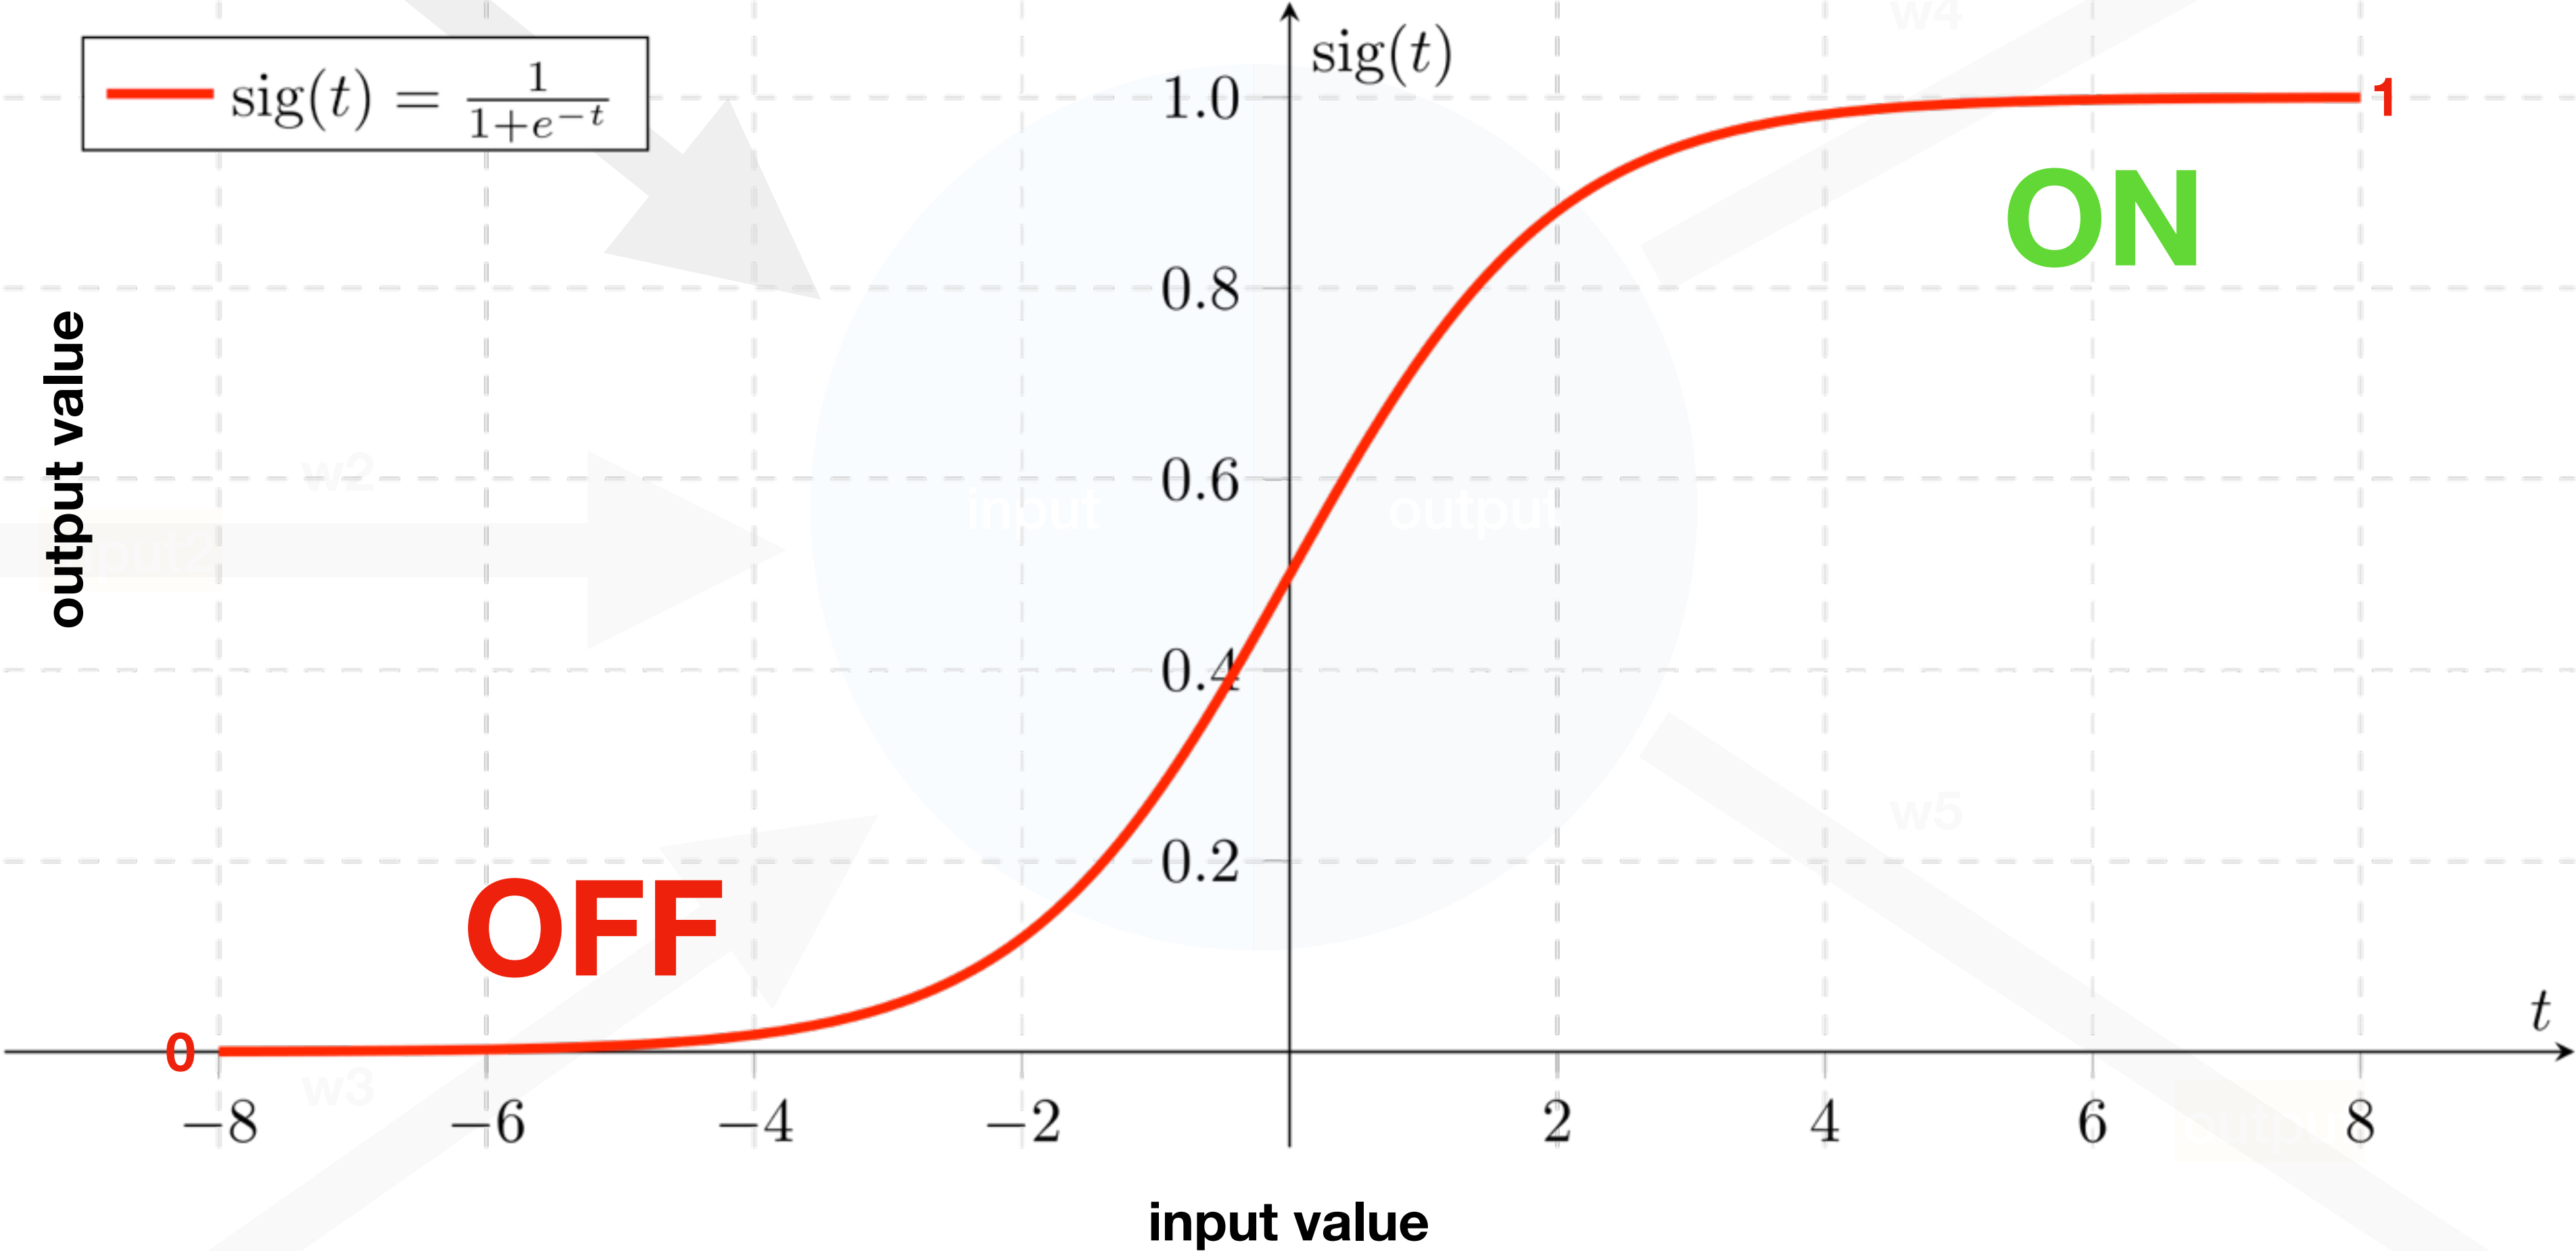
input → output

We could just use the input value unchanged, but we gain lots of nice properties if we convert in a non-linear way from input value to output value

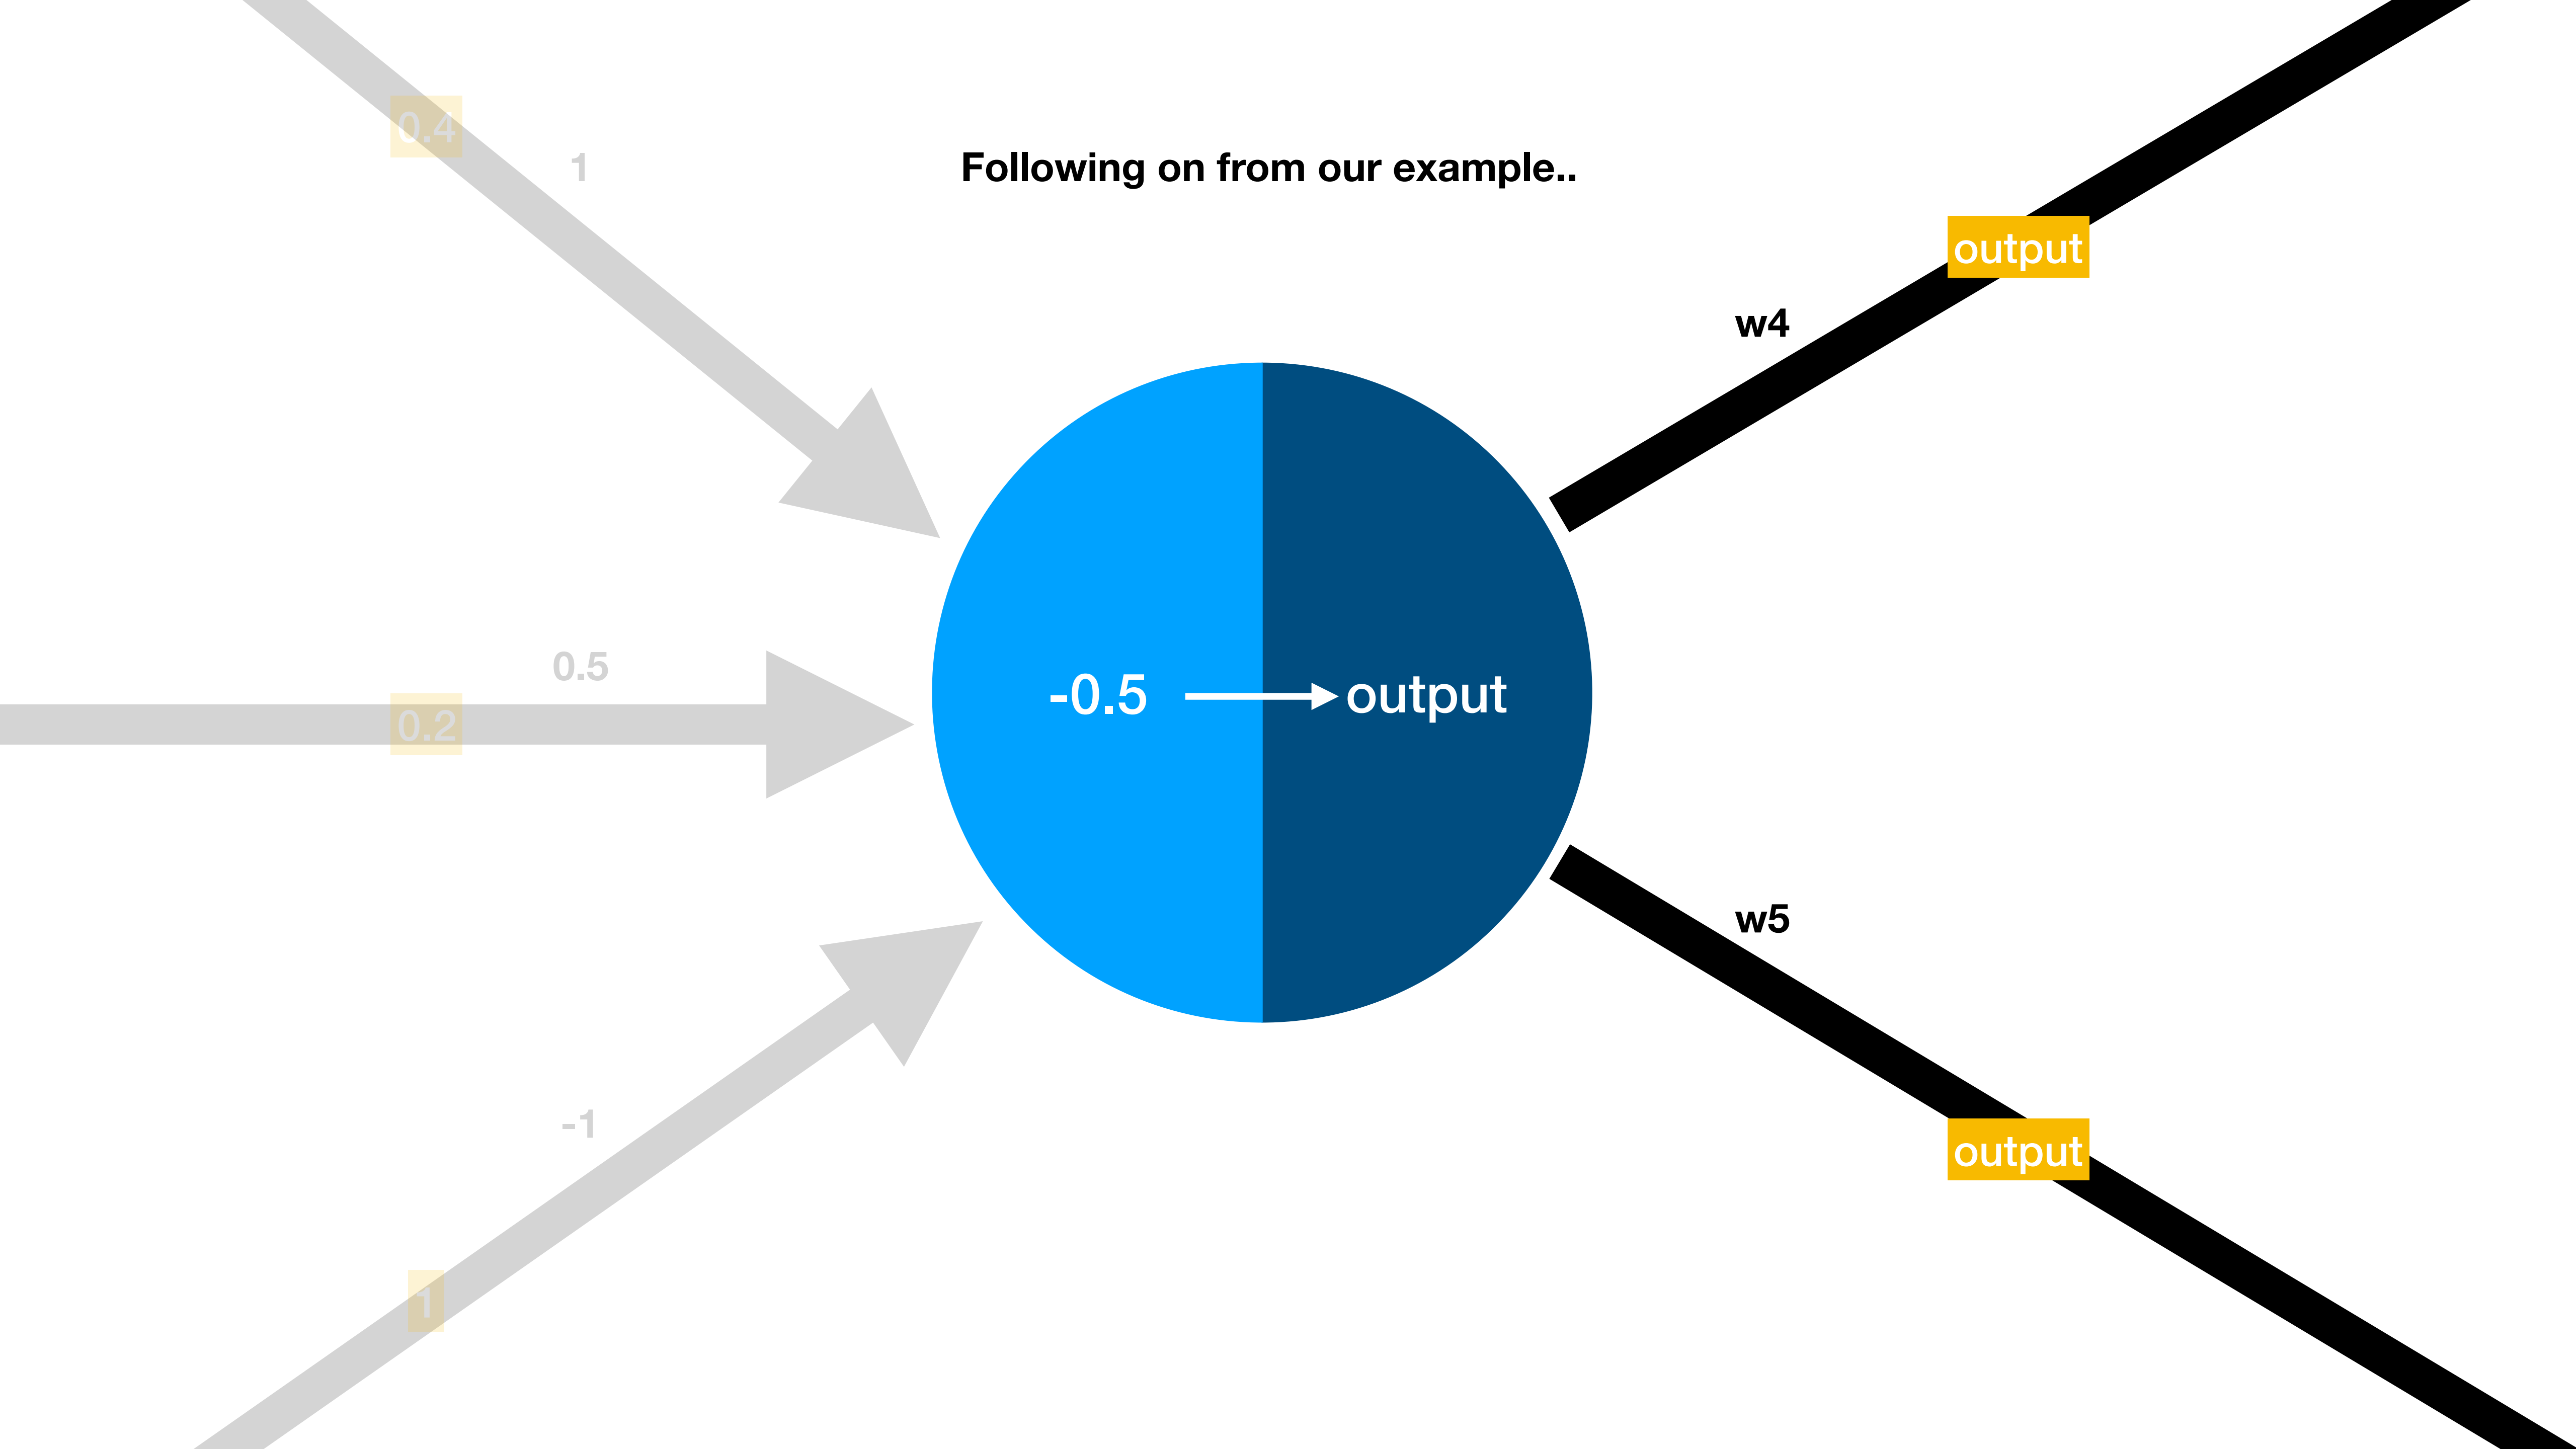
We call how we convert from input to output an *activation function*.
This is one possible choice (almost anything can be used in theory):

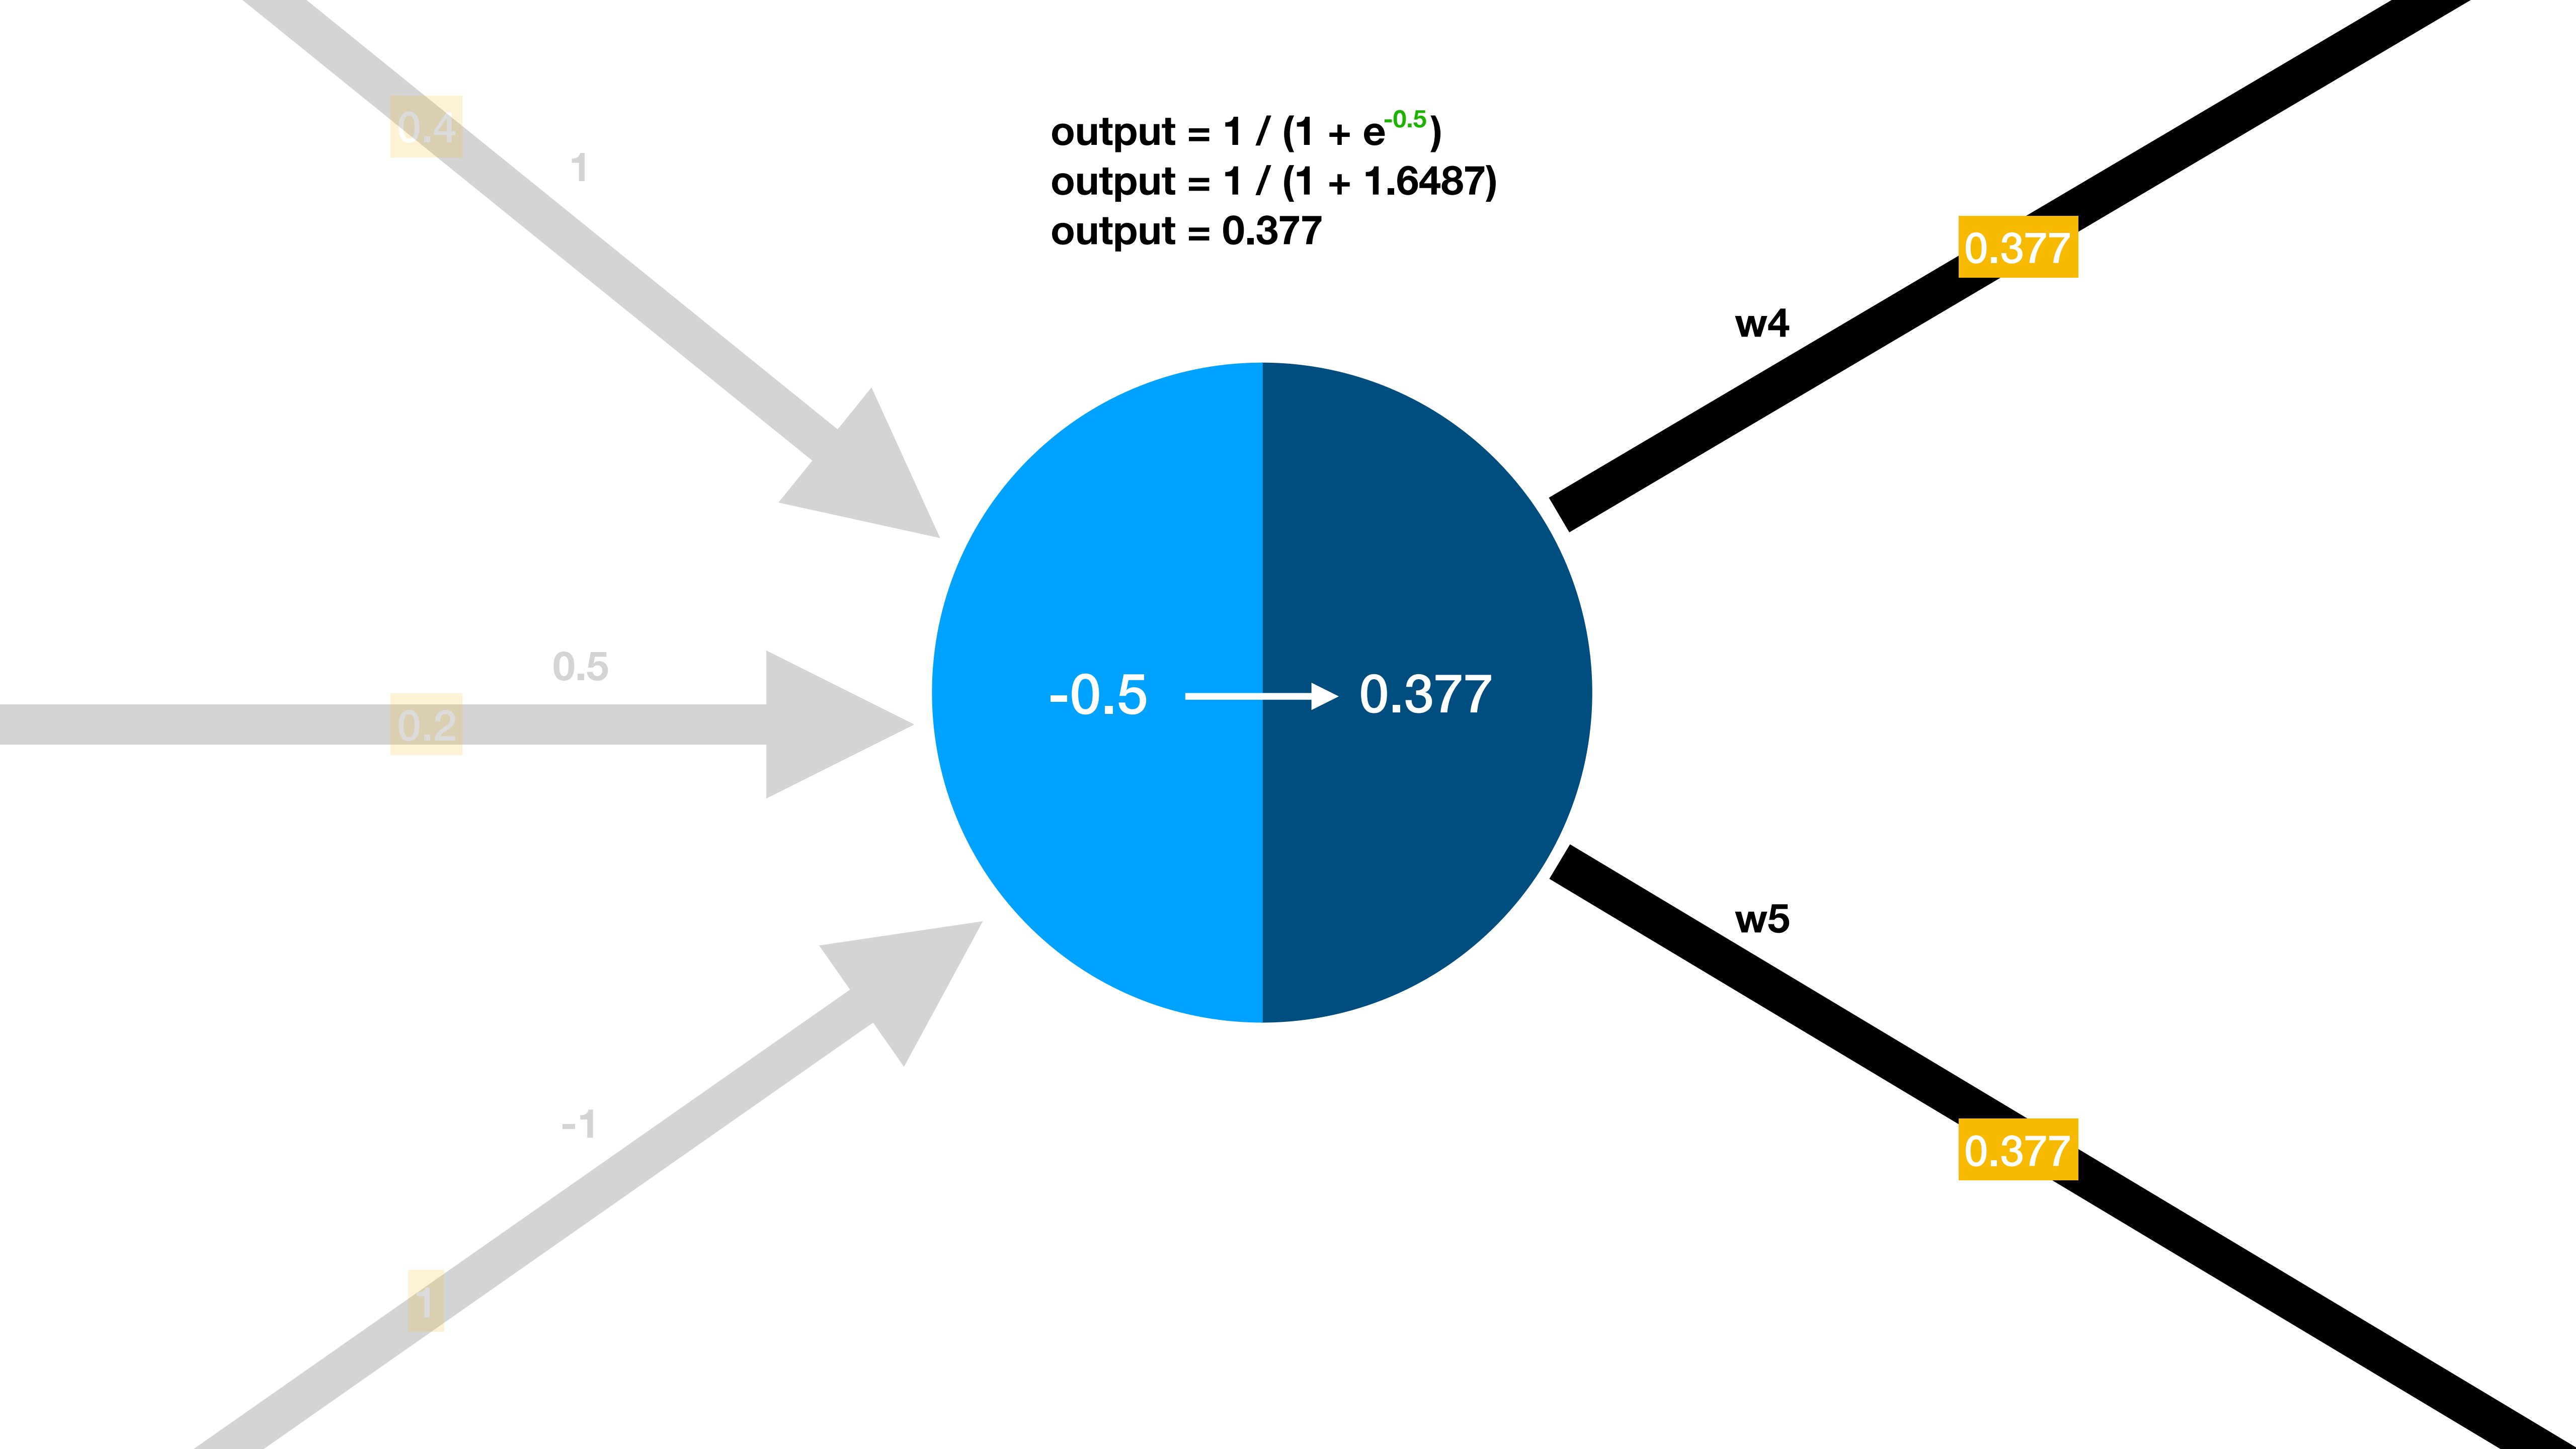


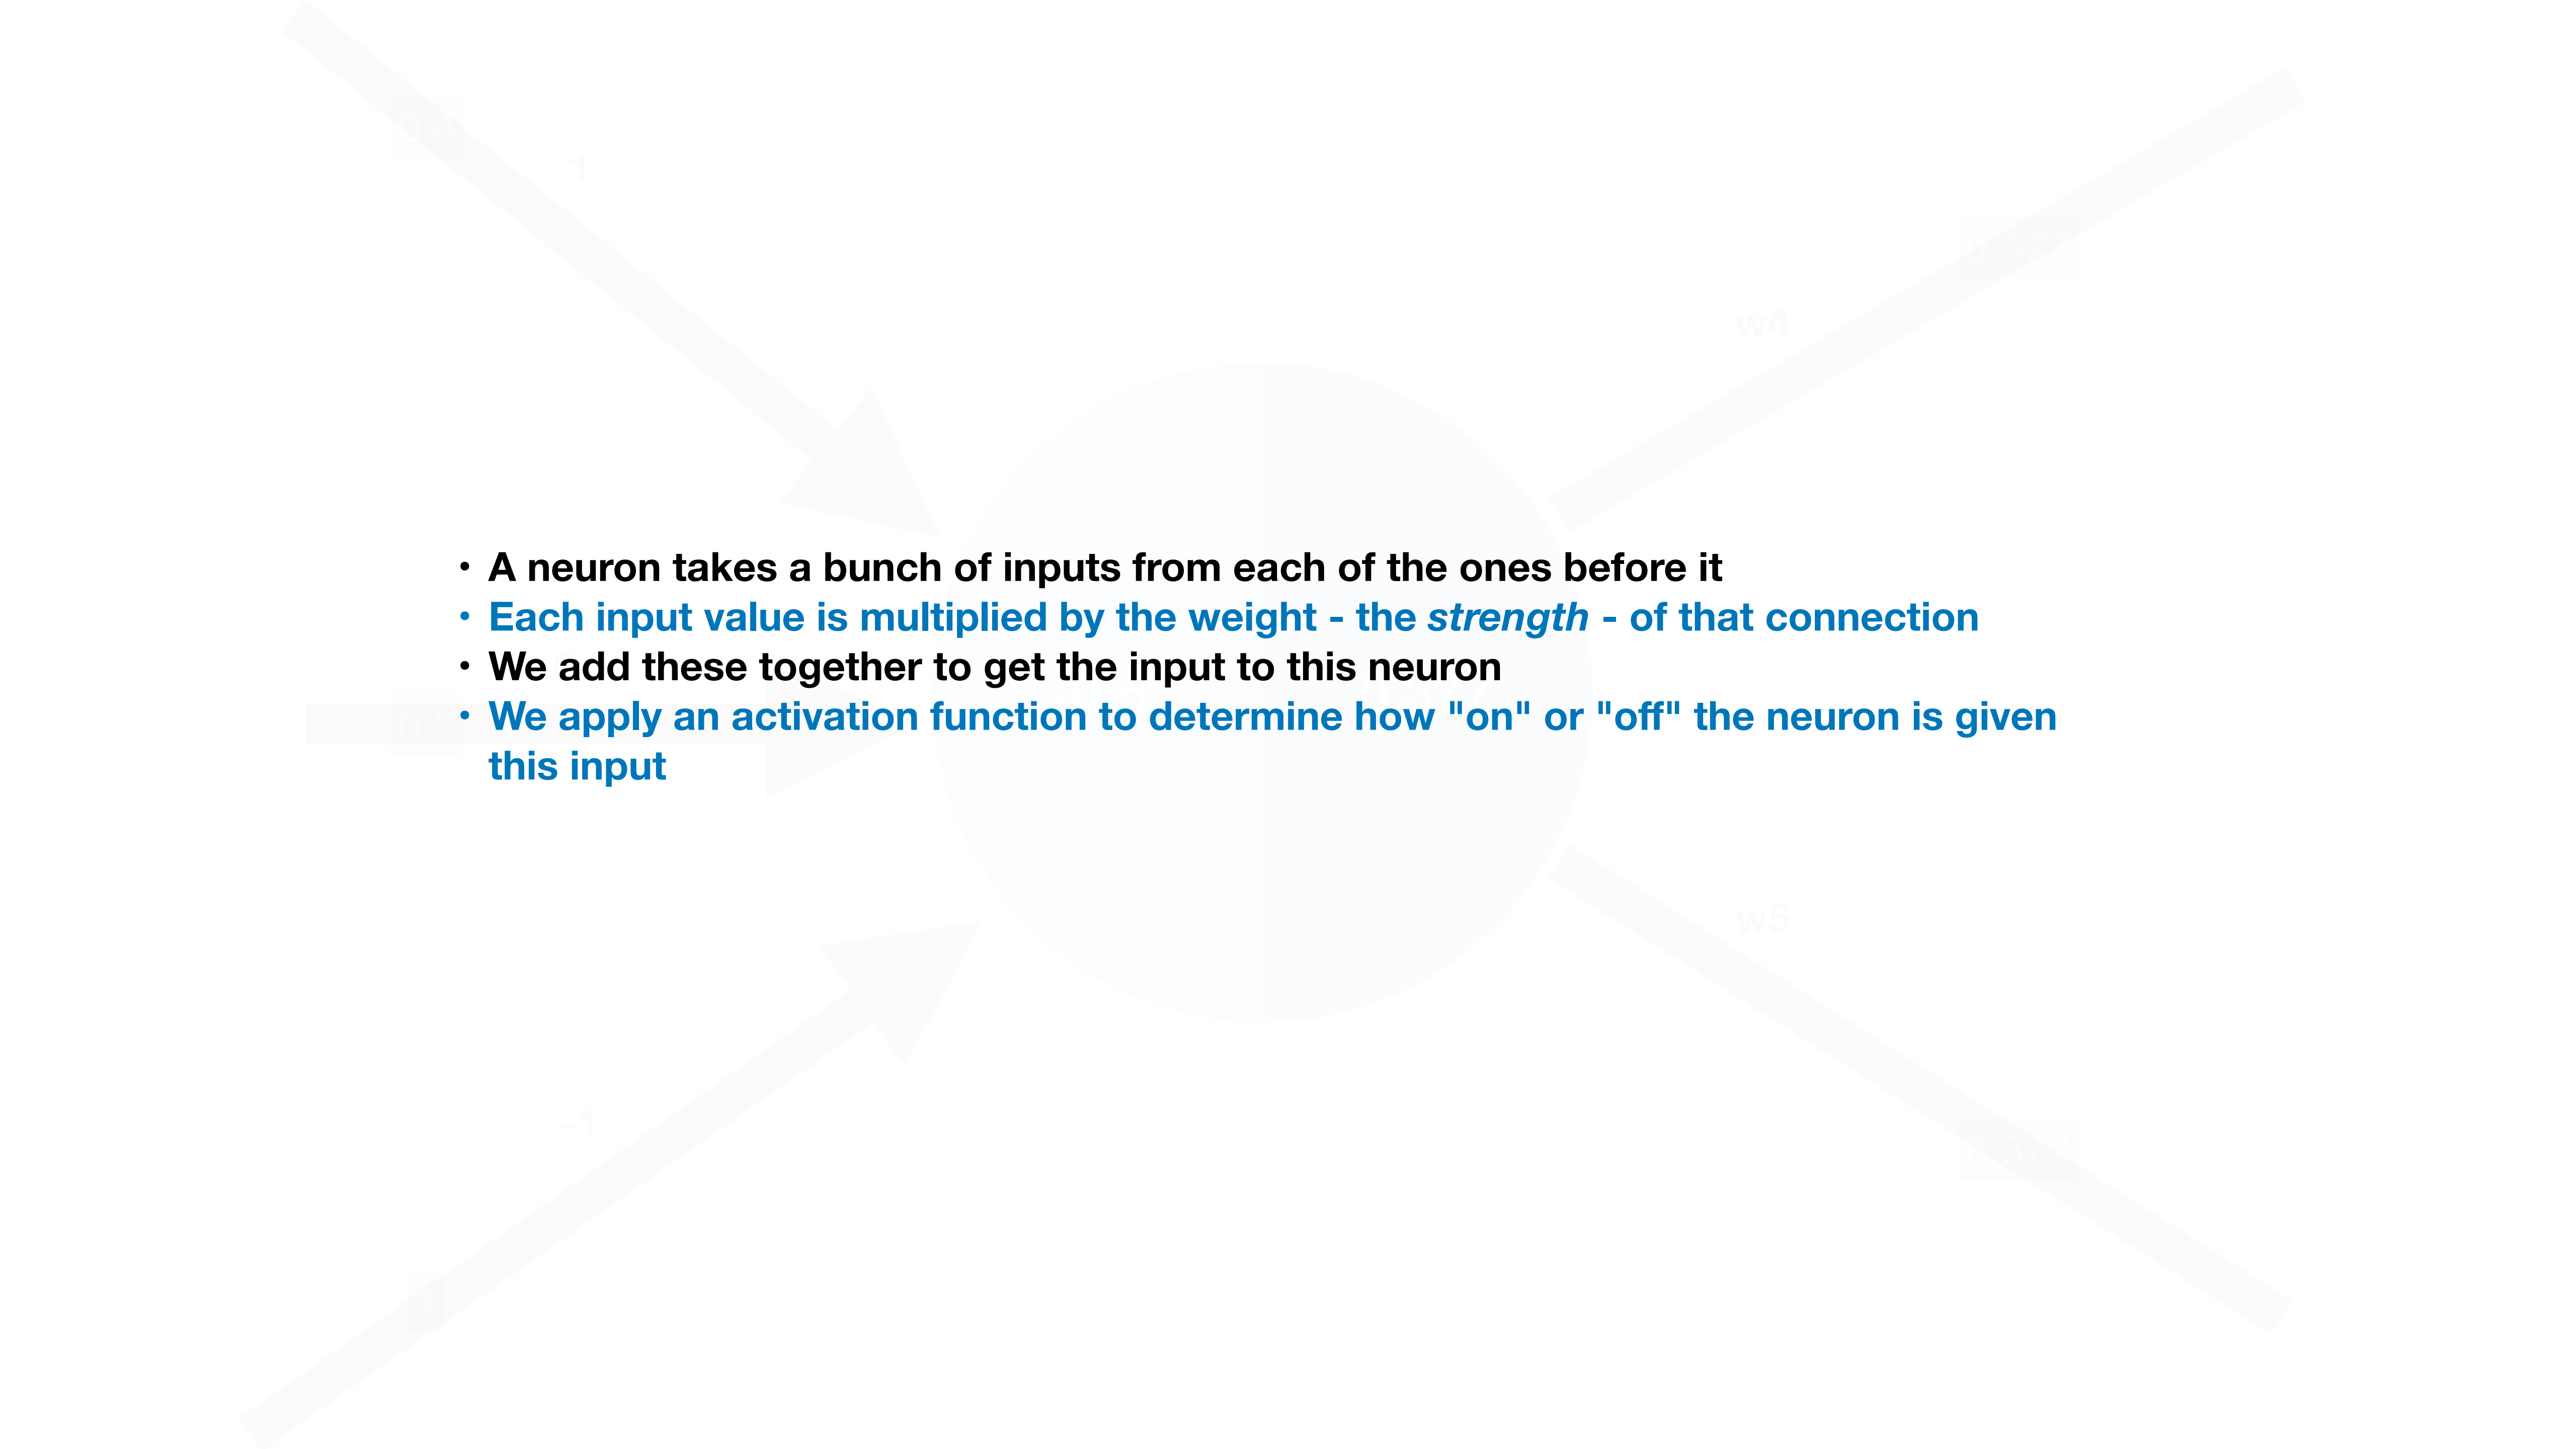
The activation function determines how "on" or "off" the neuron is, given the input to it



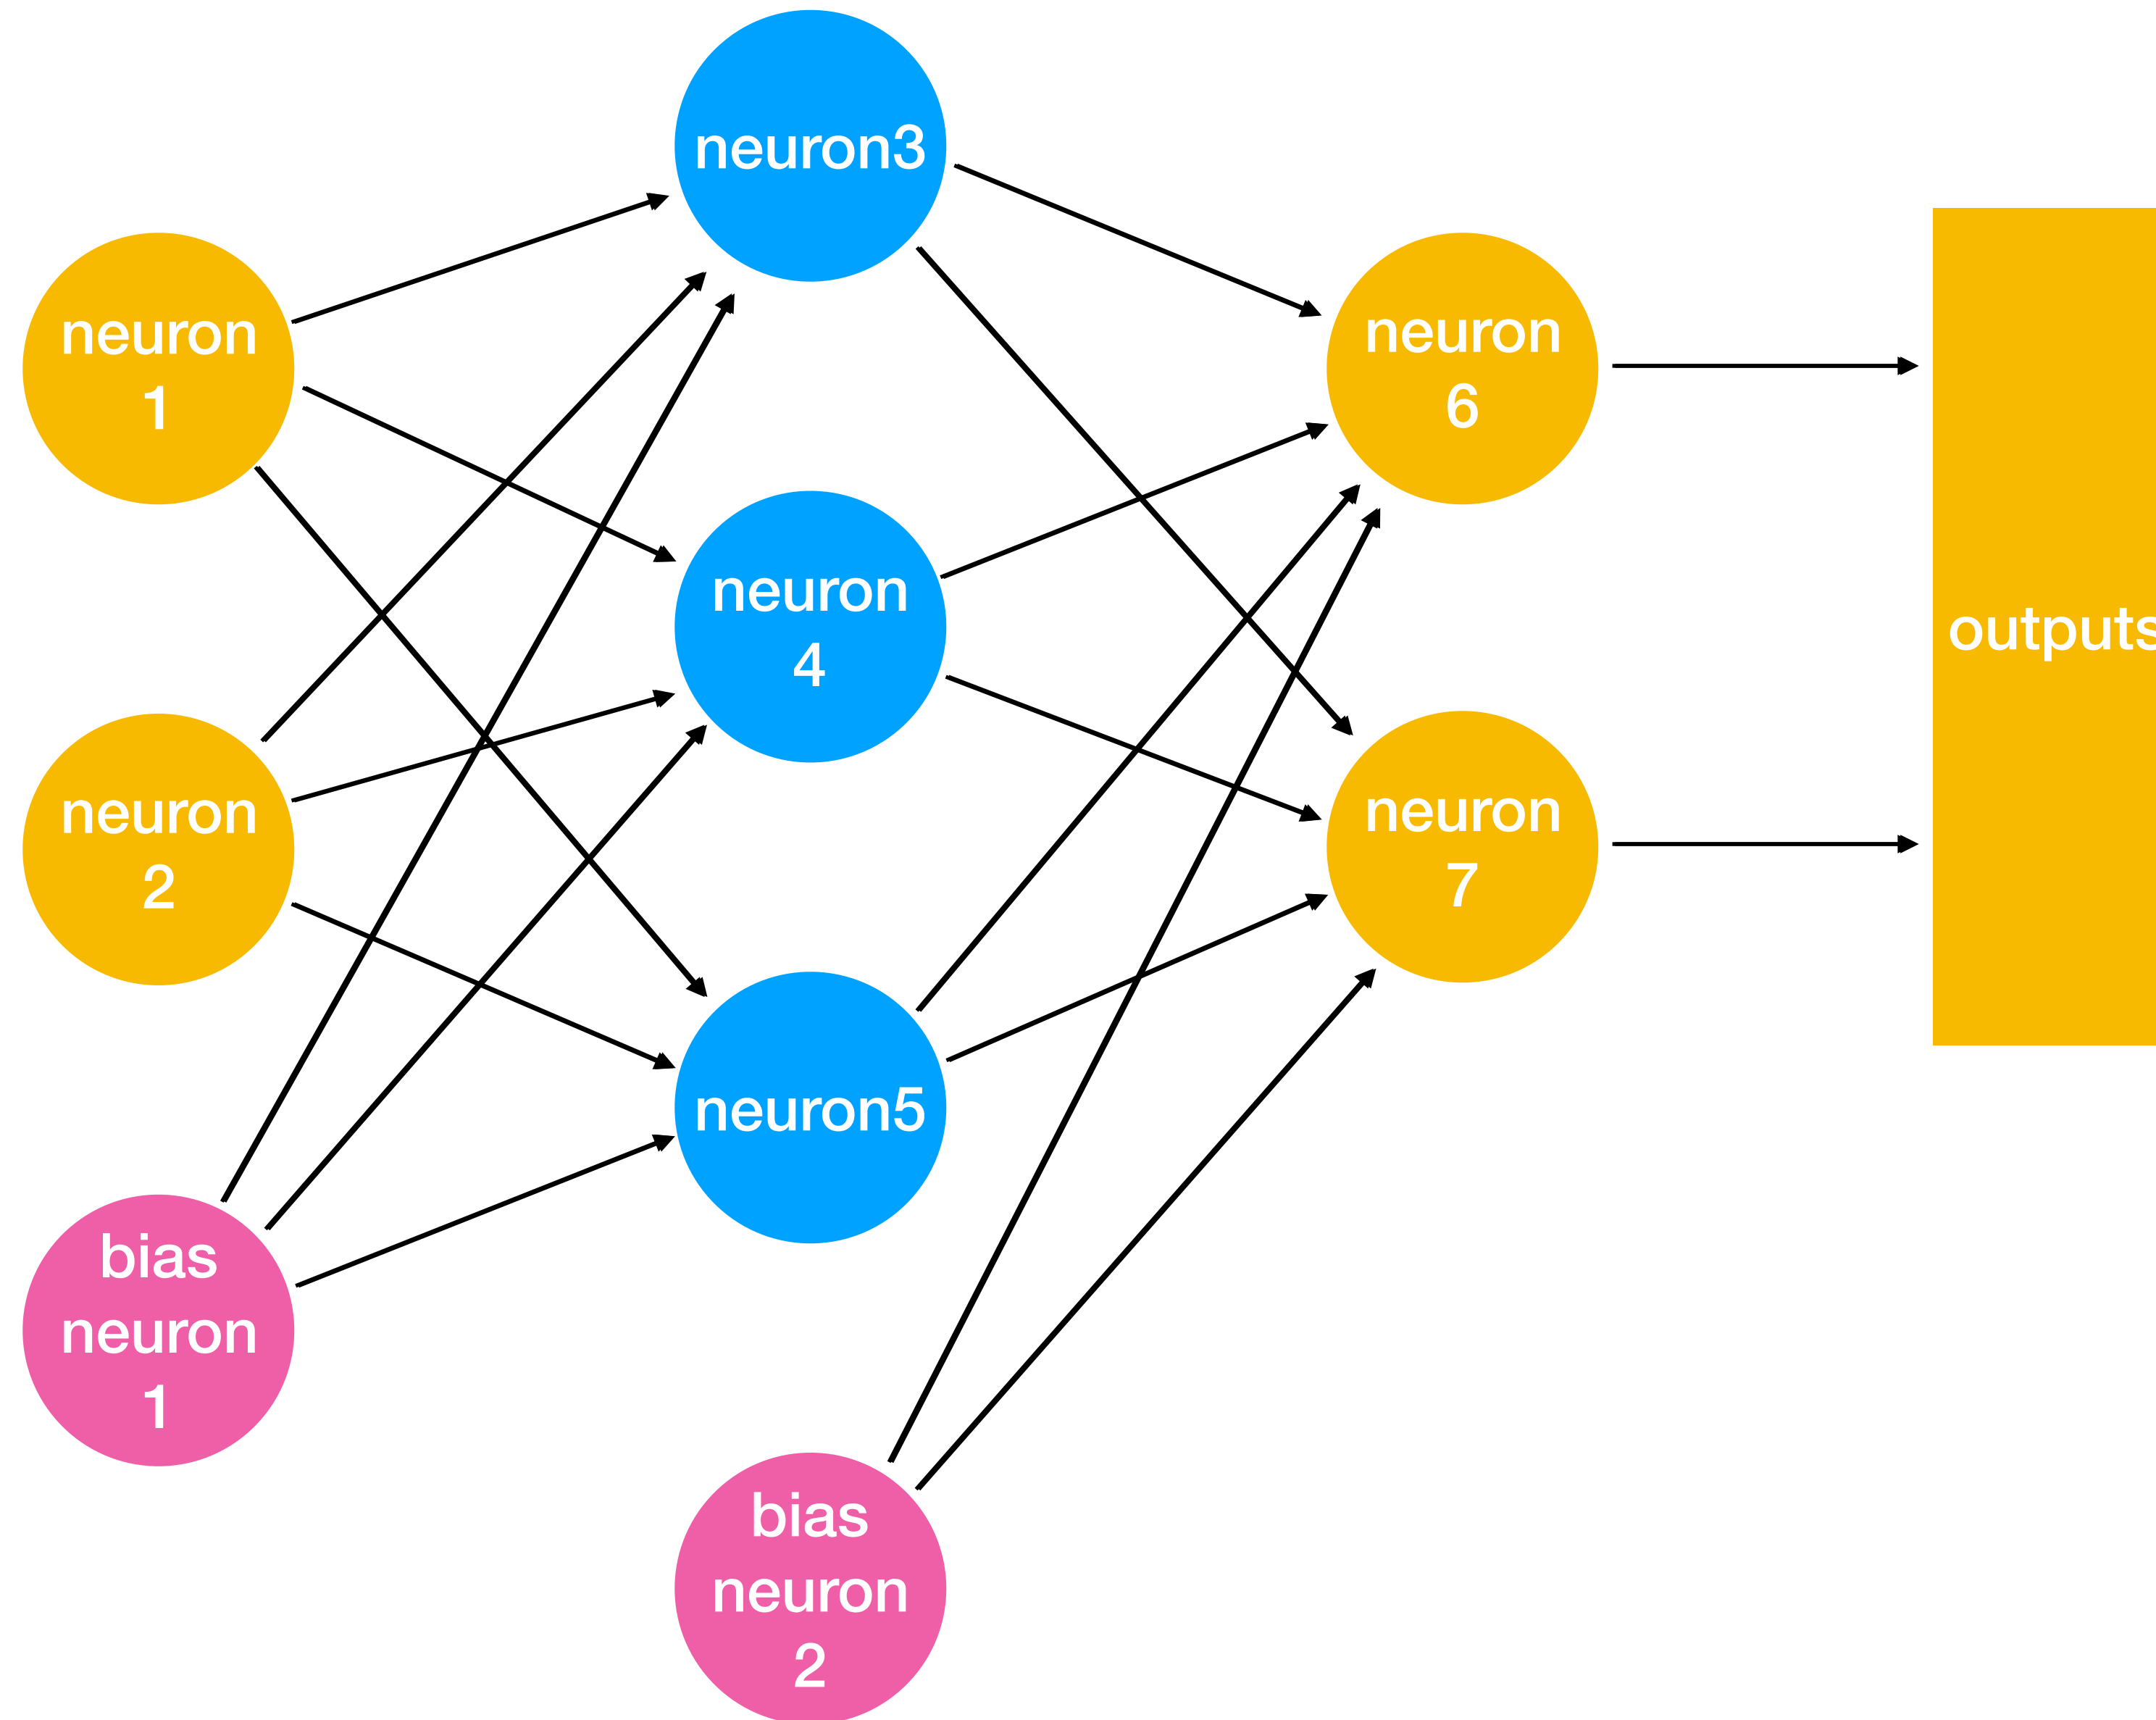
Following on from our example..



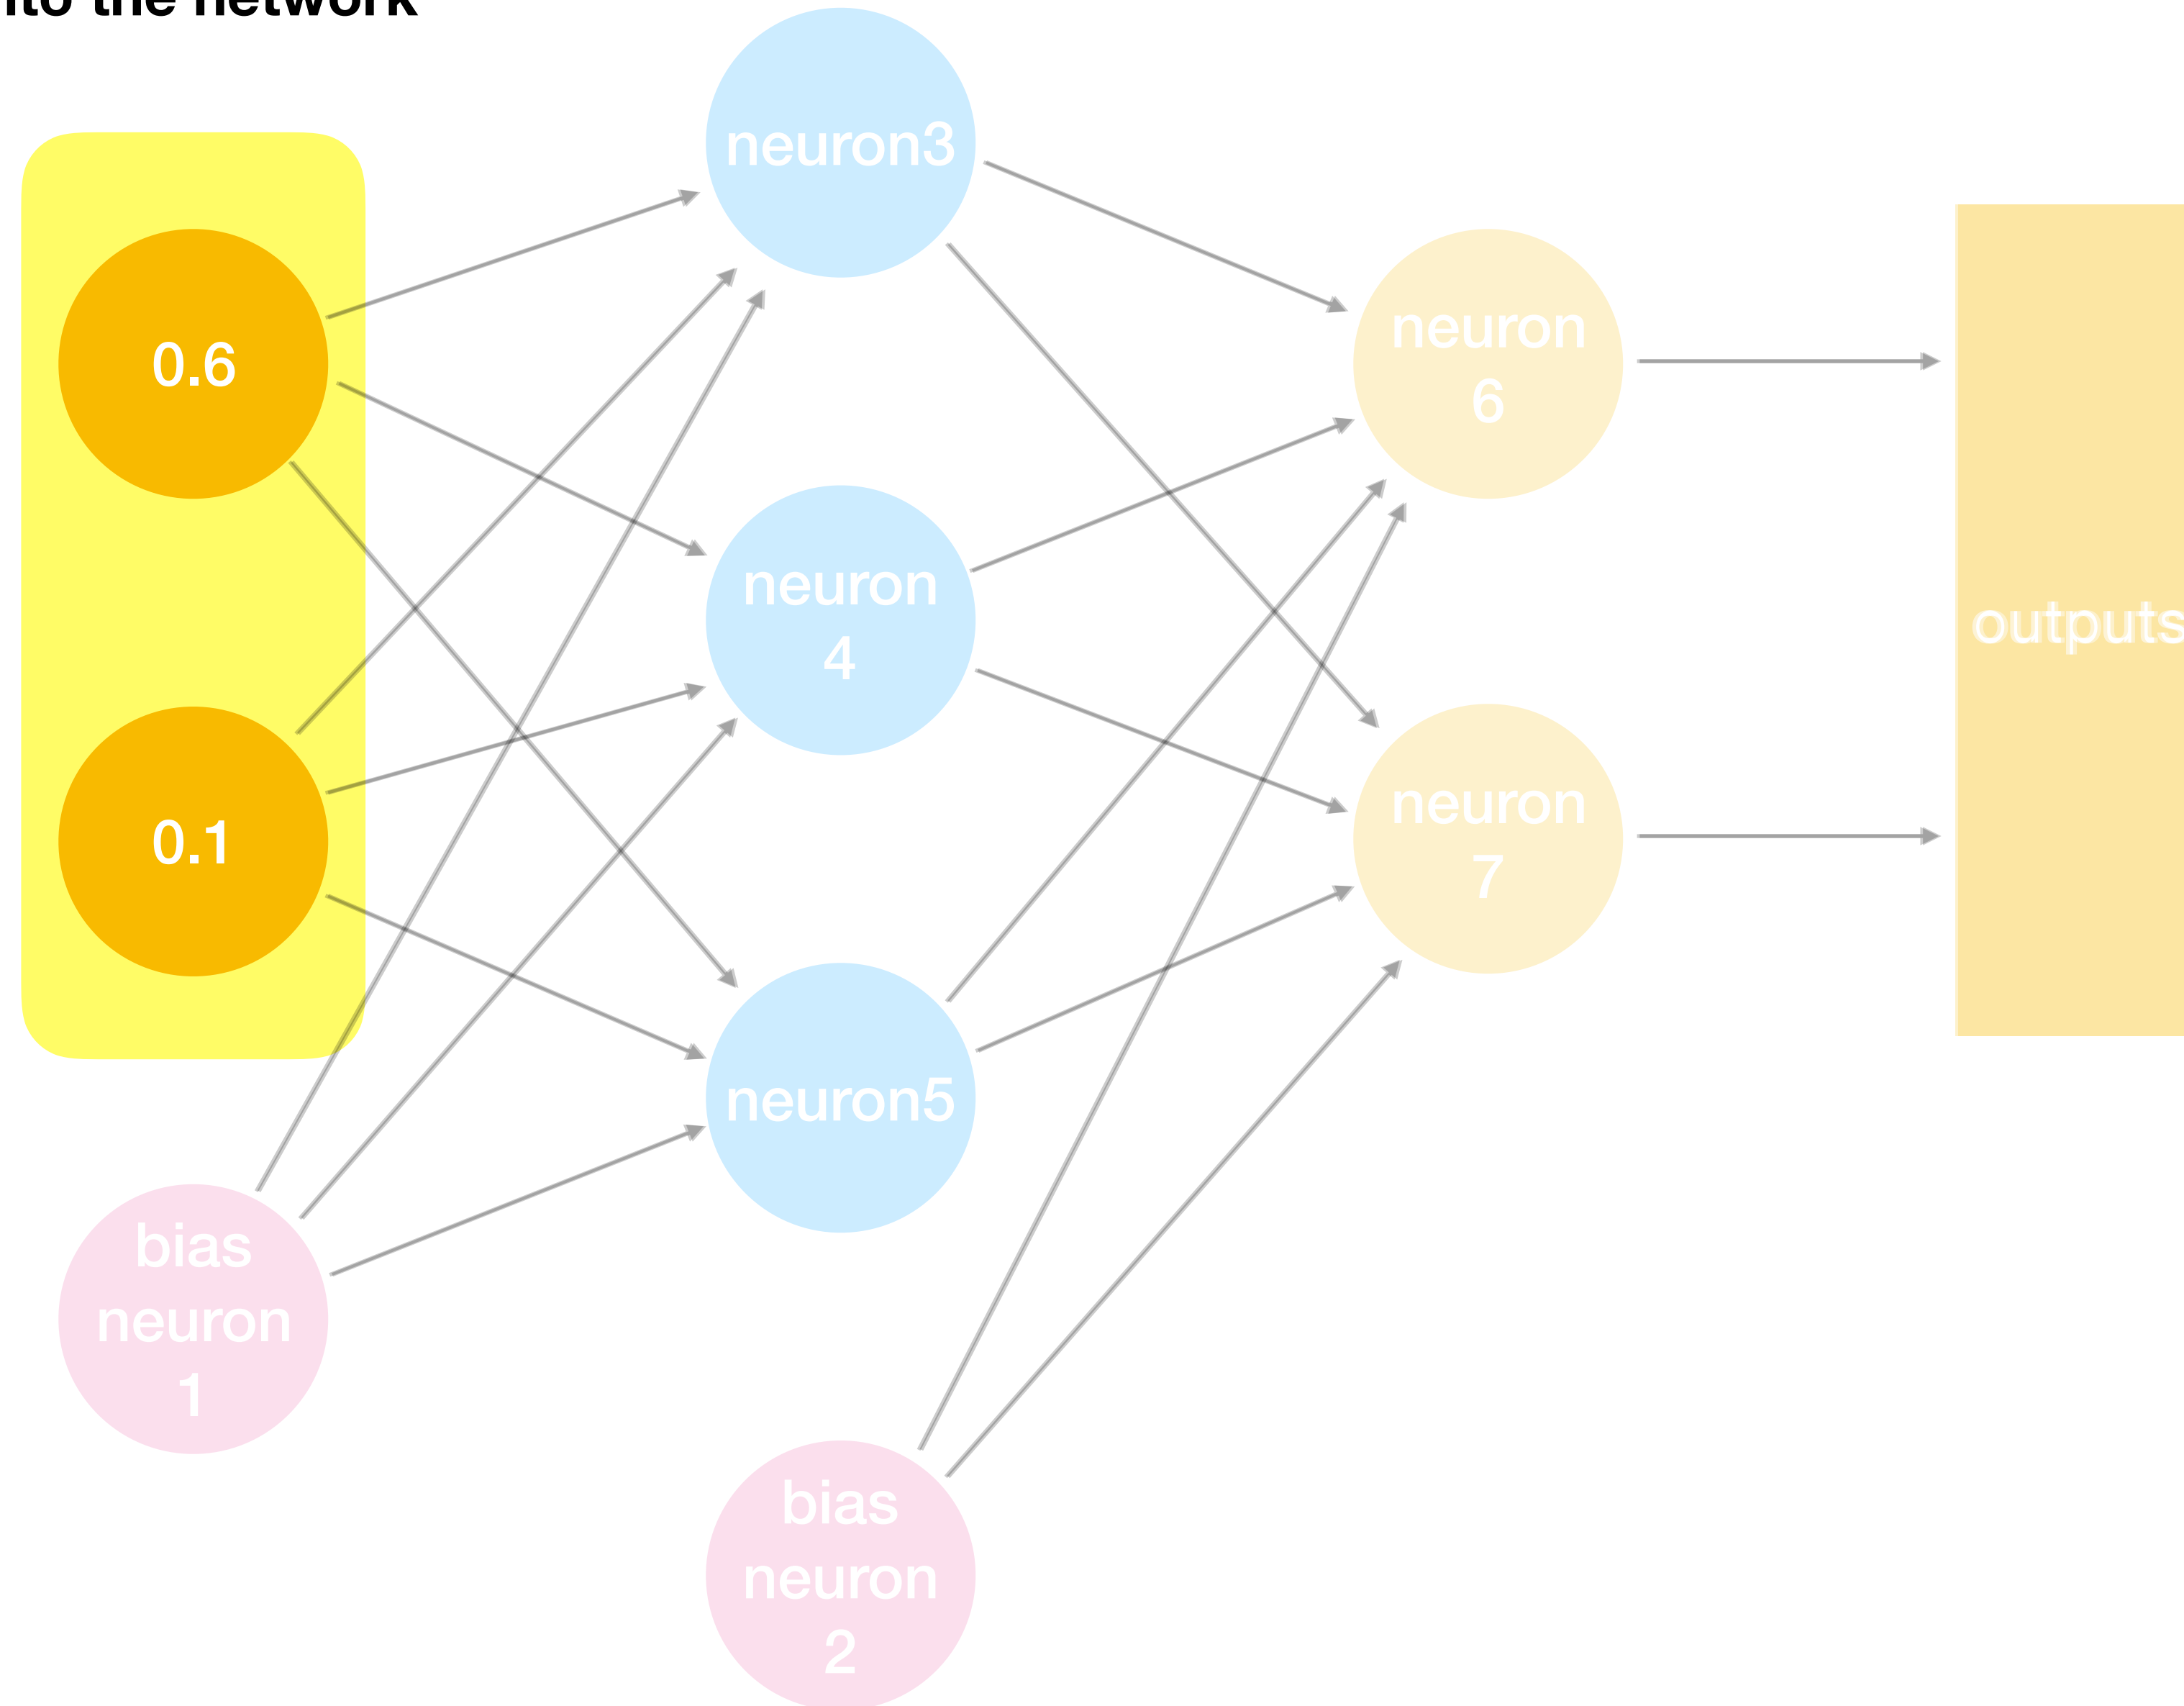


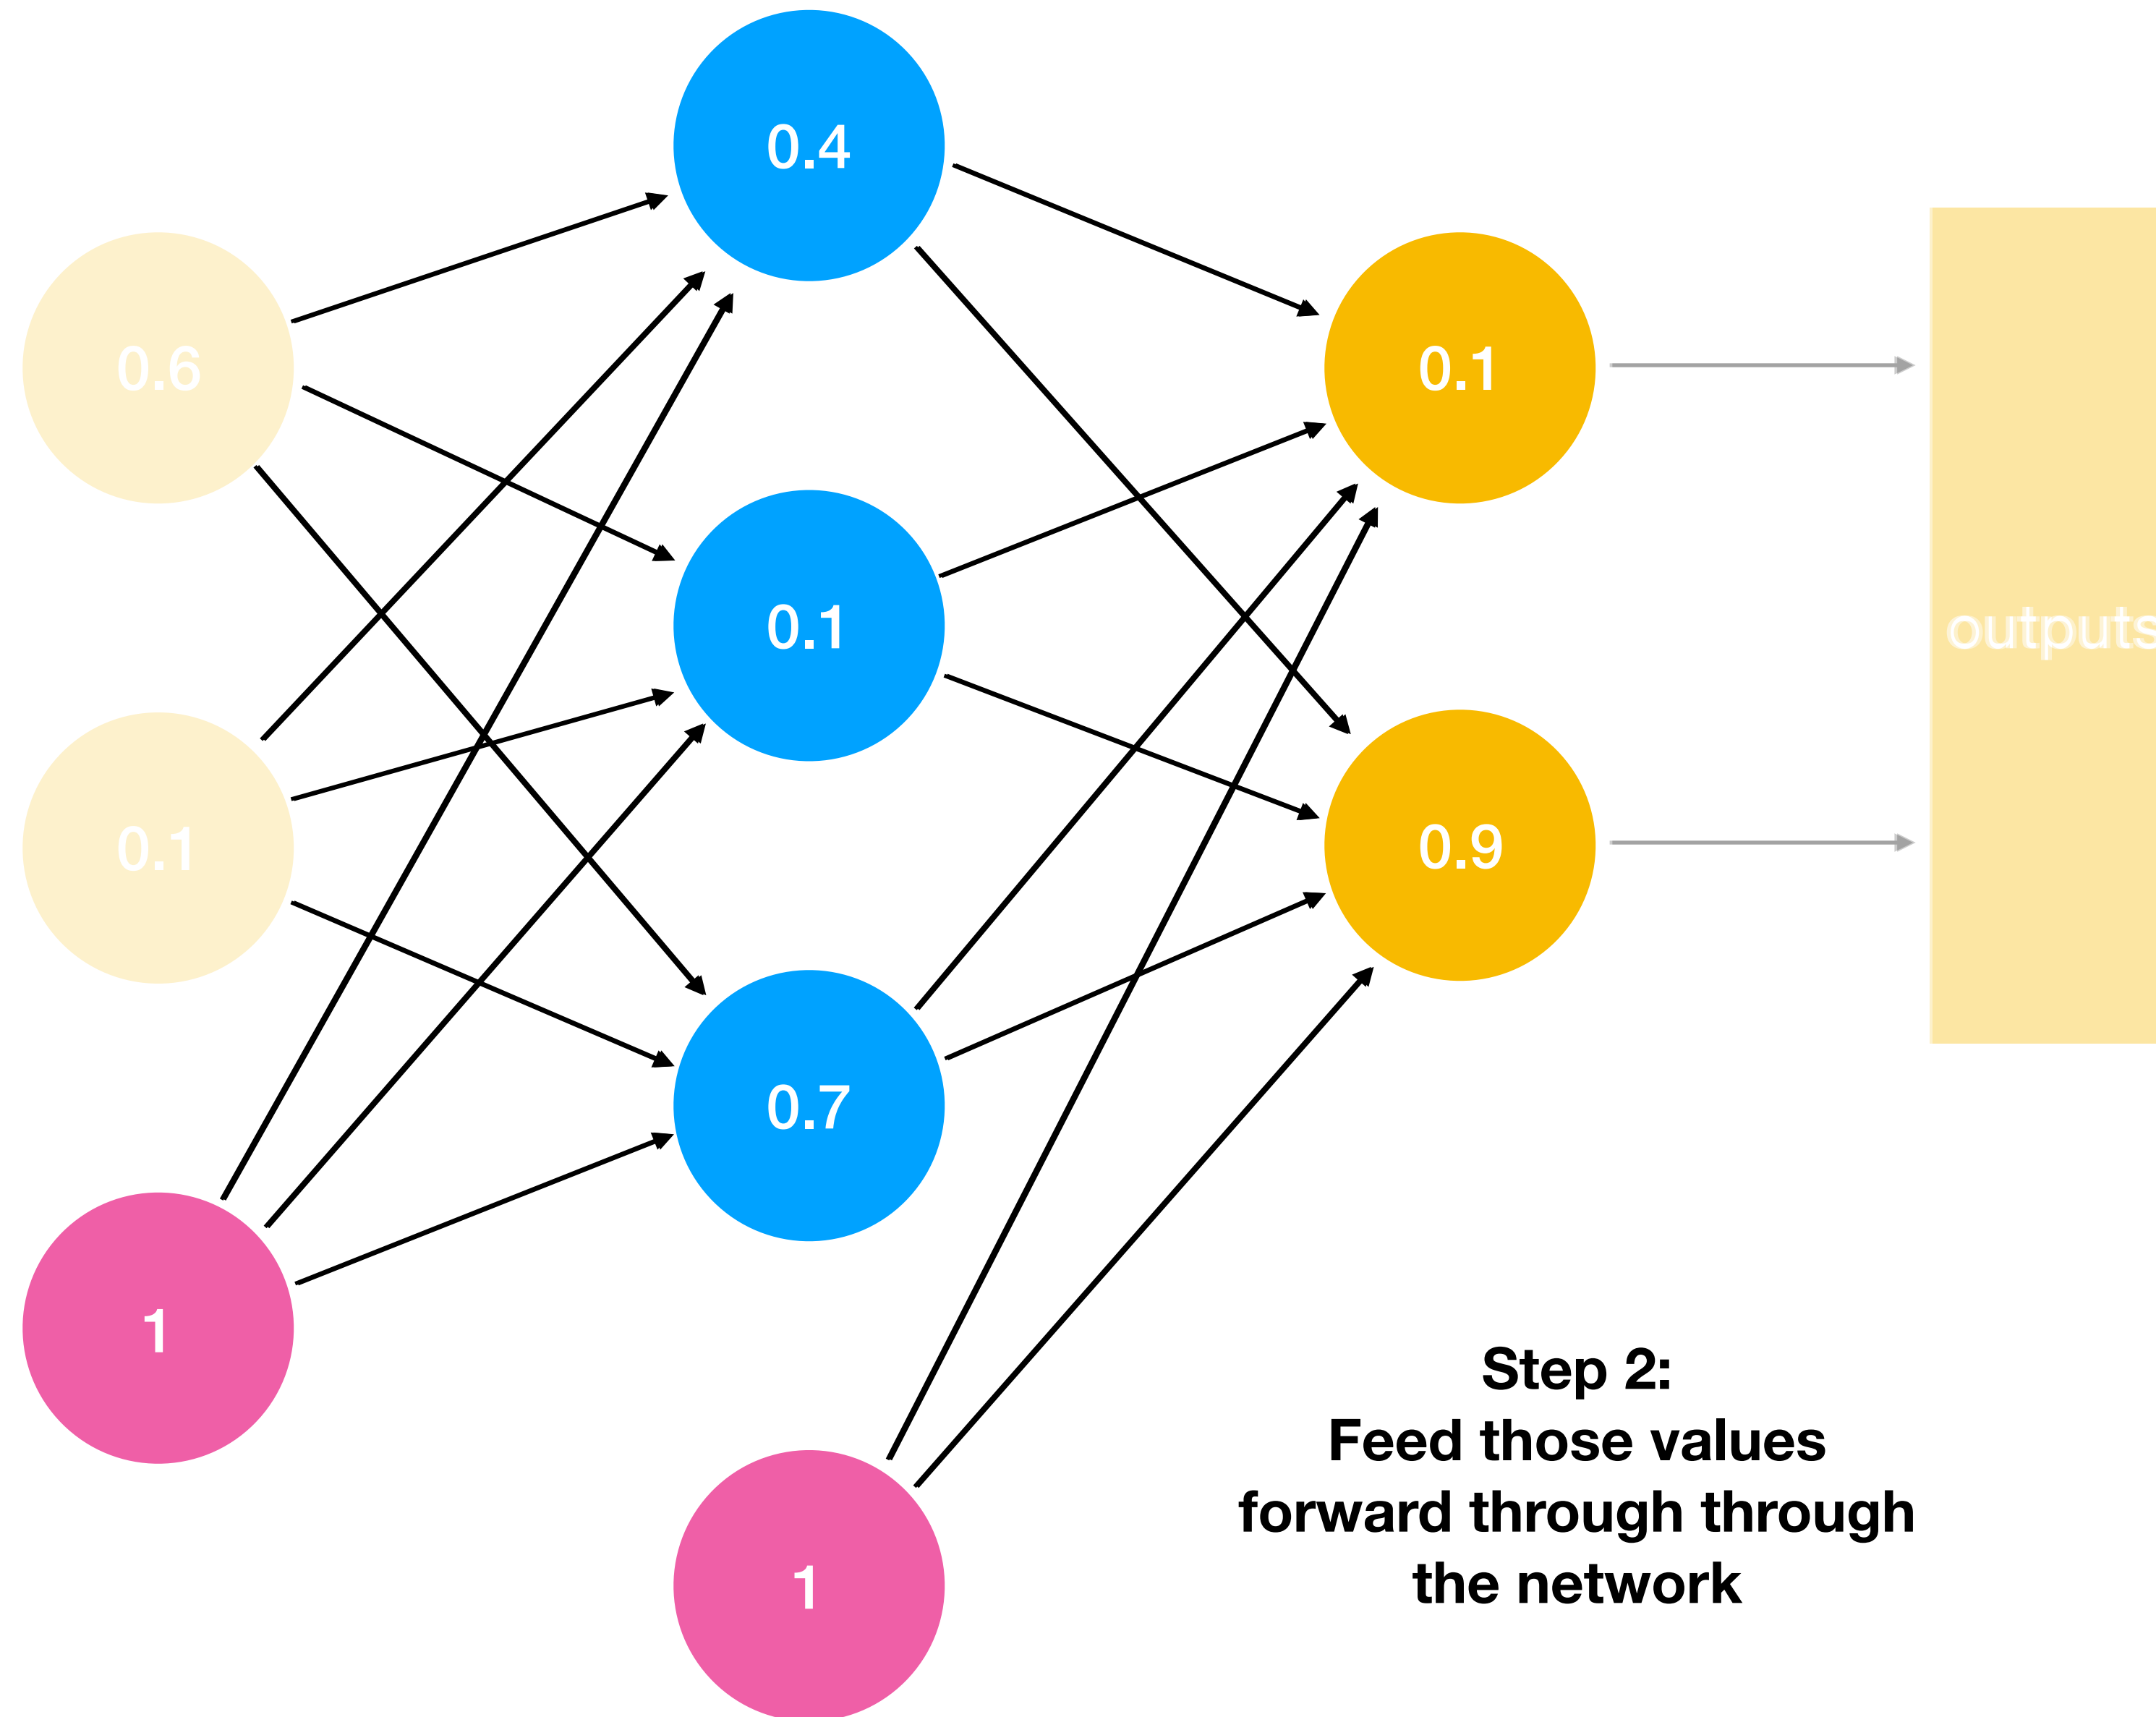
- 
- The diagram shows a central light blue circle representing a neuron. Four arrows point towards it from the corners, representing inputs. The top-left arrow is labeled '1' and has a small orange box with '0.2' next to it. The bottom-left arrow is labeled '-1' and has a small orange box with '0.2' next to it. The top-right arrow is labeled 'w4' and has a small orange box with '0.875' next to it. The bottom-right arrow is labeled 'w5' and has a small orange box with '0.375' next to it. The text in the center explains the process of a neuron taking inputs, multiplying them by weights, summing them, and applying an activation function.
- A neuron takes a bunch of inputs from each of the ones before it
 - Each input value is multiplied by the weight - the *strength* - of that connection
 - We add these together to get the input to this neuron
 - We apply an activation function to determine how "on" or "off" the neuron is given this input

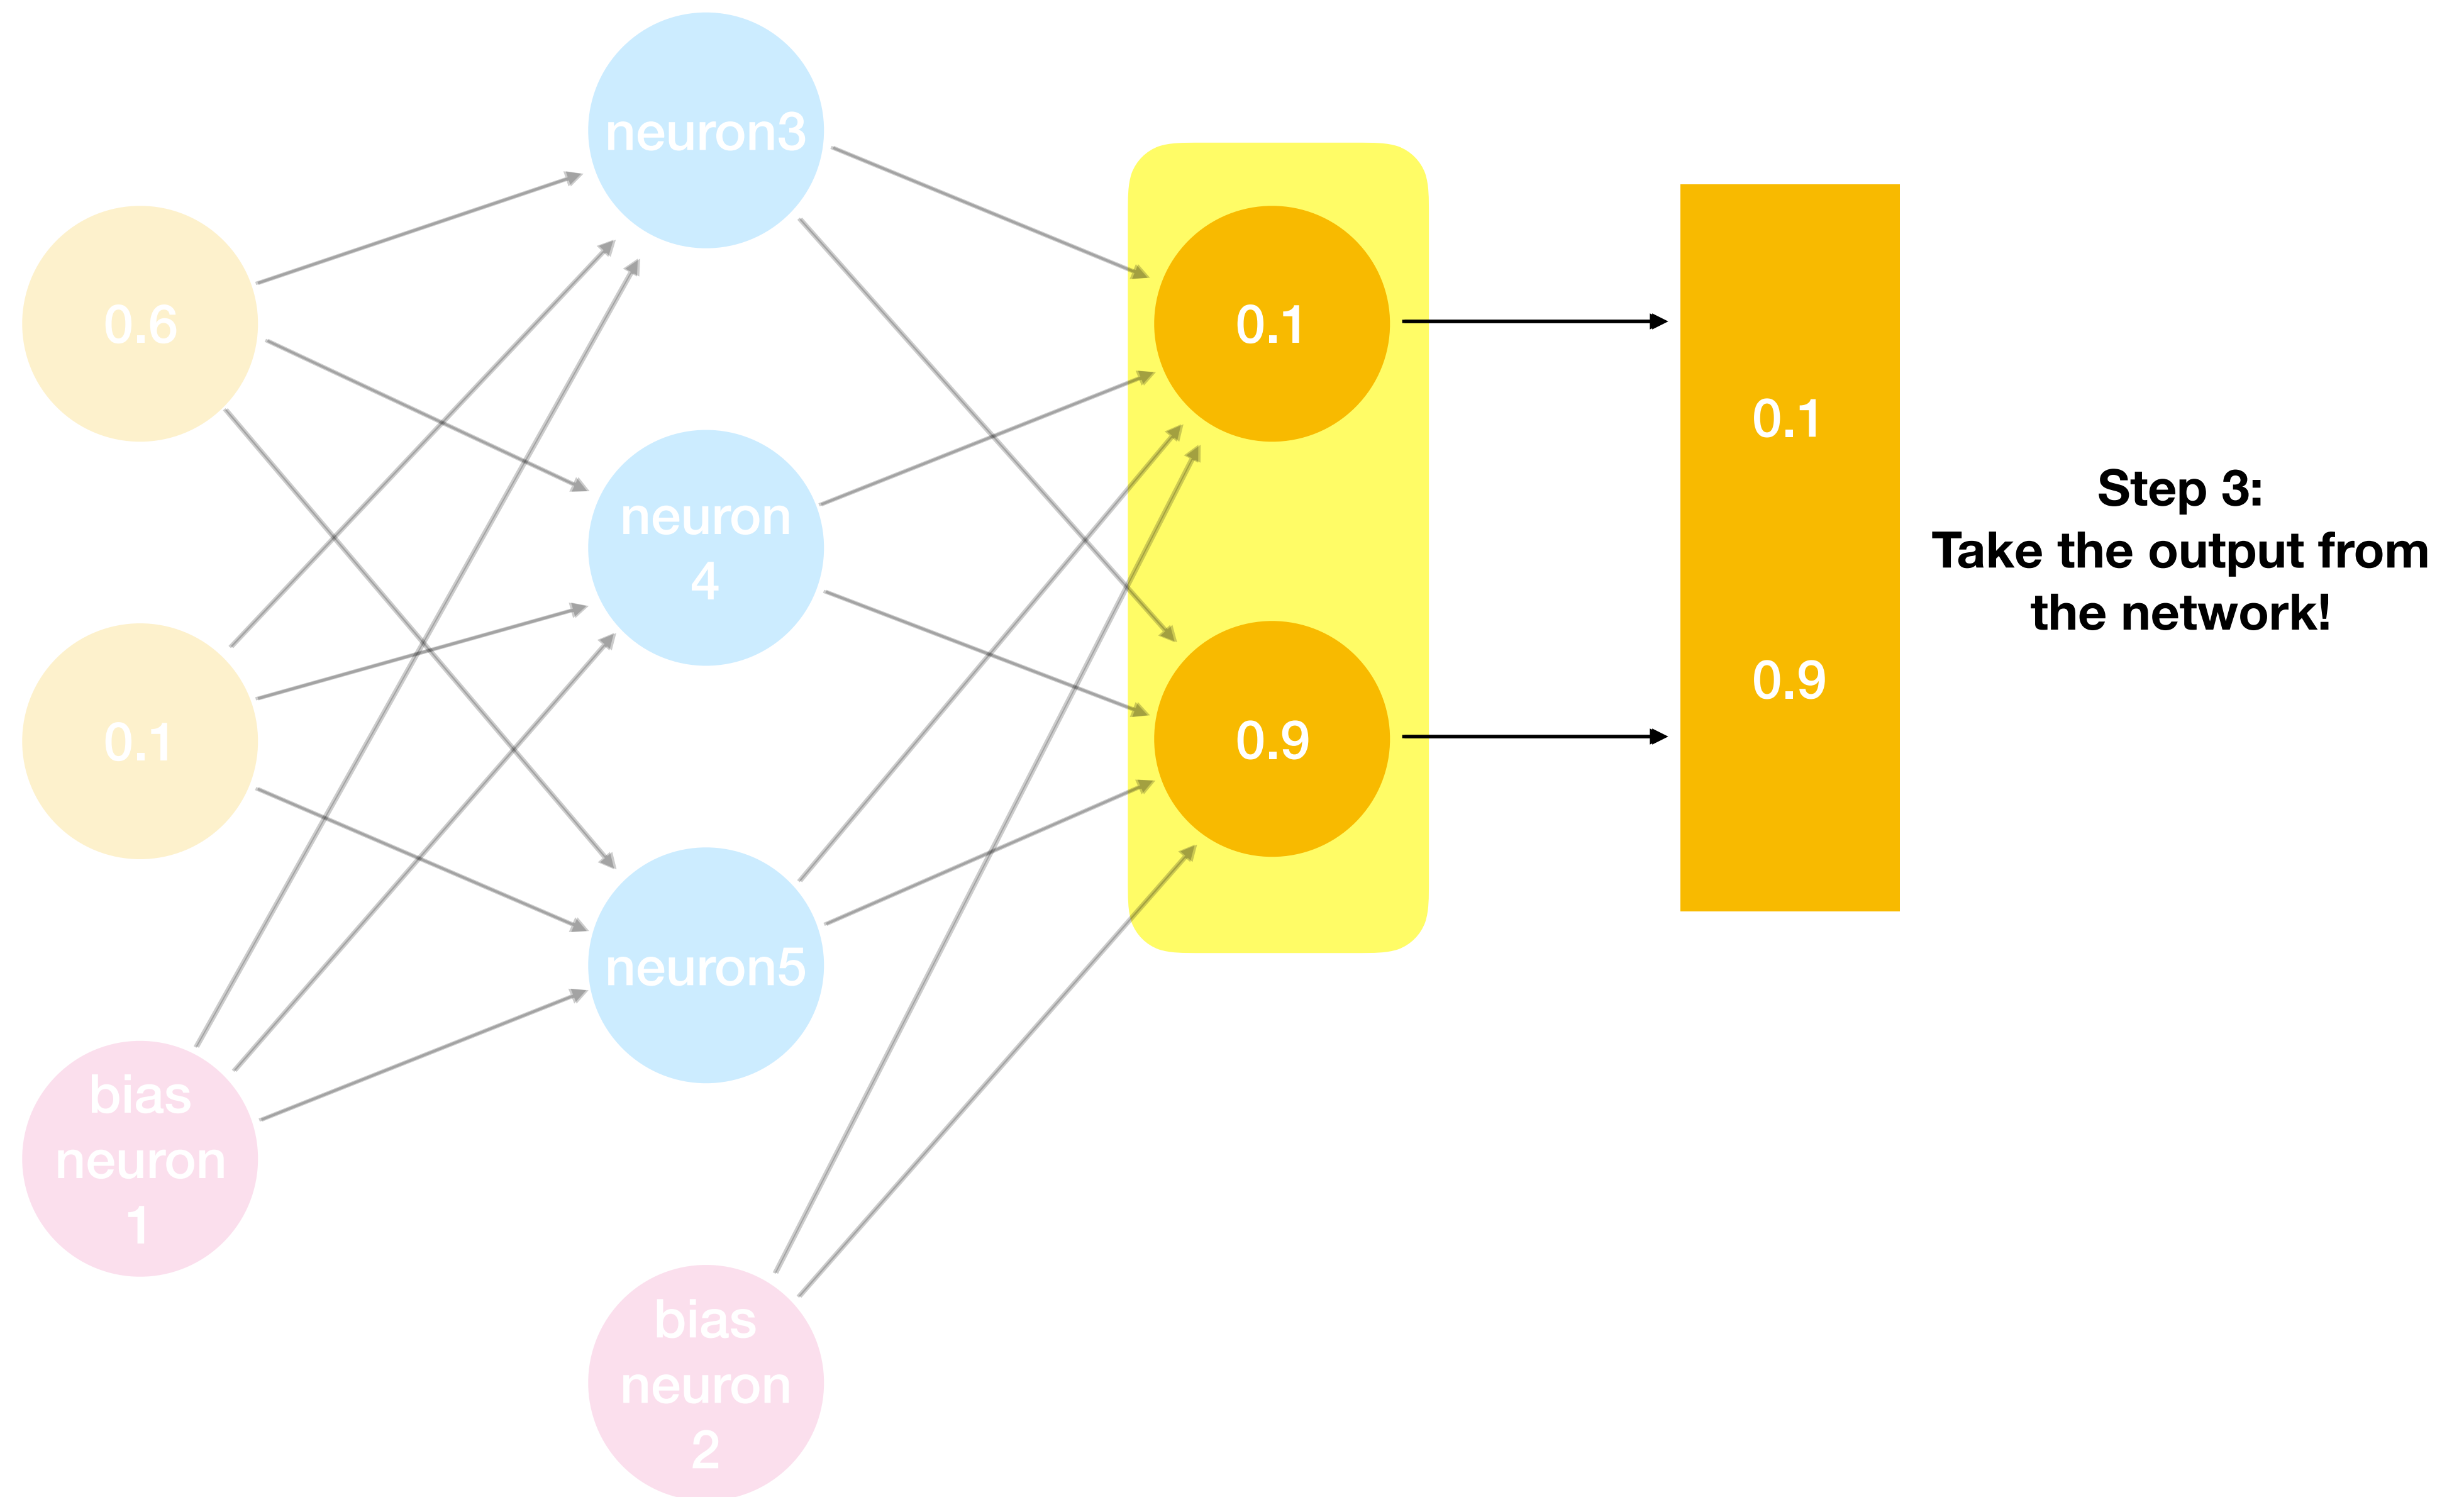
Back to the bigger picture..




Step 1:
Put some values
into the network








Training the network



The Goal

Given some input, we want to update each of the *weights* that connect neurons to each other *a wee bit* in order to reduce the difference between what we **want** the output to be, and what it **is**.




What we need

What we need

1.

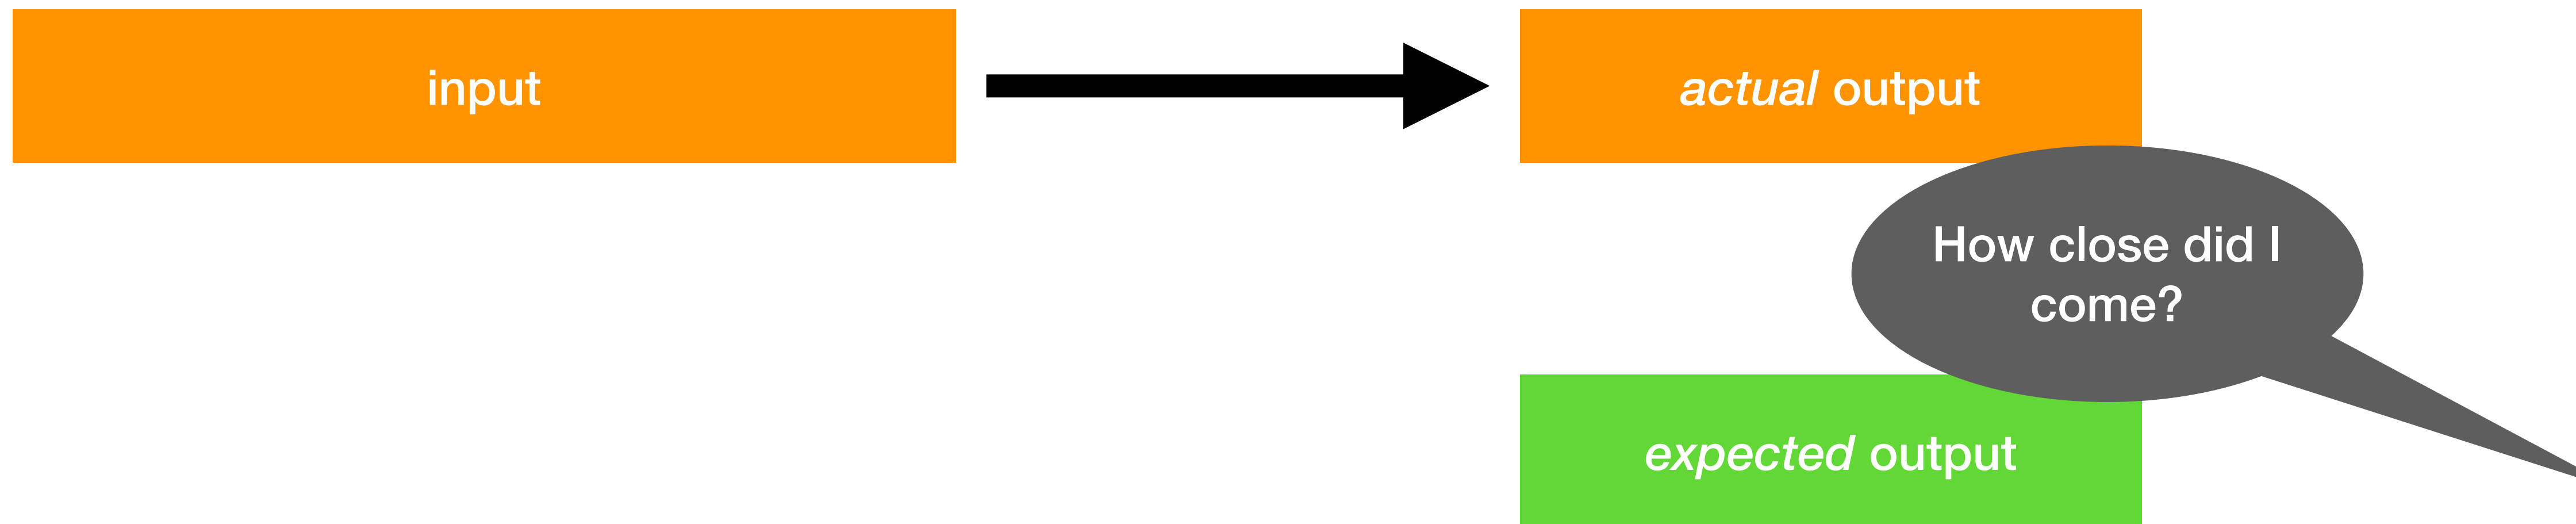
A training set. That is, a set of inputs to the network, and the output we want back for each one.



What we need

2.

A way to calculate the error given some *expected* output values and some *actual* output values



What we need

2.

A way to calculate the error given some *expected* output values and some *actual* output values. This is one possibility:

$$error = \sum \frac{1}{2}(expected - actual)^2$$

This roughly says "find the difference between each actual output value and the one we expected, and square it, then add each of these together to get our final total"

Squaring the difference means that larger differences are treated more aggressively than small ones. As the difference gets smaller, we work much less to change it. It also means that the best output value to converge on if there are 2 possibilities is one in the middle of those, since that has the smallest average error with a simple `abs()` function, arr values between the two would be equally appealing.

What we need

2.

A way to calculate the error given some *expected* output values and some *actual* output values. This is one possibility:

$$error = \sum \frac{1}{2} (expected - actual)^2$$

inputs

0, 0.5, 0.1, 0, 0, 1, 0



actual

0, 0.3, 0.9, 0.1

$$error = \frac{(0 - 0)^2 + (0 - 0.3)^2 + (1 - 0.9)^2 + (0 - 0.1)^2}{2} = 0.045$$

0, 0, 1, 0

expected




What we need

3.

A way to figure out how to update each weight in order to reduce this error. The method we'll use here is called *back propagation*.


Back Propagation

Given a training sample, for each weight we want to know how it changes with respect to the error between actual and expected output




Back Propagation

Given the current value of a weight, we can work out what the slope is at that point
(we differentiate the error with respect to one of our weights)

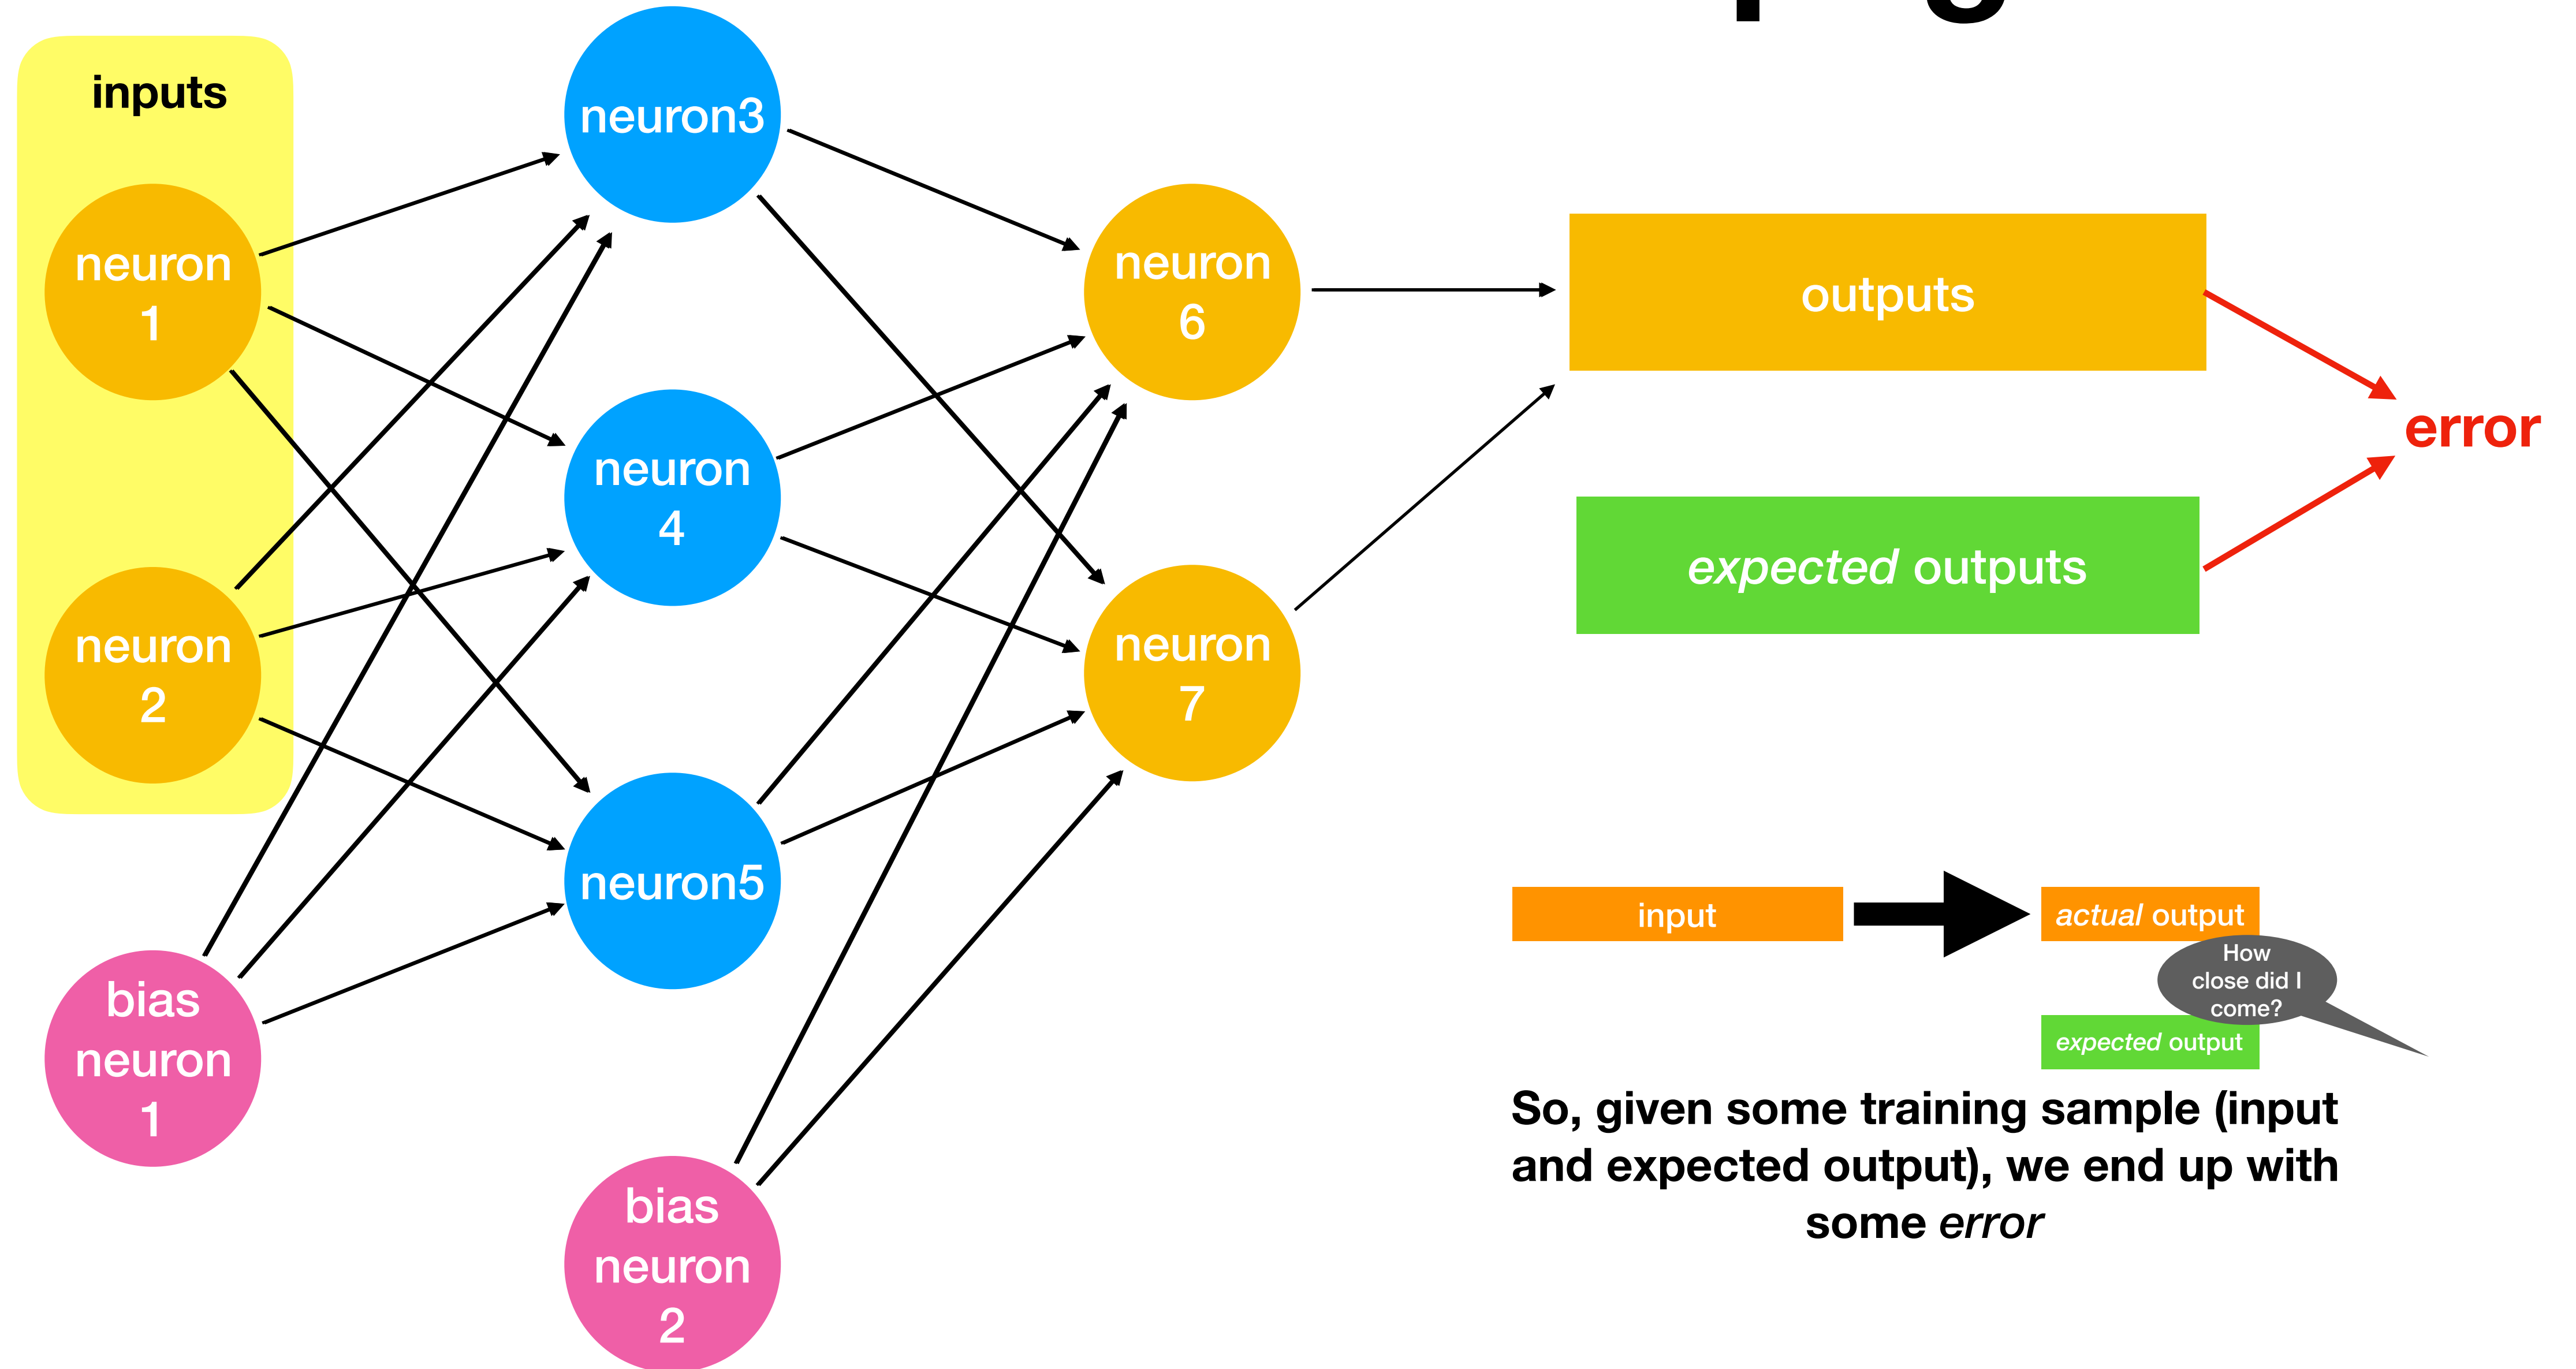


Back Propagation

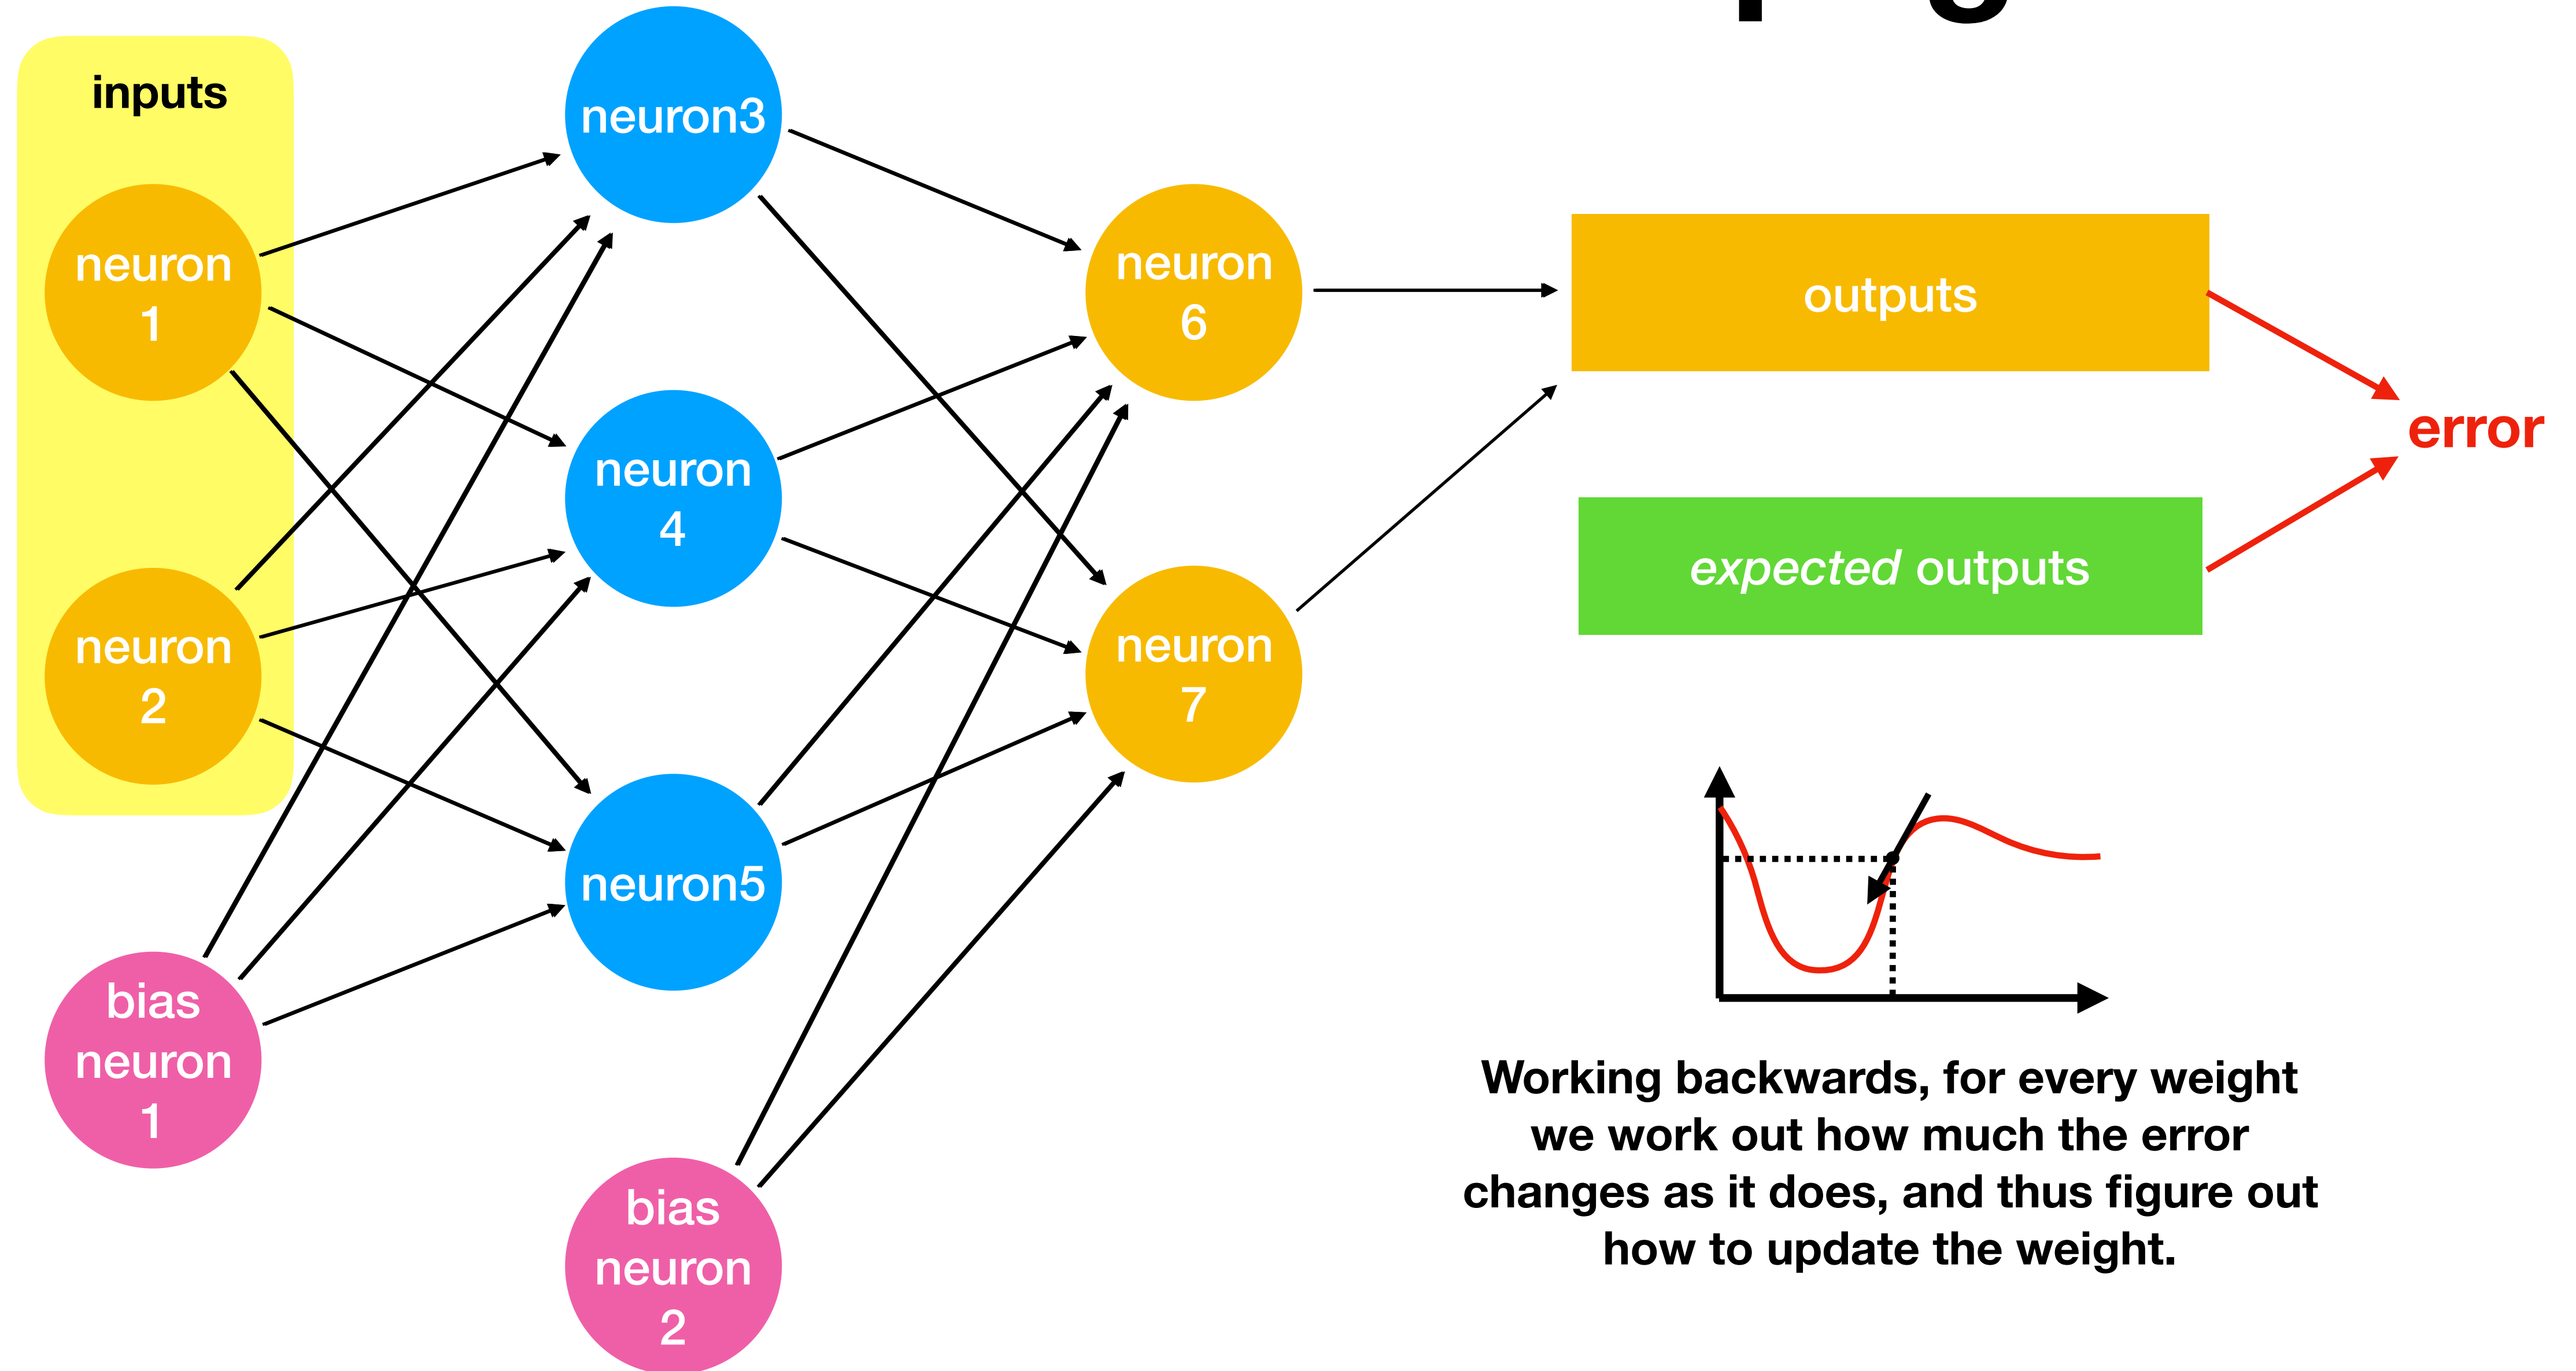
And based on this, we can decide what the new weight should be by moving a little in the direction that reduces our error



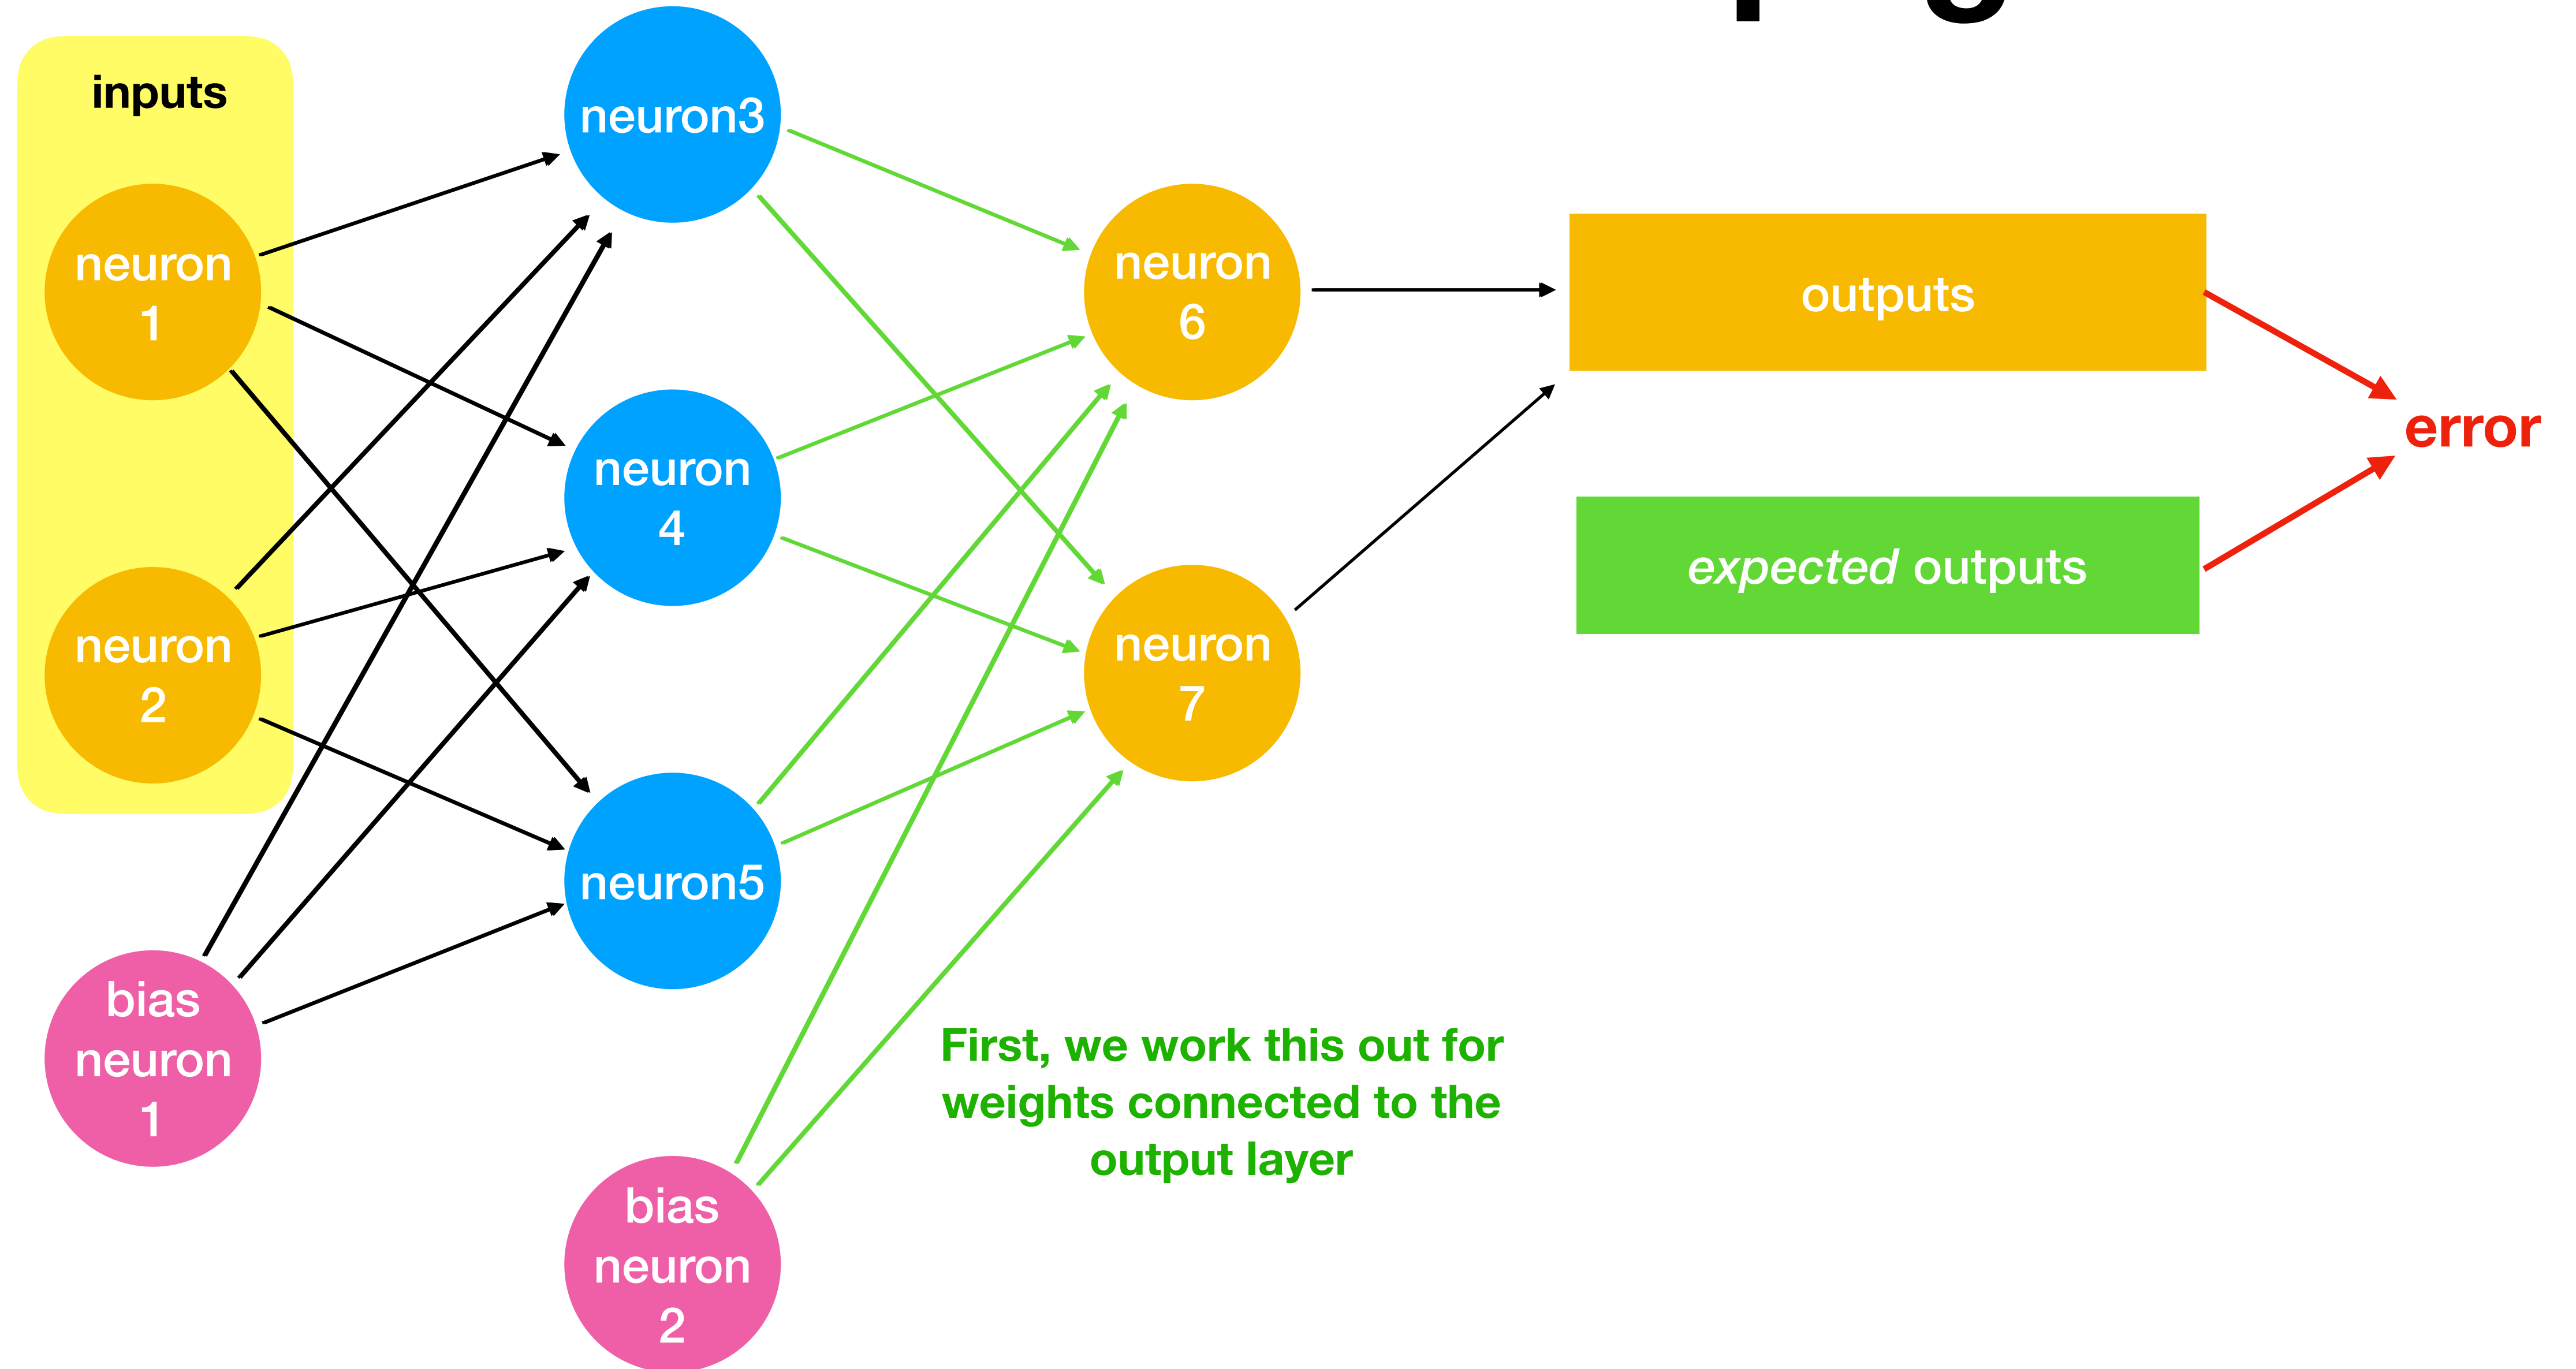
Back Propagation



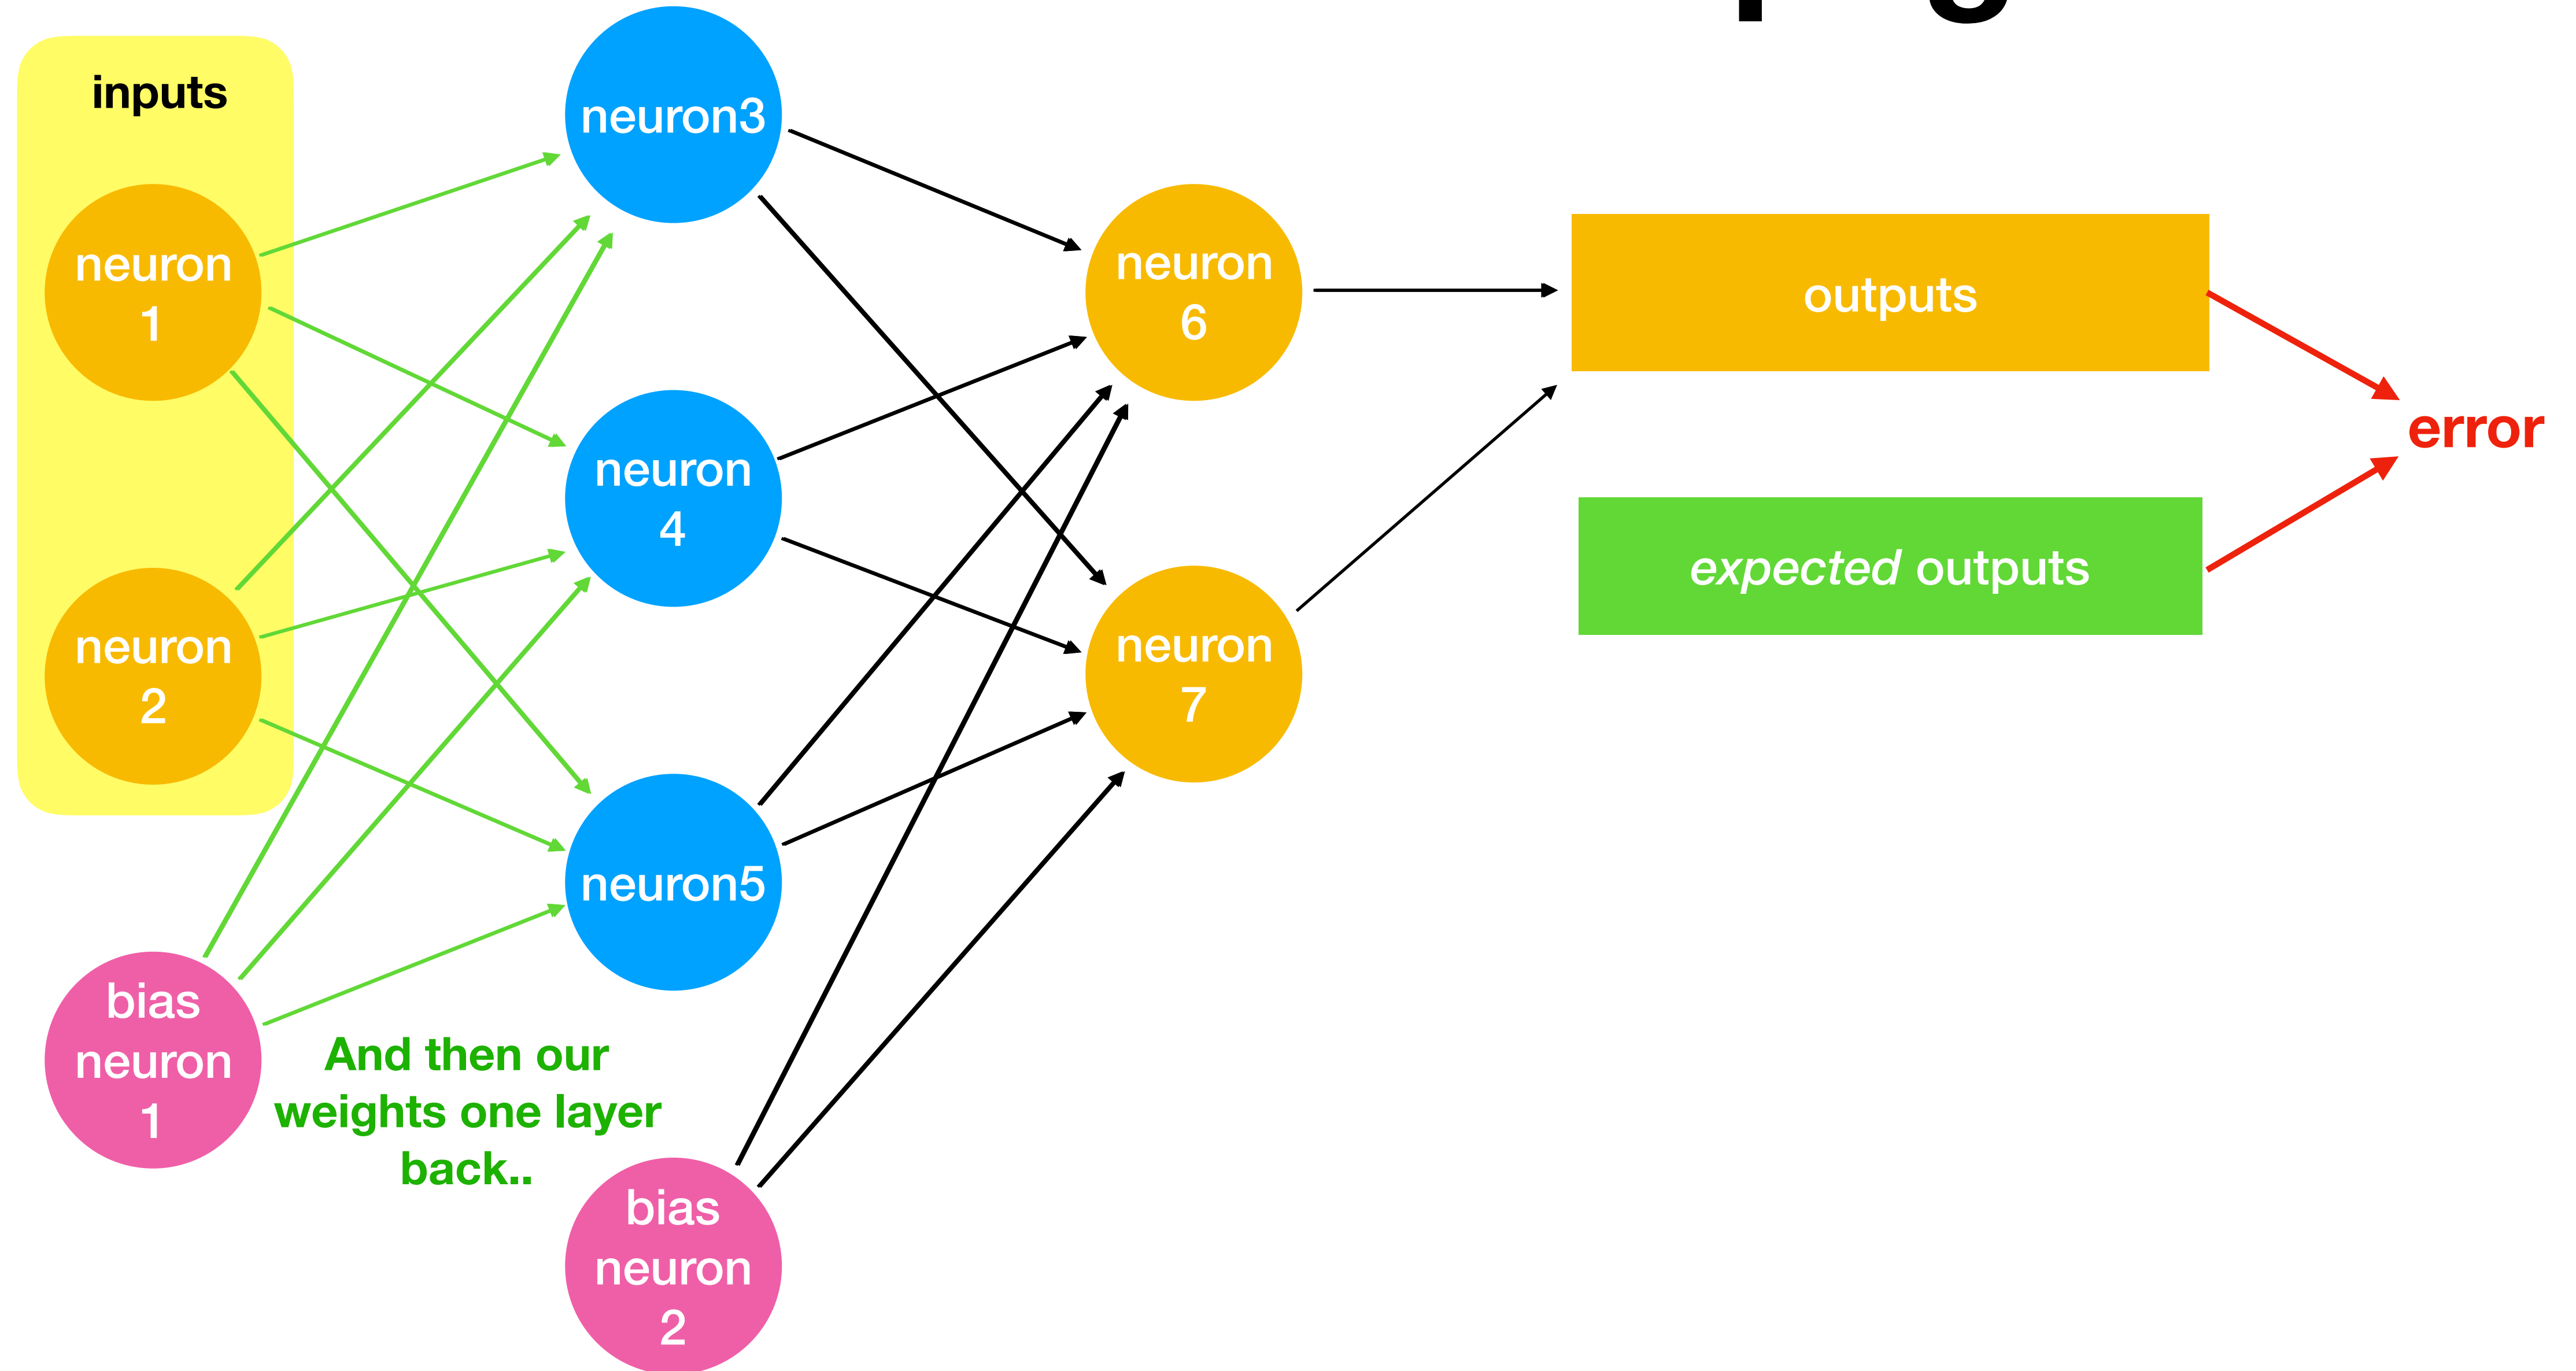
Back Propagation



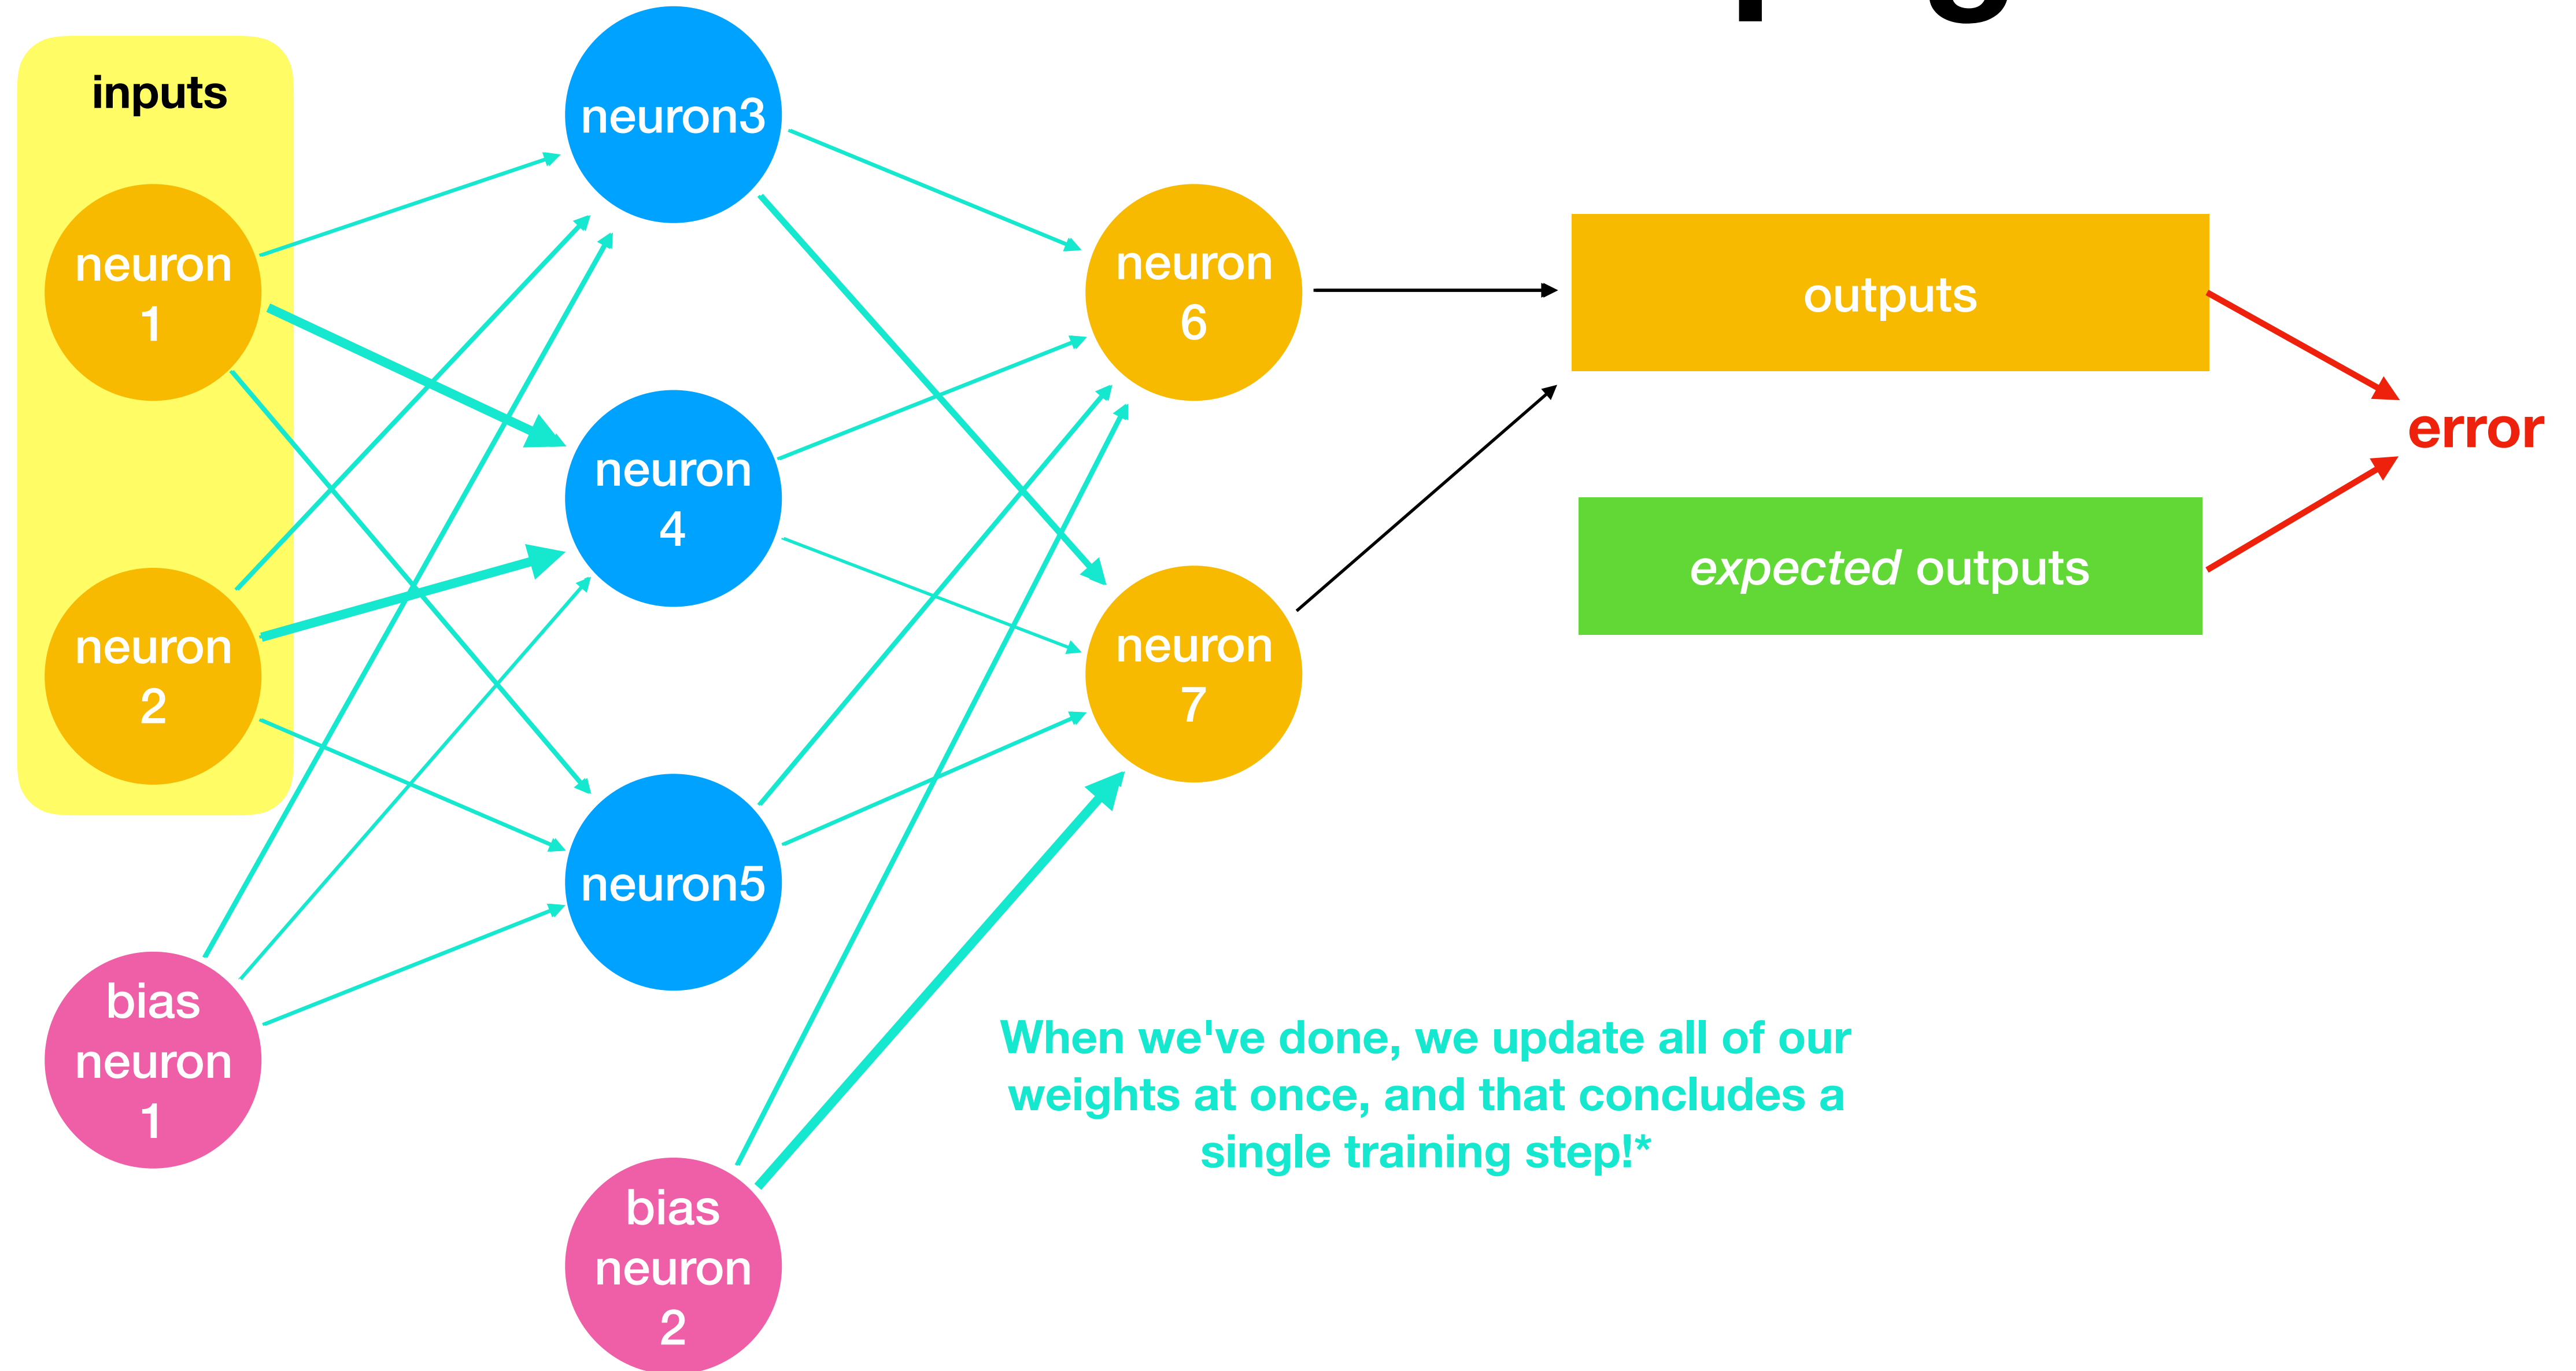
Back Propagation




Back Propagation





Back Propagation



Training



To train the network, you simply repeat this process over and over with each of the input-output pairs in your training set until the network stops improving



Demo

<http://playground.tensorflow.org>

“Deep Learning”

Is based on the exact same principles described.

The End

<https://jsdw.me/posts/neural-nets>

<https://github.com/jsdw/neural-net-example>